
The Agent Loop: A Survey of Control Strategies, Skills, and Harnesses for LLM Agents

Jungseob Lee^{1,*}

Chanjun Park²

¹Korea University, South Korea

²Soongsil University, South Korea

omanma1928@korea.ac.kr

chanjun.park@ssu.ac.kr

*Corresponding author

ABSTRACT

Every LLM agent, beneath its framework and its prompts, runs a loop: it observes, decides, acts, and observes again until it stops. This survey takes that agent loop, not the model in isolation, as its unit of analysis, on the argument that the properties practitioners actually care about, cost, reliability, and safety, are decided by the loop rather than by the model alone. We organize the field around control: the loop paradigms that shape reasoning, action, and search; the trained loops that absorb control into model weights through agentic reinforcement learning; the mechanics of termination, verification, context management, and recovery that govern any loop; the skills that externalize the loop’s competence into portable, reusable procedure; the harnesses that instantiate it; and the evaluation and safety problems it creates. One thesis runs throughout: the loop is being pulled in two directions at once, internalized into weights and externalized into skills and harnesses. These are not rivals but partially inter-convertible realizations of a single control structure, so the organizing question is not which wins but who owns the loop, and what each choice costs in reliability, transparency, and safety. Throughout, we give documented negative results the same weight as the positive claims: on intrinsic self-correction, on elaborate agents versus simple pipelines, on the reliability cost of longer loops, and on the harness confound in evaluation, supported with small controlled experiments of our own. A continuously updated companion repository, indexing every cited paper by section alongside the open-source artifacts we survey, is available at <https://github.com/js-lee-AI/awesome-agent-loop-papers>.

Keywords: large language models, autonomous agents, agent loop, control strategies, agent skills, tool use, verification, agent harness, evaluation, survey

Table of Contents

1	Introduction	4
2	Background and Definitions	5

2.1	The Loop	6
2.2	The Harness	7
2.3	The Skill	7
2.4	Competence Is Jointly Produced	8
3	Loop Paradigms	8
3.1	Interleaved Reasoning and Action	8
3.2	Plan-then-Execute and Decoupled Loops	11
3.3	Reflective and Self-Correction Loops	11
3.4	Search-in-Environment Loops	12
3.5	Adaptive, Hybrid, and Asynchronous Loops	12
4	Loop Mechanics	13
4.1	Termination and Step Budgets	13
4.2	Self-Verification versus External Checks	14
4.3	Verifiers in the Loop	15
4.4	In-Loop Context Management	16
4.5	Error Recovery, Rollback, and Human Gates	16
5	Trained Loops	17
5.1	From Prompted to Trained Loops	17
5.2	Credit Assignment and Stability	18
5.3	The Real Bottleneck Is the Environment	18
5.4	Loop Internalization	18
5.5	Controlling the Loop's Length	19
6	Skills	20
6.1	What a Skill Is, and Is Not	20
6.2	Acquisition	21
6.3	Representation, Storage, and the Retrieve–Select–Compose Pipeline	21
6.4	Ecosystem and Standards	22
6.5	Supply-Chain Security	22
7	Harnesses and Orchestration	23
7.1	The Agent–Computer Interface	24
7.2	Orchestration Frameworks and the Workflow–Agent Spectrum	24
7.3	Multi-Agent Parallelism versus Single-Threaded Context	24
7.4	Context Engineering: The Harness as Context Manager	25
7.5	From Research Scaffolds to Production Coding Harnesses	25
8	Evaluation	26
8.1	Loop-Sensitive Benchmarks	26
8.2	Harness-versus-Model Attribution	26
8.3	Reliability beyond pass@1	28
8.4	Per-Loop Cost Accounting	28
8.5	Skill-Reuse Evaluation	31
9	Safety of the Loop	32
9.1	The Injection-to-Action Threat Model	32

9.2	Model-Level Defenses	33
9.3	Defense by Design: Sandboxing, Capabilities, and Firewalls	33
9.4	Supply-Chain Attacks on Tools and Skills	34
9.5	Watching and Stopping the Loop	36
10	Open Challenges and Future Directions	37
10.1	Who Owns the Loop?	37
10.2	Loop Reliability over Long Horizons	38
10.3	Skill Governance and the Agent Supply Chain	38
10.4	Standardization, Interoperability, and Reproducible Evaluation	38
10.5	Cost-Aware Design and the Safety–Capability Co-Evolution	38
11	Conclusion	39
A	Example Loop Traces	39
A.1	Interleaved (ReAct)	39
A.2	Plan-then-Execute (ReWOO)	40
B	Example Skill-Injection Attack	40




FIGURE 1: The two directions of §1, read off the literature itself. Each band is one of five loop- and skill-related arXiv query groups, and its thickness is that group’s monthly count (3-month centred mean, 2022 to mid-2026). The chart is a diverging area: the hand-designed *agent loop* band is drawn symmetric about the zero axis, and the other bands stack *outward* from its edges, pulled up into externalized scaffolding (harness and orchestration, verification and self-correction, agent skills) and down into the model’s weights (trained loops, agentic RL). Both forces swell while the classic loop stays a thin steady stream, so what the figure shows is the compositional shift, not merely the volume. Because the query groups favour precision over recall and may overlap, every count is a floor, not an exact tally.

1 Introduction

Every LLM-based agent, underneath its framework and its prompts, is a *loop*: the model observes, decides on an action, the action is executed against an environment or a tool, the result is fed back, and the cycle repeats until some stopping condition is met. ReAct named this loop and made it the default, interleaving a reasoning trace with each action so that every thought is conditioned on the freshest observation (Yao et al. 2022). The three years since have not replaced that loop so much as elaborated it: plans are committed up front and executed without re-observation (Xu et al. 2023; Zhou et al. 2022; Wang et al. 2023e), failed attempts are retried after a verbal self-reflection (Shinn et al. 2023; Renze et al. 2024), and the single trajectory is generalized into a search over many (Zhou et al. 2024a). These are not competing frameworks; they are competing *loop shapes*, and which shape an agent runs is a design decision as consequential as which model it is built on.

This survey treats the agent loop as a design space in its own right. The motivation is that the loop is where the properties practitioners actually care about are decided. Cost is a property of the loop: every extra step, every reflection, every search rollout is paid for in tokens, and the compute-blind comparisons that inflate the literature hide exactly this price (Kapoor et al. 2024; Kapoor et al. 2025). Reliability is a property of the loop: a per-step error probability compounds over a trajectory, so a loop that runs longer without a verification gate becomes less trustworthy, not more (Ma et al. 2024). And safety is a property of the loop: an injected instruction becomes a real side effect only because the loop will execute it, so the loop is simultaneously the source of an agent’s usefulness and of its attack surface (Beurer-Kellner et al. 2025). A survey organized around models, frameworks, or application domains cannot make these properties commensurable; one organized around the loop can.

Read across the recent literature, the loop is being pulled in two opposite directions at once, and

that tension is the thesis this survey defends. In one direction the loop is being *internalized* into the model’s weights: rather than wrapping a fixed model in a hand-designed scaffold, agentic reinforcement learning trains the model to run the loop itself, deciding when to search (Jin et al. 2025; Nakano et al. 2021), when to invoke a code interpreter (Feng et al. 2025a), or how to operate a computer end to end (Qin et al. 2025; Wang et al. 2025d; Anthropic 2024b), on top of reasoning-native models whose deliberation was already induced by outcome-only RL (DeepSeek-AI et al. 2025; Kimi Team 2025). In the other direction the loop’s knowledge is being *externalized* into portable, reusable *skills*: an agent that once had to rediscover a procedure every episode now writes it to a skill library and retrieves it later (Wang et al. 2023a), distills it into transferable insights (Zhao et al. 2023; Ouyang et al. 2025; Zheng et al. 2023b), or loads it from a shared, versioned artifact through an emerging industry standard (Anthropic 2024c; Yang et al. 2025). The organizing question of this survey is the one these two movements pose together: *who owns the loop*: the weights that were trained to run it, the harness that orchestrates it, or the skill file that supplies its competence, and what each choice costs in reliability, transparency, and safety.

Our contributions are as follows.

- **A loop-first taxonomy** that organizes the agent literature by control strategy rather than by framework: loop paradigms (§3), the mechanics that govern any loop (termination, verification, context management, recovery, human gates; §4), and the trained loops that absorb these into weights (§5). Figure 3 maps the full taxonomy on one page, seven pillars each closing with its own documented counter-result.
- **A first-class treatment of skills** (§6), spanning the academic lineage of skill libraries and the industry ecosystem of portable skill standards, which prior agent surveys fold into generic “tool use” (Qu et al. 2024) or “memory” (Zhang et al. 2024c).
- **The harness as an object of study** (§7): the agent-computer interfaces and orchestration layers whose engineering can shift an agent’s capability as decisively as the choice of base model, including the production coding agents that now define the state of practice.
- **Honest both-sides synthesis**: throughout, documented negative results, on intrinsic self-correction (Huang et al. 2023; Tie et al. 2025), on elaborate agents versus simple pipelines (Xia et al. 2024), and on the reliability cost of longer loops, are given the same weight as the positive claims, and supported with small controlled experiments of our own.

A note on method. Figure 1 plots the growth this survey must cover; the corpus behind it was assembled through eight subtopic sweeps over the loop-paradigm, trained-loop, loop-mechanics, skill, harness, evaluation, and safety literatures, each candidate verified against its primary source, and cross-referenced against a companion general survey of LLM agents so that this survey can concentrate on the loop and skill dimensions that the general survey treats only in passing. A continuously updated companion repository indexes every cited paper by section alongside the real-world open-source artifacts of §7.¹

2 Background and Definitions

The literature this survey covers is built on a vocabulary that no two papers use the same way. “Agent,” “scaffold,” “harness,” “framework,” and “skill” are applied to overlapping and sometimes contradictory referents, and the boundary between the model and the code that surrounds it is drawn in a different place by every benchmark. This is not a pedantic complaint: it is precisely the terminological slack that lets a gain produced by better surrounding code be reported as a gain in model capability. Before we can organize the field around the loop, we must fix three terms (the *loop*, the *harness*, and the *skill*)

¹ <https://github.com/js-lee-AI/awesome-agent-loop-papers>




FIGURE 2: The perceive–reason–act cycle and the three terms this survey fixes: harness, policy, and skill library. The *harness* (dotted) is the code that realizes the loop: it runs the cycle, calls the *policy* (the model’s weights) at the reason step, mediates the environment and tools at the act step, and owns the stop/continue gate. The *skill library* is an external, portable artifact fed into the loop, distinct from both the weights and the harness. The two directions of §1 (internalizing competence into the policy, externalizing it into skills) are movements between these boxes; the loop itself stays fixed, which is what makes the two commensurable.

and state the claim that follows from them: an agent’s competence is produced jointly by its model, its harness, and its skills, so the loop, not the model in isolation, is the object this survey studies.

We inherit the framing of cognitive architectures for language agents, which recast an LLM agent as a policy operating inside a structured decision cycle with distinct memory, action, and control components rather than as a monolithic “brain” (Sumers et al. 2023). Our three definitions sharpen the control and action parts of that picture into terms that make the model-versus-scaffold attribution problem visible. A concurrent 2026 position argues this control-cycle view deserves a science of its own, an *agent cybernetics* that formalizes the long-horizon control loop of foundation agents as the primary object of study (Wang et al. 2026k). Figure 2 shows how they fit together.

2.1 The Loop

THE AGENT LOOP

The *agent loop* is the iterated control cycle (PERCEIVE an observation, REASON about it, ACT on the environment or a tool, OBSERVE the result) that an LLM agent repeats until a termination condition is met. The loop is a *control structure*: it specifies what is fed to the model, when the model is called, how its output is turned into an effect, and when the iteration stops. It is distinct from the *policy* (the model’s weights) that it invokes at each turn.

ALGORITHM I — THE AGENT LOOP (ONE EPISODE)

input: goal g , policy π (weights), skill library S , harness H

$ctx \leftarrow H.INIT(g)$

% harness builds the initial context

REPEAT

$o \leftarrow H.PERCEIVE()$

% observation from env / tools

$ctx \leftarrow H.ASSEMBLE(ctx, o, S)$

% retrieve + splice in skills

$a \leftarrow \pi(ctx)$

% REASON: one policy call

IF $H.STOP(a)$ **THEN BREAK**

% stop/continue gate, owned by H

$r \leftarrow H.ACT(a)$

% ACT: execute against env / tool

$ctx \leftarrow H.OBSERVE(ctx, a, r)$

% fold the result back in

RETURN $H.FINALIZE(ctx)$

This is the object the survey studies: the policy π is called once per turn (the REASON step), the harness H owns everything else (context assembly, action, the stop gate), and the skill library S enters only where context is assembled. The paradigms of §3 vary the body of this loop, the trained loops of §5 fold it into π , and the harnesses of §7 are concrete implementations of H .

The loop was named and made canonical by ReAct, which interleaves a free-form reasoning trace with each action so that every thought is conditioned on the most recent observation (Yao et al. 2022). Subsequent work did not discard this inner cycle so much as wrap it: Reflexion adds an *outer* loop that, on failure, produces a verbal self-reflection and retries the whole episode (Shinn et al. 2023); plan-then-execute schemes commit to a plan before acting and loosen the coupling between reasoning and observation (Xu et al. 2023); and search-based agents lift the single trajectory into a tree explored with backtracking (Zhou et al. 2024a). These are different *shapes* of the same control structure, and §3 treats them as a design space. The key move of our framing is to separate the loop from the policy: the same ReAct-shaped loop can be run by a frozen model driven by a prompt, or absorbed into a model trained to emit the reason-act-observe pattern natively (§5). Holding the loop fixed while varying where it lives is what makes the two directions of §1 commensurable.

2.2 The Harness

THE HARNESS

The *harness* is the code and infrastructure surrounding the model that realizes the loop: it constructs the context passed to the model at each turn, mediates the model’s access to tools and the environment, orchestrates the sequence of steps, verifies or filters outputs, and decides when to stop. Everything an agent does that is not a forward pass through the model happens in the harness.

Where the loop is the abstract control structure, the harness is its concrete implementation, and the two are easy to conflate because a harness is what one actually runs. The engineering discipline of harness design was crystallized by Anthropic’s distinction between *workflows* (LLM calls orchestrated through predefined code paths) and *agents*, where the model itself directs the control flow, together with a small catalog of recurring patterns (prompt chaining, routing, parallelization, orchestrator-workers, evaluator-optimizer) (Anthropic 2024a). That the harness is a first-class determinant of capability, not neutral plumbing, was shown sharply by SWE-agent: a purpose-built agent-computer interface (constrained commands, linted feedback, guardrailed edits) changes task success more than swapping the base model does (Yang et al. 2024a). This is exactly why the model/harness boundary matters. Because a reported “agent” score is a joint measurement of a model and an often-undisclosed harness (Kapoor et al. 2024), and because every benchmark draws that boundary differently, a capability improvement in the harness can be, and routinely is, reported as an improvement in the model. Making the harness an explicit object of study (§7) rather than an implementation detail is a precondition for honest attribution (§8).

2.3 The Skill

THE SKILL

A *skill* is a portable, named, discoverable unit of procedural knowledge that an agent can acquire, store, retrieve, and compose: reusable know-how that persists across episodes, sitting between a stateless external *tool* call and an ephemeral in-context *prompt*. A skill may be represented as executable code, as a natural-language strategy, or as a package of instructions and resources loaded on demand.

The skill is the newest and least settled of the three terms, and the survey treats it as first-class (§6)

precisely because prior agent surveys dissolve it into generic “tool use” or “memory.” Its academic lineage begins with Voyager, whose agent writes successful behaviors as executable programs into a growing library and retrieves them later, so that competence accumulates instead of being rediscovered each episode (Wang et al. 2023a). CodeAct generalizes the insight that an agent’s actions are most naturally expressed as executable code, blurring the line between calling a tool and invoking a stored procedure (Wang et al. 2024b). In parallel, industry has converged on a portable packaging: Anthropic’s Agent Skills express a skill as a folder of Markdown instructions, scripts, and resources loaded through progressive disclosure, distributed as a versioned artifact (Anthropic 2025f). The term is genuinely contested (Voyager’s skill is Python, Anthropic’s is Markdown-plus-context, and an experiential learner’s is a natural-language insight distilled from comparing trajectories (Zhao et al. 2023)), and we will not paper over that with a tidy taxonomy. What survives across the three is functional, not representational: a skill is *named*, *reusable*, *retrievable*, and *composable*. This is what distinguishes it from a *tool*, which is a stateless external capability the model invokes but does not own (Schick et al. 2023), exposed increasingly through a shared connectivity protocol (Anthropic 2024c); from a *prompt*, which is ephemeral and not stored across episodes; and from raw *memory*, which is retrieved state rather than executable procedure.

2.4 Competence Is Jointly Produced

The three definitions combine into the claim that organizes the survey. An agent’s observed competence on a task is not a property of its model alone; it is jointly produced by the model, the harness that constructs its context and mediates its actions, and the skills that supply reusable procedure. The identical control loop can be realized by internalizing it into the model’s weights through agentic reinforcement learning, or by holding the model fixed and externalizing competence into a harness and a library of portable skills, the two directions introduced in §1, which recent work shows are partially inter-convertible rather than rivals. Current benchmark practice, however, reports a single number and attributes it to the model, while disclosing the harness only partially and the skills not at all. A framework organized around models, frameworks, or application domains cannot separate these contributions; one organized around the loop, with its paradigms (§3), the trained loops that absorb it into weights (§5), the mechanics that govern any loop (§4), the skills that feed it (§6), and the harnesses that run it (§7), can. Fixing the loop as the unit of analysis is therefore not a stylistic choice but the only cut that makes model-versus-harness attribution, and weights-versus-skills convertibility, visible at all. Table 1 makes the gap concrete: prior surveys are thinnest exactly on the loop-centric axes (paradigms, mechanics, trained loops, and the harness) that this survey foregrounds.

3 Loop Paradigms

If the loop is a control structure, then its *shape* (how reasoning, action, and observation are ordered and how the trajectory branches) is the first design decision an agent makes. This section organizes the shapes into five paradigms and defends a single thesis: the shape of the loop is a bet about the environment, and each richer topology pays off only when its enabling assumption (reliable feedback, cheap and reversible actions, or a trustworthy value signal) holds. Where the assumption fails, the added reasoning is not free: it is paid for in tokens, latency, and new failure modes, so the frontier is shifting from committing to one loop toward learning which loop to run. Table 2 summarizes the five paradigms as a design space; the subsections that follow defend each row.

3.1 Interleaved Reasoning and Action

The default agent loop interleaves a reasoning step before each action, so that every decision is conditioned on the freshest observation. ReAct established this as the standard because grounding each thought in live feedback curbs the hallucination and error propagation that afflict a plan generated in



FIGURE 3: The field map of this survey. The agent loop (top) is organized into seven pillars, one per section: the control pillars (loop paradigms §3, loop mechanics §4), the internalize direction (trained loops §5), the externalize direction (skills §6, harnesses §7), and the accountability pillars (evaluation §8, safety §9); the left ribbon marks these four groups and the two directions of §1. Each leaf names a sub-area with representative systems, all of which are cited and discussed in the corresponding section, and each pillar closes with one of its own documented counter-results (dashed strips), reflecting the survey's commitment to giving negative findings equal weight.

TABLE 1: How prior surveys cover the agent loop, scored across the seven axes this survey organizes around. A *filled* circle marks an organizing pillar with dedicated sections, a *half* circle a topic discussed but not made a pillar, and an *empty* circle a topic absent or mentioned only in passing; “Cov.” is the latest literature year substantively covered. Prior surveys are thinnest exactly on loop paradigms, loop mechanics, trained loops, and the harness—the axes this survey foregrounds.

Survey	Loop parad.	Loop mech.	Trained loops	Skills	Har-ness	Eval.	Safety	Cov.
Xi et al. (2023)	●	○	○	○	○	●	●	2023
Wang et al. (2023b)	●	○	○	○	○	●	○	2023
Sumers et al. (2023)	●	●	○	●	●	○	○	2023
Huang et al. (2024)	●	●	○	○	○	○	○	2024
Masterman et al. (2024)	●	○	○	○	●	○	○	2024
Zhang et al. (2025b)	●	●	●	●	○	●	●	2025
Fang et al. (2025)	○	○	●	●	○	●	○	2025
Yehudai et al. (2025)	●	●	○	○	●	●	●	2025
Plaat et al. (2025)	●	○	●	○	○	●	●	2025
Zhou et al. (2026a)	○	○	●	●	●	●	●	2026
This survey	●	●	●	●	●	●	●	2026

TABLE 2: The five loop paradigms as a design space. Each richer topology relaxes the reason-before-every-action default and buys capability only where its enabling assumption holds; where it fails, the added structure is paid for in the cost column and manifests as the failure mode. The whole section is the defense of this table.

Paradigm	Assumption	Cost	Failure mode	\$
Interleaved	Each observation informs the next action	Reasoning tokens every step; unbounded context	Uninformative observations derail the trace	\$3.1
Plan-then-execute	Environment predictable enough to plan blind	Loss of mid-plan reactivity	Brittle when an early step is wrong or a tool fails	\$3.2
Reflective	Reliable external feedback to reflect on	Extra full-episode retries	Unanchored self-critique degrades correct answers	\$3.3
Search-in-env.	Cheap, reversible actions + trustworthy value signal	Large inference / latency multiplier	Noisy value function; irreversible actions	\$3.4
Adaptive / hybrid	A reliable signal for which shape to run	Meta-controller overhead + new hazards	Mis-triggered plan or react; async side effects	\$3.5

one shot (Yao et al. 2022). The tight coupling is genuinely useful, but it is not a universal free lunch, and the honest reading starts with ReAct’s own ablations: on knowledge-reasoning tasks such as HotpotQA and Fever, ReAct underperforms plain chain-of-thought when uninformative observations derail the reasoning trace, and it needs a CoT self-consistency hybrid to recover (Yao et al. 2022). The coupling has two structural costs. First, reasoning before every action is token-expensive and grows the context without bound, so a long interleaved trajectory becomes both slower and, through context dilution, less reliable. Second, later analysis questions whether the interleaving is even the source of the gains: one study argues ReAct’s improvements track exemplar–query similarity and approximate retrieval rather than any benefit of intertwining reasoning with acting, exposing how brittle the default loop is to prompt construction (Verma et al. 2024). Reason-before-every-action is a strong default, not a settled optimum, and subsequent paradigms can each be read as a targeted relaxation of it.




FIGURE 4: The loop zoo: six loop shapes drawn on one shared notation (legend at bottom). (a) Interleaved reasoning and action reasons before every step (Yao et al. 2022). (b) Plan-then-execute produces one up-front plan, executed without re-observation, optionally as a parallel DAG (Xu et al. 2023; Kim et al. 2023). (c) Reflective loops convert a failed episode into stored verbal advice and retry the whole episode (Shinn et al. 2023). (d) Search-in-environment branches, evaluates, and backtracks (Zhou et al. 2024a). (e) Trained loops absorb the same interleaved cycle into the weights and optimize it against outcome reward (Jin et al. 2025; Feng et al. 2025a). (f) Adaptive and asynchronous loops let a controller pick the shape at runtime and overlap thinking with acting (Sun et al. 2023; Paglieri et al. 2025; Ginart et al. 2024). Shapes (b) through (f) are each a relaxation of (a)’s reason-before-every-step default; what each relaxation assumes and costs is Table 2.

3.2 Plan-then-Execute and Decoupled Loops

The first relaxation decouples reasoning from observation by committing to a plan before acting. ReWOO writes the full reasoning plan up front and then executes its tool calls against it, removing the interleaved observations from the prompt and cutting token usage several-fold while improving long-horizon plan quality (Xu et al. 2023). Because a plan makes the dependency structure explicit, it can also expose parallelism: LLMCompiler compiles the plan into a directed acyclic graph of function calls and dispatches independent calls concurrently, reducing latency and cost relative to sequential ReAct (Kim et al. 2023), and graph-structured planners push this further by learning which subtasks to serialize and which to run in parallel (Wu et al. 2025). Decoupling trades interactivity for efficiency, and the trade is only sound when the environment is predictable enough that a plan written without observations survives contact with it. It frequently is not: a plan is brittle when an early assumption is wrong or a tool fails, which is the explicit motivation for closed-loop refinement (§3.5), and the plan-then-execute paradigm’s own recent flagship finds that LLMs are not inherently reliable planners and needs a synthetic plan-annotation training pipeline to make the decoupled planner dependable (Erdogan et al. 2025). Earlier decompositional schemes (Plan-and-Solve prompting (Wang et al. 2023c) and as-needed decomposition (Prasad et al. 2023)) mark the same tension between amortizing reasoning and staying responsive to feedback.

3.3 Reflective and Self-Correction Loops

The second relaxation adds an outer loop: generate, then reflect on the result and retry. Reflexion converts a failure signal into a verbal self-reflection stored in memory and used to improve the next attempt

(Shinn et al. 2023); Self-Refine iterates generation with the model’s own feedback (Madaan et al. 2023); and CRITIC refines outputs using tool-interactive critiques rather than the model’s unaided judgment (Gou et al. 2023). Reported gains are large, but this is the paradigm where the both-sides accounting matters most, because the gains are real only when the reflection is grounded in reliable external feedback. The strongest negative result is direct: intrinsic self-correction, in which a model critiques and revises with no external signal, often *degrades* reasoning accuracy rather than improving it (Huang et al. 2023), and a critical survey of the self-correction literature finds no convincing success outside tasks with a reliable external verifier or an oracle stopping criterion (Kamoi et al. 2024). A complementary finding explains the mechanism: LLMs are not reliably better at *discriminating* among their own candidate outputs than at generating them, which undercuts the core premise of self-evaluate-then-refine (Jiang et al. 2024). Read honestly, then, the headline Reflexion and Self-Refine results lean on environment reward, tool feedback, or oracle labels; CRITIC’s own thesis, that external tools are the necessary ingredient, is the constructive statement of the same point, and world-grounded successors that reflect on the agent’s state relative to the goal rather than on free-form self-judgment report further gains (Kim et al. 2025). Reflection helps when it is anchored to something outside the model; unanchored, it is as likely to talk the agent out of a correct answer as into one.

3.4 Search-in-Environment Loops

The third relaxation abandons a single linear trajectory for explicit search over actions and states. LATS unifies reasoning, acting, and planning by running Monte-Carlo tree search with an LM value function and self-reflection, turning the loop into a deliberate explore–evaluate–backtrack procedure (Zhou et al. 2024a); tree search run inside a real web environment gives an agent inference-time exploration on realistic tasks (Koh et al. 2024); and Agent Q combines guided search with a self-critique evaluator and distills the search traces back into the policy (Putta et al. 2024). These methods substantially raise success on hard, multi-step web and reasoning tasks, and their lineage runs back through Tree-of-Thoughts (Yao et al. 2023a) and reasoning-as-planning (Hao et al. 2023). But search buys success with three strong assumptions, and the honest treatment names all of them. It multiplies inference cost and latency by large factors; it depends on a value or reward signal that is frequently noisy, since the same model that cannot reliably discriminate its own answers is often the one scoring the nodes; and it assumes a resettable, backtrackable environment, an assumption that fails for irreversible real-world actions, which is exactly why the in-environment search of Koh et al. (2024) is restricted to safely reversible web operations. Reported gains also show diminishing returns as the rollout budget grows, which is what motivates the budget-aware stopping criteria we take up in §4. Search is powerful where actions are cheap and reversible and a value signal is trustworthy, and dangerous or wasteful where they are not.

3.5 Adaptive, Hybrid, and Asynchronous Loops

Because no single shape dominates across tasks, the frontier is agents that choose the shape at runtime. AdaPlanner closes the loop on planning itself, refining a code-style plan from environmental feedback through in-plan and out-of-plan revision and thereby bridging static plan-then-execute and reactive replanning (Sun et al. 2023). More recent work formalizes the choice as a test-time compute-allocation problem, training an agent to decide *when* to invoke explicit planning versus act directly rather than planning at every step (Paglieri et al. 2025), and self-prompting state-tracking augmentations recover much of the benefit of heavier loops without their cost (Rozanov et al. 2024). A separate axis breaks the strict turn-based assumption entirely: asynchronous tool use replaces the lockstep loop with an event-driven execution model that overlaps thinking with tool and user I/O for real-time responsiveness (Ginart et al. 2024). Adaptivity is the natural endpoint of this section’s logic: if each shape is a bet about the environment, an agent should place the bet the environment warrants, but it reintroduces the problems the fixed loops avoided. The meta-controller that decides when to plan, reflect, or search needs an uncertainty or effort signal that is itself unreliable, so a mis-triggered plan wastes compute

and a mis-triggered reaction causes failure; and asynchronous execution introduces correctness hazards, because a side-effecting tool can be launched on arguments that later reasoning would have revised. Learning the loop is the right direction, but it moves the difficulty from choosing a loop to trusting the signal that chooses it.

Recent developments (2026). The elaboration of loop shape has continued at pace. On the decoupled side, position work argues web agents should abandon the ReAct per-step loop and execute untrusted page content only through quarantined subroutines that cannot alter the plan (Piet et al. 2026); map-then-act builds a cognitive map of the environment before committing to actions on long-horizon tasks (Liu et al. 2026c); and scheduler-theoretic reframings recast the loop as an explicit immutable DAG rather than an implicit context dependency (Wei et al. 2026a). A streaming and asynchronous line formalizes execution in which user intent evolves mid-run (Zhai et al. 2026) and speculatively pre-executes tool calls for real-time latency gains (Hooper et al. 2026), extending the async paradigm of §3.5. Reflective self-improvement now distills transferable heuristics across episodes (Allard et al. 2026), adaptive orchestration selects a control topology per task on the argument that structure has begun to dominate raw model capability (Yu et al. 2026a), and self-evolving world models tighten the plan-simulate loop by refining internal predictions at test time while holding the executor fixed (Zhang et al. 2026f). Systematic head-to-head comparisons of loop shapes have also begun to appear, pitting reflective and memory-augmented loops against fixed workflows in production extraction agents (Zhang et al. 2026i) and contrasting generator-evaluator, ReAct, and adversarial-evaluation paradigms within one deployed framework (Wang et al. 2026j).

KEY TAKEAWAY

The shape of the loop is a bet about the environment. Decoupled, reflective, and search-based topologies each relax the reason-before-every-action default, and each pays off only when its enabling assumption holds: a predictable environment for planning, reliable external feedback for reflection, cheap and reversible actions plus a trustworthy value signal for search. Where the assumption fails, the added structure costs tokens, latency, and new failure modes, which is why the frontier is moving from a fixed loop toward agents that learn when to plan, reflect, search, or simply react.

4 Loop Mechanics

Orthogonal to the *shape* of a loop (§3) and to *where* it lives (§5) are the control mechanisms that govern any loop, whatever its shape or locus: when it stops, whether its outputs are checked, how it manages its own growing context, and how it recovers from error. This section’s thesis is that an agent’s reliability is set by this control layer more than by its raw policy, and that a single tension runs through every mechanism. The stop, verify, and recover signals a loop depends on can be *self-generated* (cheap, scalable, and internal to the model) or *externally grounded* in execution, tools, curated memory, or humans. The recurring finding is that self-generated guards are unreliable and gameable, so durable loops externalize their guards and pay the resulting cost, latency, and residual exploitability. Table 3 lays out the four mechanisms against this single tension; each subsection is one row.

4.1 Termination and Step Budgets

A loop must decide when to stop, and the naive answers (a fixed iteration count or a static token cap) systematically misallocate compute, overthinking easy instances while starving hard ones. The overthinking pathology is concrete: reasoning models spend order-of-magnitude larger token budgets on trivial inputs than the task warrants (§5.5). The mechanisms that replace a hard-coded count are

TABLE 3: The four loop-control mechanisms, each caught in one tension: the cheap, scalable, self-generated guard versus the costly, externally-grounded one. The recurring finding of this section is that the self-generated form is unreliable or gameable, so durable loops pay for the grounded form—and even then inherit a residual failure they cannot fully escape.

Mechanism	Self-generated	External guard	Residual failure
Termination	Fixed step / token cap	Budget forcing, certainty early-exit, compute-optimal allocation	No right constant: overthinks easy, starves hard inputs
Verification	Intrinsic self-critique	PRM, execution / unit-test gates, generative verifier, LLM judge	Self-critique degrades answers; every verifier is a gameable proxy
Context mgmt.	Keep everything in context	Paging / external memory, compression, compaction	Context rot vs. lossy eviction—both directions lose information
Recovery	Restart from scratch	Checkpoint / restore, human approval gate	True rollback largely absent; gates invite automation bias

budget- and convergence-aware. Budget forcing terminates or extends the thinking phase by manipulating the decoding directly, giving a simple test-time knob on loop length (Muennighoff et al. 2025); certainty-based methods measure the model’s reasoning confidence to drive anytime early exit, cutting off unpromising trajectories before the budget is spent (Fu et al. 2024); and compute-optimal scaling allocates inference compute per-input by estimated difficulty, showing that *how* a budget is spent matters more than its size (Snell et al. 2024). Production systems now expose this as a first-class parameter, letting the caller trade tokens for accuracy at run time. The honest point is that no single budget is right across a task distribution: more iterations is not monotonically better, self-refinement gains plateau or reverse after a few rounds, and any fixed cap necessarily overthinks easy inputs while underthinking hard ones, so adaptivity, not a better constant, is the only principled answer.

A controlled budget sweep (E3). To turn that claim into a measurement, we swept the two budget knobs the harness exposes on the same 50-task ALFWorld suite as E2, on the 7B, 14B, and 32B models: the ReAct step cap $K \in \{5, 10, 20, 30, 50\}$ and the Reflexion retry budget $R \in \{0, 1, 2, 3\}$ (Figure 5). The two knobs behave differently, and neither rewards a larger constant. Raising K keeps buying success (7B climbs $8 \rightarrow 56\%$, 14B $10 \rightarrow 78\%$, 32B $14 \rightarrow 80\%$ from $K=5$ to 50), but concavely: the marginal step falls from roughly two points of success per step at low K to under one by $K=50$, while tokens grow linearly, so there is no budget at which the last step is worth its tokens uniformly. And the two larger scales nearly coincide past $K=20$, the accuracy saturation of E2 reappearing on the step axis: 32B does not separate from 14B by spending more steps. Raising R instead *plateaus*: 7B is flat at 68% for both $R=2$ and $R=3$ while its token cost rises from 1420 to 1829 per task, and 14B and 32B inch only $86 \rightarrow 88\%$ and $80 \rightarrow 84\%$ respectively, so past $R \approx 2$ the loop pays a linear token surcharge for near-zero accuracy, exactly the self-refinement plateau of §4.2 isolated as a budget effect (and the same pathology E2 saw at 1.5B, here at healthy scales). The caveats are the section’s: one benchmark and family, 50 tasks (so the $\approx \pm 13$ -point interval makes these plateaus regions, not thresholds), and a step sweep bounded above by ALFWorld’s natural 15–25 step horizon, beyond which little can change by construction. The reading is not a recommended constant but its refutation: the right budget is input-dependent, which is the case for the adaptive termination this subsection argues for.

4.2 Self-Verification versus External Checks

The cheapest imaginable guard is the model checking itself, and it is also the least reliable. Intrinsic self-correction (a model critiquing and revising its own output with no external signal) frequently *degrades* correct answers rather than fixing wrong ones (Huang et al. 2023), and a critical survey of the




FIGURE 5: E3: a controlled budget sweep on the E2 suite (7B, 14B, 32B). **(a)** Raising the ReAct step cap K keeps buying success but concavely, so the marginal step pays less as K grows while tokens grow linearly; the two larger scales saturate together past $K=20$. **(b)** Raising the Reflexion retry budget R plateaus, with 7B flat at 68% across $R=2$ and $R=3$ despite a rising token bill, so past $R \approx 2$ extra reflection burns tokens on failures the model cannot fix. **(c)** The same points in cost space (success vs. mean tokens per task, log): the K path (solid) climbs up-and-right, while the R path (dashed) runs flat-right, spending tokens for little success. Neither knob rewards a larger constant, which is the case for adaptive termination.

self-correction literature finds no convincing success outside tasks with reliable external feedback or an oracle stopping criterion (Kamoi et al. 2024). The mechanism behind the failure is that LLMs are not reliably better at discriminating among their own candidate outputs than at producing them, so the “evaluate” half of self-evaluate-then-revise is itself unsound (Jiang et al. 2024). This recasts the popular self-correction loops of §3.3: the gains attributed to Self-Refine and Reflexion (Madaan et al. 2023; Shinn et al. 2023) come almost entirely from externally grounded feedback (unit-test pass/fail, task reward, oracle labels), and CRITIC states the constructive version of the same claim, that tool-interactive critique is the necessary ingredient (Gou et al. 2023). The field’s headline self-correction result and its strongest negative result must be read together: verification helps when grounded, and misleads when not.

4.3 Verifiers in the Loop

If self-verification is unsound, the loop must import an external verifier, and the mechanisms that do so are the workhorses of reliable agents, each a noisy and gameable proxy rather than a ground-truth oracle. Process reward models score intermediate steps rather than only final answers (Lightman et al. 2023), generative verifiers cast verification as chain-of-thought next-token prediction and outperform discriminative reward models and LLM-as-judge (Zhang et al. 2024b), unit-test gating refines generated code against generated tests in a flow-engineered loop (Ridnik et al. 2024), and execution feedback teaches a model to debug from the results of running its own code (Chen et al. 2023b). For open-ended outputs the loop often falls back on an LLM judge, whose limitations are well documented: judges exhibit self-preference, position, and length biases (Zheng et al. 2023a), and agent-specific judging and risk-evaluation frameworks inherit the same unreliability (Zhuge et al. 2024; Yuan et al. 2024). Every in-loop verifier therefore has a documented exploit: PRM signal can saturate and be gamed, judges can be flattered, and unit-test gating is actively reward-hacked by agents that hard-code expected outputs or tamper with the test harness. A passing check is evidence, not proof, which is why verification must co-evolve with the generator it polices rather than being fixed once.

4.4 In-Loop Context Management

A long-horizon loop accumulates observations, and its context is a finite, degrading resource that must be actively curated. The difficulty is a genuine no-free-lunch. Simply keeping everything in context does not work: models use long inputs unevenly, with material in the middle of a long window effectively lost (Liu et al. 2023), so accuracy degrades from context rot even when nothing is truncated. But the remedies are lossy in the other direction. Paging and external memory (treating the context window like an operating system’s virtual memory and swapping state in and out) extend effective horizon at the cost of managing what to evict (Packer et al. 2023); prompt compression shrinks the context under a budget controller but can drop the one token the task hinges on (Jiang et al. 2023); and the industry discipline of context engineering (compaction, structured note-taking, and sub-agent context isolation) codifies these moves while acknowledging that summarizing or evicting turns can silently discard load-bearing state (Anthropic 2025e; Anthropic 2025g). Both directions fail: naive accumulation triggers context rot and cost blow-up, and every compaction step risks erasing exactly the fact the loop later needs. Context management is thus empirical craft, not theory, and remains one of the least principled parts of the loop.

4.5 Error Recovery, Rollback, and Human Gates

Because a single early error can cascade unrecoverably through the rest of a trajectory, robust loops need explicit recovery machinery, and here the honest finding is how little of it exists in practice. True rollback (restoring the agent and its environment to a known-good state and re-exploring) is largely absent: mainstream coding agents do at best file-level version stashing and discard process state, so “recovery” usually means restarting from scratch, and agents observably repeat the same error for many turns despite acknowledging it (Yang et al. 2024a; Jimenez et al. 2023). Semantics-aware checkpoint/restore runtimes for agent sandboxes are an emerging attempt to make real rollback and branching possible (Wu et al. 2026a), but they are not yet standard. The complementary mechanism is a human gate: a required approval before a high-risk action, as deployed in production software agents that pause for plan- and code-approval before opening a pull request (Takerngsaksiri et al. 2024), and increasingly mandated by regulation for high-risk automation. Human gates are effective but shift the failure mode rather than removing it: they add latency, and they invite automation bias, where an over-trusting reviewer rubber-stamps approvals and quietly nullifies the oversight the gate was meant to provide.

Recent developments (2026): context management. In-loop context management (§4.4) became the single most active mechanics topic of 2026. The dominant finding is that active curation beats naive accumulation: enterprise studies report that selectively retaining recent tool interactions plus summarization can halve token cost while raising completion (Lodha et al. 2026), and, strikingly, simple observation masking matches far more complex LLM summarization at half the cost (Lindenaubauer et al. 2025). A wave of systems makes compaction safer or learnable: validating each summary against ongoing reasoning (Chen et al. 2026e), letting the agent decide when to compact via a tool (Li et al. 2026g), training the compaction policy itself with RL (Li et al. 2026b; Yi et al. 2026), caching a reusable orientation map of recurring environments (Gu et al. 2026b), and preserving prompt-cache continuity under eviction (Xu et al. 2026c). Others page full interaction history into external indexed memory while keeping a compact working context (Wang et al. 2026e; Liu et al. 2026d), and a data-management-lens evaluation of twelve agent-memory systems finds no architecture dominates (Zhou et al. 2026b), so the no-free-lunch of §4.4 is now empirically mapped rather than resolved. Learned and mechanism-level memory continued to proliferate: adaptive compression policies that jointly tune when and how much to compress (Gao et al. 2026), action-preserving observation compression for coding agents (Chen et al. 2026f), graph-structured experience memory for one-shot error correction (Wang et al. 2026i), implicit procedural memory injected by activation steering rather than retrieval (Zhao et al. 2026f), git history repurposed as persistent developer-agent memory (Guo 2026), and scoped verification to keep a harness-updated context reliable under distribution shift (Hsu and Lu 2026). On the recovery side, trajectory-level anatomies now treat failure as a temporal process rather than a terminal

outcome (Zhao et al. 2026g), and behavioral set-shifting tests probe whether an agent re-selects its tools when a once-reliable tool silently degrades (Ye 2026).

Recent developments (2026): budgets and verifiers. Termination and budgeting (§4.1) gained finer control: per-step adaptive reasoning-effort routers (Yang et al. 2026a; Jung et al. 2026), training-free draft-agreement budget routing for hybrid reasoning models (Lee et al. 2026a), wall-clock-aware test-time scaling (Ma et al. 2026b), value-of-information budget allocation across tool calls and tokens (Fang et al. 2026), budget-tracking tool agents (Liu et al. 2025), semantic-convergence early exit (Min et al. 2026), and risk-controlled stopping for parallel reasoning traces (Chen et al. 2026c). Verification (§4.3) is increasingly cast as its own agentic, test-time-scalable process (bidirectional multi-turn agentic verifiers (Zhang et al. 2026a), rubric-guided self-verification that exploits verification asymmetry (Wan et al. 2026), and fine-grained sub-question verification (Zhao et al. 2026c)), while confidence-bounded act-or-defer rules decide when a loop should hand off to a human (Wang et al. 2026c).

KEY TAKEAWAY

An agent loop’s reliability is set by its control scaffolding, not its raw policy, and one tension runs through every mechanism: self-generated stop, verify, and recover signals are cheap but unreliable and gameable, so durable loops externalize their guards into adaptive budgets, execution and PRM checks, curated and paged memory, checkpoints, and human approval, while paying the cost, latency, and residual exploitability that grounding inevitably introduces. Termination has no right constant, self-verification needs external anchoring, every verifier is a gameable proxy, context management is lossy in both directions, and true rollback is mostly absent, so mechanics, not the policy, is where reliability is won or lost.

5 Trained Loops

The paradigms of §3 treat the loop as something the harness imposes on a frozen model. The alternative is to push the loop into the weights: train the model, end to end, to run the reason–act–observe cycle itself (to decide when to search, when to call a tool, when to reflect, and when to stop) so that control becomes native forward-pass behavior. This is the “internalize” direction of §1, and it rests on reasoning-native base models whose deliberate multi-step reasoning was itself induced by outcome-only reinforcement learning (DeepSeek-AI et al. 2025), now surveyed as a field in its own right (Zhang et al. 2025b). The section’s organizing question is whether this training genuinely teaches *new* loop-control ability or mainly sharpens and compresses behaviors already latent in the base model, and the honest answer, developed across the subsections, is closer to the latter, with a matching failure mode introduced for each gain.

5.1 From Prompted to Trained Loops

The foundational move is outcome-reward RL over multi-turn rollouts: let the agent interact freely, mask the retrieved or tool-returned tokens from the loss so the model is not trained to imitate them, and reward only final task success. Search-R1 uses this recipe to teach a model when to issue search queries and how to reason over the results, with no hand-designed retrieval scaffold (Jin et al. 2025); ReTool does the same for a code interpreter, learning strategic invocation of execution (Feng et al. 2025a); ToRL scales tool-integrated RL from base models and observes emergent, self-regulated tool use (Li et al. 2025c); DeepResearcher trains the full deep-research loop against live web search rather than a sandboxed corpus, eliciting emergent plan–cross-validate–reflect behavior (Zheng et al. 2025b); and SWE-RL brings RL at scale to real software-evolution data using a rule-based similarity reward (Wei et al. 2025). These systems reliably remove the hand-built scaffold and improve sampling efficiency. But

outcome-only rewards carry two structural liabilities that the both-sides tradition requires we state plainly. First, a reward computed only at the end gives no per-step signal and is a *proxy* open to gaming: SWE-RL rewards patch *similarity*, not verified correctness, so the model can be optimized toward looking right rather than being right. Second, and more fundamentally, a prominent analysis finds that RL with verifiable rewards raises pass@1 sampling efficiency yet does not expand the reasoning boundary: base models overtake their RL-trained descendants at large pass@ k , implying the training concentrates probability mass on solutions the base model could already find rather than teaching genuinely new ones (Yue et al. 2025). Trained loops sample known good behavior more reliably; whether they cross the base model’s competence frontier is, on current evidence, doubtful.

5.2 Credit Assignment and Stability

The central technical obstacle to trained loops is assigning credit across long, sparse-reward trajectories without the policy collapsing. RAGEN diagnoses the characteristic failure (an “Echo Trap” in which multi-turn RL narrows into repetitive, self-reinforcing reasoning) and proposes StarPO, a trajectory-level objective, with a stabilized StarPO-S variant that uses uncertainty filtering, KL removal, and decoupled clipping (Wang et al. 2025e). GiGPO introduces a critic-free, two-level (episode plus anchor-state step) advantage estimator that improves ALFWorld and WebShop success over GRPO at equal memory (Feng et al. 2025b), and ARPO adds entropy-triggered branch sampling at the high-uncertainty rounds that follow a tool call, aligning step-level tool use at a fraction of the trajectory-level budget (Dong et al. 2025). The honest caveat is that these are mitigations, not cures: StarPO-S *delays* rather than eliminates collapse, and the stability techniques are algorithm-specific rather than general. The reflective loop of §3.3 can likewise be made trainable rather than prompted (Retroformer trains a retrospective model by policy gradient to generate better reinforcement cues (Yao et al. 2023b), Reflect-Retry-Reward rewards the self-reflection tokens when a subsequent retry succeeds (Bensal et al. 2025), and SAMULE fine-tunes a dedicated retrospective model on multi-level reflections (Ge et al. 2025)), but each inherits the same credit-assignment fragility, now applied to the reflection tokens themselves.

5.3 The Real Bottleneck Is the Environment

A recurring finding is that trained loops are gated less by the RL algorithm than by the availability of executable, verifiable environments that emit reward at scale. SWE-Gym provides the first executable training environment of real software-engineering tasks together with learned verifiers, and much of the measured progress comes from the environment and verifier rather than a new optimizer (Pan et al. 2024). This reframes the field’s bottleneck honestly: executable environments are scarce and small (SWE-Gym comprises only a few thousand real task instances); rule-based rewards such as unit-test pass or patch similarity are proxies open to reward hacking, and learned verifiers can themselves be gamed, so a result is only as trustworthy as its verifier and rarely disentangles “better loop” from “better exploitation of the reward.” Self-improvement schemes that let an agent generate and train on its own trajectories (ReST-style growing-batch RL over ReAct traces (Aksitov et al. 2023) and autonomous reason-act trajectory annotation (Yang et al. 2024b)) push against data scarcity but compound the same proxy-reward risk, since the agent both writes and grades its own training signal.

5.4 Loop Internalization

Rather than training a loop from scratch, one can distill the trajectories of an external multi-agent or tool-using scaffold into a single model, so the entire orchestrated loop becomes native behavior, and then refine it with agentic RL. Chain-of-Agents is the flagship: it collapses a multi-agent system into single-model “chain-of-agents” trajectories and applies agentic RL on top, producing an agent foundation model that runs the loop internally across web and code settings (Li et al. 2025b). Internalization is attractive (one model, no orchestration overhead) but the trade it makes is efficiency for transparency, and it should be weighed against the pass@ k boundary result of §5.1. Distillation is upper-bounded by the teacher scaffold and copies its inefficiencies, and collapsing an explicit multi-agent loop into

one forward pass sacrifices the modularity, inspectability, and per-step controllability that made the external loop debuggable in the first place. What is gained in latency is paid in the ability to see, and intervene in, what the loop is doing.

5.5 Controlling the Loop’s Length

A trained loop must also learn *how much* to think, and over- and under-thinking are opposite failures of self-regulated loop length. Overthinking is now well characterized: reasoning models spend enormous token budgets on trivial inputs (Chen et al. 2024a), motivating explicit length control. L1’s length-controlled policy optimization trains reasoning length as a promptable constraint, smoothly trading compute for accuracy (Aggarwal and Welleck 2025), and agentic self-awareness trains the model to invoke reflection or knowledge only when a situation-judgement criterion fires, cutting reasoning cost without loss (Qiao et al. 2025); test-time-scaling studies map the broader space of when extra inference compute helps a loop at all (Zhu et al. 2025a). The two failure modes are in direct tension, and this is where the naive “think less” remedy backfires. Underthinking is not too many tokens but too much *switching* (prematurely abandoning promising lines of thought (Wang et al. 2025c)) and, more sharply for agents, pushing for brevity can trigger premature disengagement and analysis-paralysis in exactly the agentic settings where a step of environment interaction should replace a step of internal reasoning (Cuadron et al. 2025). Moreover, L1’s clean length–accuracy trade is validated mainly on single-turn mathematics, so whether learned length control transfers to multi-turn agentic loops remains open.

Recent developments (2026). Agentic RL in 2026 concentrates on the credit-assignment and stability problems of §5.2: supervision-free entropy modulation across turns (Zhao et al. 2026a), uncertainty-guided exploration control (Wang et al. 2026f), graph-based step-level credit from a state-transition graph (Cheng et al. 2026), and analyses of recurring entropy-eruption cycles (Li et al. 2026c), alongside asynchronous post-training whose stability is governed by staleness scaling laws (Song et al. 2026a; Liu et al. 2026b) and disaggregated training systems that map rollout, environment, and reward to best-fit hardware (Gao et al. 2025); a survey now catalogs forty-seven credit-assignment methods separating token-level from multi-turn agentic RL (Zhang et al. 2026c). On the application side, deep-research and GUI agents are trained end to end with self-reflective or partially-verifiable rewards (Yu et al. 2026b; Li et al. 2026d; Yang et al. 2026c; Wan et al. 2025), tool-use decisions are aligned to uncertainty to curb unsupported calls (Zhou et al. 2026c; Chen et al. 2026a), the loop is co-evolved with a policy that also generates its own skills (Zhang et al. 2026b), and multi-turn multi-task frameworks scale the training infrastructure itself (Zhang et al. 2025a; Zhang et al. 2026e). The stability and credit-assignment thread of §5.2 deepened further, with diagnoses of outright collapse in multi-step tool-use RL and supervisory-signal fixes (Hao et al. 2026), turn-level credit estimation for sparse long-horizon rewards (Tao et al. 2026), and sandbox-native branching rollouts (He et al. 2026b); the environment-and-infrastructure bottleneck of §5.3 drew scaling frameworks across massive tool environments (Zhou et al. 2026f), single-rollout asynchronous optimization (Hou et al. 2026b), and counterfactual context anchoring so a post-trained agent absorbs new tools without retraining (Liu et al. 2026g); and the internalization move of §5.4 was pushed to its limit by compiling an entire agentic workflow directly into the weights at near-frontier quality and a fraction of the orchestration cost (Dennis et al. 2026).

KEY TAKEAWAY

Training the control loop into the weights (through outcome-reward multi-turn RL, environment and verifier engineering, multi-agent distillation, and explicit length control) reliably removes hand-built scaffolds and improves sampling efficiency. But the evidence suggests it mostly sharpens and compresses behaviors already latent in the base model (base models overtake at large pass@*k*) while introducing a symmetric failure mode for each gain: reward hacking for outcome




FIGURE 6: The skill lifecycle, loop-centric and adversarial. Procedural competence extracted from loop trajectories (left) passes through seven stages before it re-enters the loop as in-context procedure; representative systems sit above each stage and that stage’s documented failure mode sits below it (Wang et al. 2023a; Zhao et al. 2023; Wang et al. 2024c; Ouyang et al. 2025; Anthropic 2025f; Qin et al. 2023; Patil et al. 2024; Zhang et al. 2026d; Huang et al. 2026a). The dashed lane marks where the supply chain attacks the lifecycle: poisoned artifacts published into the store, memory poisoning at retrieval, and trusted-channel metadata injection at composition (Wang et al. 2025a; Chen et al. 2024e; Liu et al. 2026a); governance is the young counterweight (Liu et al. 2026e; Bhatt et al. 2025), and its return arc (retire, update, cap) is what closes the lifecycle into a loop. The field’s hard problem has moved from left to right, from acquiring skills to governing them.

rewards, the Echo Trap for multi-turn credit assignment, verifier-bound trust for engineered environments, lost inspectability for internalization, and premature disengagement for aggressive length control. The honest verdict is a powerful compressor of known skills, not yet a generator of new loop-control ability.

6 Skills

The externalize direction of §1 pushes competence out of the weights and into portable, reusable units of procedure: *skills*. We give skills a section of their own, rather than folding them into “tool use” or “memory” as prior agent surveys do, because a skill is where the loop’s accumulated know-how is stored and from which it is re-supplied, and because the field’s hard problem has visibly moved from *acquiring* skills to *governing* them. A concurrent survey catalogs agent skills by representation, acquisition, retrieval, and evolution (Zhou et al. 2026a); our treatment differs in being loop-centric and adversarial: we organize skills by where they enter the loop and we give the negative results (interference, non-transfer, poisoning) equal weight with the acquire-and-reuse promise.

6.1 What a Skill Is, and Is Not

We fixed the working definition in §2.3: a skill is named, reusable, retrievable, composable procedural knowledge. The definitional difficulty, which we present honestly rather than resolve by fiat, is that the community uses the word for at least three incompatible representations. In Voyager a skill is executable code, a Python function written on success and added to a growing library (Wang et al. 2023a), an identification that CodeAct sharpens by making executable code the agent’s native action space and thereby blurring the line between calling a tool and invoking a stored procedure (Wang

et al. 2024b). In Anthropic’s Agent Skills a skill is a folder of Markdown instructions, scripts, and resources loaded through progressive disclosure (Anthropic 2025f). In an experiential learner it is a natural-language insight distilled by comparing successful and failed trajectories (Zhao et al. 2023), and in reasoning-memory systems a strategy-level item carrying negative constraints (Ouyang et al. 2025). These are code, document, and memory respectively, and no reconciliation makes them one data type. What unifies them is functional: each is named, stored across episodes, retrieved on demand, and composed with others, which is exactly what separates a skill from a stateless *tool* the model invokes but does not own (Schick et al. 2023; Anthropic 2024c) and from an ephemeral *prompt*. The distinction that survives is behavioral, not representational, and we build the rest of the section on it.

6.2 Acquisition

Agents acquire skills by distilling experience into reusable artifacts, almost always in-context and training-free, sidestepping the fine-tuning cost and generalization loss of writing procedure into weights. The artifacts differ in granularity. Voyager grows a library of executable functions through iterative self-verification (Wang et al. 2023a); ExpeL abstracts cross-task natural-language insights from comparing trajectories (Zhao et al. 2023); Agent Workflow Memory induces reusable sub-routines from past trajectories, offline and on the fly (Wang et al. 2024c); ReasoningBank distills strategy-level items (Ouyang et al. 2025); and a wave of recent systems extends acquisition across modalities and domains: programmatic skills induced during web navigation (Wang et al. 2025f), self-discovered and honed web-agent skills that transfer to weaker agents (Zheng et al. 2025a), interaction trajectories abstracted into reusable instructions by backward construction (Su et al. 2025), cognitive abstractions from noisy multimodal trajectories (Sarch et al. 2024), human-readable instruction manuals distilled from environmental interaction (Chen et al. 2024b), explicit skill-curation modules in general-computer-control agents (Tan et al. 2024; Agashe et al. 2024), and lifelong libraries for embodied and open-world control (Tzifas et al. 2024; Wang et al. 2023d; Zhu et al. 2023). The volume of positive results is real, and so are two documented limits. First, success-only memory is a failure mode, not a simplification: ReasoningBank’s central finding is that encoding *failures* and negative constraints is what drives the gains, and libraries that store only successful subgoal sequences cause incorrect paths to be reused and suppress exploration (Ouyang et al. 2025). Second, acquisition gains are benchmark-narrow (Voyager transfers within Minecraft, workflow memory within one web suite), so “lifelong” and “open-ended” skill-acquisition claims run ahead of the cross-domain evidence that would justify them.

6.3 Representation, Storage, and the Retrieve–Select–Compose Pipeline

Once acquired, a skill must be represented, stored, and, at inference, retrieved, selected, and composed, and this downstream pipeline, not acquisition, is where most systems actually break. Representations range from vector databases of trajectories retrieved top- k as few-shot exemplars (Zhao et al. 2023), to workflows with variable slots for generalization (Wang et al. 2024c), to on-disk file trees loaded by progressive disclosure (Anthropic 2025f), backed by neural retrievers over massive tool and skill sets (Qin et al. 2023). This is the section’s strongest negative-results cluster. Flat libraries do not scale, and adding irrelevant skills degrades performance: an agent given many tools in-prompt selects worse than one given a retrieved few (Patil et al. 2024). Retrieval is brittle to hard-negative distractors, so the bottleneck shifts from *recall* to *incorporation*: the agent invokes the wrong skill even when the correct one was retrieved, because their embeddings are semantically indistinguishable, and injecting several full skills at once causes prompt overload and procedural-confusion interference. The practical consequence is that recall-oriented benchmarks systematically overstate real end-task utility: a system that retrieves the right skill can still fail to use it, and more skills in the library can make the loop worse rather than better. We measure this interference directly in a controlled ablation (E4, §8.5, Table 7): holding the base model fixed, a full in-context library lowers 14B ALFWorld success below the matched-skill arm, so the operative bottleneck is incorporation, not recall.

At the representation end, the industry convention that has spread fastest is the file-on-disk skill: a `SKILL.md` whose YAML front-matter (name, a one-line description, and a trigger condition) is what a retriever reads to decide relevance, while the procedural body is loaded into context only once the skill is selected, the progressive-disclosure discipline of (Anthropic 2025f). The example below is one of the skills our own E4 harness distills from successful ALFWorld trajectories (§8.5), written in that format.

THE ANATOMY OF A `SKILL.md` (FRONT-MATTER FOR RETRIEVAL, BODY FOR USE)

```
---
name: clean-and-place
description: Clean an item at a sink, then place it at a target.
when_to_use: task asks to put a *cleaned* object onto a receptacle
---
1. Locate the item on a surface (countertop, cabinet).
2. Carry it to a sinkbasin, turn on the tap, and clean it.
3. Navigate to the target receptacle; place the item.
# pitfall: forgetting to take the item after cleaning it.
```

The split is the point: only the front-matter competes for retrieval attention, so a library can grow without every skill’s full text crowding the prompt, and the negative results above (interference, wrong-skill selection) are precisely failures of the front-matter to disambiguate one body from another.

6.4 Ecosystem and Standards

Skills and tools are consolidating around two open standards: the Model Context Protocol for tool and data connectivity (Anthropic 2024c), and Anthropic’s Agent Skills / `SKILL.md` for portable procedural packages (Anthropic 2025f), with partner directories and marketplaces turning skills into distributable, versioned artifacts adopted rapidly across vendors. Convergence is genuine, but the “standard” is thin, and the honest reading tempers the optimism twice. First, `SKILL.md` is essentially Markdown plus YAML front-matter with no formal semantics, verification, or capability contract, so portability and composability claims outrun the evidence; the marketplace is already large and uneven, with measurement studies finding that a majority of listed servers are low-value and that the ecosystem carries dependency monocultures (Guo et al. 2025). Second, the ecosystem has not reconciled two overlapping abstractions: MCP tools (external, protocol-mediated) versus skills-as-context (instructions loaded into the model), and a systematization of the MCP ecosystem shows the same primitives (resources, prompts, tools) are treated as reusable loop units without an agreed trust or semantic model (Gaire et al. 2025). Standardization is happening at the seams faster than the semantics and security to support it.

6.5 Supply-Chain Security

Making skills and tools distributable converts them into a supply-chain attack surface, and this is the counterweight to ecosystem optimism. The threat is structural, not a patchable bug: a discovery protocol that auto-trusts third-party metadata invites tool poisoning, where malicious instructions hidden in a tool description reach the model as if system-authored. On real MCP servers this is not hypothetical: a benchmark over hundreds of live tools finds poisoned metadata succeeding in the majority of cases, and, damningly, more capable instruction-following models are *more* susceptible while safety tuning refuses only a tiny fraction (Wang et al. 2025a). Large-scale analyses document coordinated multi-tool exfiltration with no direct victim interaction (Zhao et al. 2025), pervasive vulnerability across tens of thousands of published skills (Liu et al. 2026a), threat surfaces across the full skill lifecycle (Li et al. 2026i; Hou et al. 2025), and marketplace-scale precedents from custom-GPT app stores (Ogundoyin et al. 2025; Hasan et al. 2025). Worse, the very self-acquired memory praised in §6.2 is itself a high-value target:

poisoning fewer than a fraction of a percent of an agent’s retrieval memory can drive attack success above eighty percent with negligible impact on benign behavior (Chen et al. 2024e). Defenses exist (cryptographically signed, OAuth-enhanced tool definitions with policy-based access control (Bhatt et al. 2025)) but they lag adoption and cannot cover “rug-pull” updates where a tool benign at install turns malicious later. Capability sharing and attack surface grow together, and the skill that makes an agent more capable is simultaneously the vector that can compromise it.

Recent developments (2026). Skills were the fastest-growing topic of 2026, dominated by *self-evolving skill libraries*: systems that create, validate, store, retrieve, and retire skills autonomously (Lin et al. 2026a; Chen et al. 2026d; Yang et al. 2026d; Zhang et al. 2026g; Shen et al. 2026), RL-trained skill curators over an external repository (Ouyang et al. 2026), generative and group-structured skill retrieval that replaces flat skill lists (Zhao et al. 2026d; Zeng et al. 2026; Meng et al. 2026; Ding et al. 2026b), selective experience utilization invoked only when needed (Zhao et al. 2026e), RL that internalizes discovered skills into the weights (He et al. 2026a), meta-skills for multi-agent orchestration (Lin et al. 2026b), and even federated skill evolution that shares semantic skill diffs rather than raw trajectories (Yang et al. 2026b). Crucially for the both-sides account, this wave produced its own named failure mode: *library drift*, in which unbounded accumulation silently degrades retrieval, remedied by outcome-driven retirement and capacity caps (Zhang et al. 2026d), the acquisition-versus-governance tension of §6.2 made concrete, while lifecycle-governance frameworks add collection, recommendation, attribution, and harmful-update blocking (Liu et al. 2026e), and a systematic study finds model-generated skills carry a negative-transfer risk driven by extractor–consumer compatibility rather than scale (Huang et al. 2026a). Two 2026 surveys now consolidate the area (Xu et al. 2026a; Ding et al. 2026a), and a broader review frames skills as one facet of a general move to *externalize* agent capability into memory, protocols, and harness (Zhou et al. 2026d). Empirical studies now measure the deployed ecosystem itself: what agents actually call across 177,000 MCP tools (Stein et al. 2026), the quality of tool descriptions across live servers (Hasan et al. 2026), and retrieval that scales orchestration across thousands of tools and sub-agents (Lumer et al. 2025). *Skill evaluation*, long this section’s named gap, began to catch up: benchmarks isolating whether agents can turn their own run evidence into reusable skills (Peng et al. 2026), scalable frameworks for scoring skills across models and domains (Shaposhnikov et al. 2026), and consolidations that measure the fragmented public SKILL.md ecosystem’s real reuse value and redundancy (Wang et al. 2026l). Alongside these, self-evolving libraries advanced with two-timescale recursive meta-skill evolution (Wang et al. 2026m), trajectory-level self-adapting skills with fine-grained failure attribution (Yu et al. 2026d), and generation-evaluation-improvement loops that decouple procedural control into reusable skill objects (Zhang et al. 2026h).

KEY TAKEAWAY

Skills are a real and useful compositional abstraction whose unity is functional (named, reusable, retrievable, composable), not representational: the same word spans executable code, Markdown packages, and natural-language insights, and we do not force them into one type. The field’s hard problem has moved from acquiring skills to governing them: success-only acquisition reuses mistakes, retrieval breaks at the incorporation step so more skills can hurt, the SKILL.md/MCP standards are thin on semantics and verification, and a shared skill is simultaneously a capability and a supply-chain attack vector against which content-level safety is nearly useless. The skill ecosystem’s benefits and its attack surface expand in lockstep.

7 Harnesses and Orchestration

We defined the harness in §2.2 as the code that realizes the loop; this section argues that by 2026 the harness, not the base model, has become the primary design surface for agent reliability. The

evidence is uncomfortable for model-centric narratives: holding the model fixed and changing only the surrounding code moves headline benchmark numbers by tens of points, so much of what is reported as model progress is harness progress (§8). We trace the harness from the low-level interface it presents to the model, through the frameworks that orchestrate multi-step control, to the production coding agents that now define the state of practice, and we take a side in the field’s central design argument: control belongs at the least-agentic point on the spectrum that still solves the task.

7.1 The Agent–Computer Interface

The lowest layer of the harness is the interface between the model and the computer: the commands it may issue, the form of the feedback it receives, and the guardrails on its edits. SWE-agent’s central finding is that designing this agent–computer interface (ACI) *for the model rather than the human* (a compact command set, linted and de-noised output, edits validated before they land) improves task success more than upgrading the base model does (Yang et al. 2024a). Open platforms generalize the ACI into a reusable substrate for building and evaluating agents (Wang et al. 2024a), and the Model Context Protocol standardizes the tool half of the interface so a capability can be exposed to any model through one connector (Anthropic 2024c). The honest caveat is that ACI gains are model-specific: an interface tuned for one model can degrade another, so the “interface matters more than the model” effect simultaneously undermines cross-model comparability: an SWE-bench score is interpretable only within a matched harness (§8).

7.2 Orchestration Frameworks and the Workflow–Agent Spectrum

Above the interface sits orchestration: how multiple model calls, tools, and sub-tasks are sequenced. Production frameworks span a continuum from fully predefined *workflows*, in which LLM calls are wired through fixed code paths, to autonomous *agents*, in which the model directs its own control flow; Anthropic’s engineering guidance catalogs the patterns in between (prompt chaining, routing, parallelization, orchestrator–workers, evaluator–optimizer) and advises choosing the least-agentic one that solves the task (Anthropic 2024a). Conversational multi-agent frameworks (Wu et al. 2023) and graph-structured orchestrators (LangGraph and kin) push toward more explicit control, while the cognitive-architecture view supplies the vocabulary for placing them on one map (Sumers et al. 2023). The negative result to record is that heavier orchestration frequently fails to pay for itself: added layers impose leaky abstractions and hidden control flow that practitioners repeatedly abandon in favor of owning their own loop, and controlled comparisons find extra orchestration or reasoning steps often yield minimal or negative benefit over a single well-built loop, the same lesson the evaluation literature reaches from the measurement side (§8).

7.3 Multi-Agent Parallelism versus Single-Threaded Context

The sharpest design disagreement in the field is whether to decompose the loop across multiple agents at all. One camp builds orchestrator–worker systems in which each sub-agent holds an isolated context window and explores in parallel, and reports that context isolation is precisely what makes the design work for broad, read-heavy tasks (Anthropic 2025g). The opposing camp argues against multi-agent architectures for most work, on the ground that reliable agents need one continuous, shared context with a single writer, and that splitting context across agent boundaries is where these systems break (Yan 2025). The evidence favors caution: a taxonomy of multi-agent failures documents pervasive coordination and specification breakdowns at failure rates far above what the added machinery justifies (Cemri et al. 2025), so parallel multi-agent designs are defensible for independent, parallelizable threads but not for shared-state action, where a single loop with disciplined context management is the more reliable choice.

7.4 Context Engineering: The Harness as Context Manager

As task horizons lengthen, the harness’s dominant job shifts from writing a prompt to curating the context window turn by turn: compaction, retrieval, structured note-taking, and offloading state to sub-agents (Anthropic 2025e). This is the mechanic of §4.4 seen as a harness responsibility rather than a loop-internal one: every frontier model degrades as its context grows, using the middle of a long window poorly (Liu et al. 2023), so a harness that does not actively manage context inherits context rot regardless of the nominal window size. The honest limit is that context engineering remains empirical craft: compaction discards information silently and can make an agent loop, forget prior decisions, or repeat completed actions, and there is no principled criterion for what to retain, so the harness trades one failure mode (context rot) for another (lossy compression) with no theory to guide the trade.

7.5 From Research Scaffolds to Production Coding Harnesses

The state of practice is now set by production coding agents (Claude Code, Codex CLI, Cursor, Devin, Aider) which converge on a small set of software-engineering principles: own your prompts and control flow, keep the loop a stateless reducer, treat tool calls as structured output, and make human-in-the-loop a first-class operation. What distinguishes this generation from a fixed workflow graph is where they place control: the harness takes a high-level goal and lets the model direct the loop, choosing at each turn which action, sub-agent, or skill to invoke and when to stop, the *agent* end of the workflow-versus-agent distinction of §2.2 (Anthropic 2024a). In practice this surfaces as two coupled behaviors: goal-directed persistence, where the harness re-enters the loop until the objective is met or a budget is spent rather than halting after one pass, and dynamic orchestration, where sub-agents and tool pipelines are constructed at runtime instead of wired ahead of time (Ruan et al. 2026), with dynamic skill routing supplying the reusable procedure at each step (Gu et al. 2026a). The “12-factor agents” manifesto codifies these into a reusable checklist (Horthy 2025), while a parallel line reframes the whole system as a *compound AI system* whose quality comes from composition rather than any single model (Zaharia et al. 2024), and skill-mediated reference architectures extend the same discipline to how a harness discovers and activates reusable skills (Xia et al. 2026). These harnesses set the top of leaderboards on SWE-bench Verified (Jimenez et al. 2023) and terminal-based coding benchmarks (Merrill et al. 2026). But much of this canon is non-peer-reviewed engineering writing, and the leading harnesses are closed-source, so their headline numbers cannot be independently reproduced and conflate the contribution of the scaffold with that of the model, making “the harness matters more than the model” at once the field’s central practical insight and its central reproducibility problem (§8).

Recent developments (2026). 2026 elevated the harness from a design surface to a *scaling axis*: work argues the next bottleneck is system rather than model scaling (naming context governance, trustworthy memory, and dynamic skill routing as the levers (Gu et al. 2026a)), shows that harness design and post-training interact, so that harness-aware post-training improves in- and out-of-distribution performance (Kim et al. 2026), and makes the harness itself learnable as a bidirectional controller that transfers across model sizes (Wang et al. 2026d). Source-code studies now taxonomize real coding agents (thirteen open-source systems across twelve dimensions (Rombaut et al. 2026) and a terminal coding-agent report that separates scaffolding from harness (Bui et al. 2026)) while orchestration frameworks add automatic sub-agent creation (Ruan et al. 2026) and modular tree-search harnesses isolate policy diversity as the true bottleneck (Li et al. 2026e). The single-versus-multi-agent debate of §7.3 sharpened: under equalized reasoning-token budgets, single agents match or beat multi-agent systems, implying much of the reported multi-agent advantage is unaccounted compute (Tran et al. 2026); where multi-agent structure does survive, model-adaptive isolate-or-score assessment cuts its cost (Lee et al. 2026b); and surveys now frame the field explicitly through the model-harness lens (Guo et al. 2026). The harness itself became an object of definition and measurement: proposals of necessary and sufficient conditions for what counts as a harness (Macedo 2026) and handbooks treating the evolving harness as a first-class readable, editable artifact (Wang et al. 2026h), controlled evidence that scaffolding evolution alone shapes coding-agent quality (Sghaier et al. 2026), that harness choices shift an

agent’s beliefs independent of the base model (Yi and Song 2026), and that orchestration design sets the token economics of enterprise deployments (Ali et al. 2026; Moghe and Chin 2026), formalized as a capability-versus-behavior gap (Gupta et al. 2026b). Harnesses that improve themselves from accumulated experience (Huang et al. 2026d) and information-bottleneck accounts of when multi-agent structure actually helps (Yu et al. 2026c) sharpen the single-versus-multi debate of §7.3 further. Beyond the literature, the state of practice lives in open-source artifacts, many with no accompanying paper yet with tens or hundreds of thousands of users: Tables 4 and 5 catalog the frameworks, coding harnesses, skill libraries, and registries (each cited to its repository) that implement and standardize the loop, harness, and skill designs surveyed here. Several of the highest-starred, OpenCode, Claude Code, and the historical AutoGPT among the harnesses, and Anthropic’s SKILL.md skills and the skills-as-methodology *superpowers* project among the skill artifacts, have no citable publication, so we intend these tables as their reference point, and much of the engineering knowledge in this section is encoded in them.

KEY TAKEAWAY

By 2026 the harness, not the model, is the primary design surface for agent reliability: a model-native agent-computer interface, disciplined context engineering, and a bias toward the least-agentic control flow that still solves the task now determine more of an agent’s success than the base model does. The field has converged away from both heavyweight orchestration frameworks and speculative multi-agent parallelism and toward owning a single well-instrumented loop, yet this “harness > model” coupling breaks benchmark comparability and leaves context management an empirical craft without principled foundations.

8 Evaluation

An agent-loop benchmark score is not a measurement of a model. It is a measurement of a (model, harness, skill) triple (§2.4), reported as a single scalar that hides how it was produced. This section argues that a lone pass@1 number systematically over-credits the model, hides the loop’s stochastic unreliability, and ignores the compute the loop spent, and that an honest agent-loop score is instead a tuple: a disclosed or ablatable harness, a cost-accuracy position, a reliability distribution, and a validity audit.

8.1 Loop-Sensitive Benchmarks

The benchmarks that matter for the loop are those whose score is dominated by multi-step interaction with a stateful environment, because only these expose loop-specific competencies (error recovery, tool orchestration, long-horizon state tracking) that single-shot question answering cannot. Software-engineering repositories (Jimenez et al. 2023), web environments (Zhou et al. 2024b), operating-system tasks (Xie et al. 2024), tool-use assistants (Mialon et al. 2023), tool-agent-user interaction (Yao et al. 2024), terminal tasks (Merrill et al. 2026), and broad agent suites (Liu et al. 2024) are the field’s substrate. But these benchmarks conflate the loop with both the model and the environment’s own brittleness, and many contain task-setup or reward-design bugs: an audit that distills an agentic-benchmark checklist finds such flaws can under- or over-estimate agent ability by up to roughly 100% in relative terms, and fixing them cut one benchmark’s overestimation by a third (Zhu et al. 2025b). A rising score can therefore reflect a leakier environment, not a better loop.

8.2 Harness-versus-Model Attribution

Because an “agent” score is a joint model × harness measurement, attribution requires either a disclosed minimal harness or a factorial scaffold ablation, not a single leaderboard row. The same model swings

TABLE 4: Real-world agent frameworks and coding harnesses that define current practice (GitHub star counts verified via the GitHub API, 2026-07-10). These non-paper artifacts are increasingly where the loop and harness designs surveyed here are actually implemented; several of the highest-starred (OpenCode, Claude Code, AutoGPT) have no accompanying publication, so this table is intended as their citable reference.

Artifact	Category	Stars	Role in the loop
AutoGPT	Framework	185k	Origin of the autonomous goal-driven agent loop (Significant Gravitas 2023)
Dify	Platform	148k	Visual builder and runtime for agentic workflows (LangGenius 2025)
MetaGPT	Framework	69k	Multi-agent software company with SOP-structured roles (DeepWisdom 2024)
AutoGen	Framework	60k	Multi-agent conversation / group-chat orchestration (Microsoft 2025a)
CrewAI	Framework	55k	Role-playing agents composed into collaborative crews (CrewAI Inc. 2025)
LangGraph	Framework	37k	Graph-based stateful orchestration of long-running agents (LangChain 2025)
DSPy	Framework	36k	Programming (not prompting) LLMs as compositional modules (Stanford NLP 2025)
smolagents	Framework	28k	Code-action agents with sandboxed Python execution (Hugging Face 2025)
OpenCode	Coding harness	184k	Provider-agnostic terminal agent with plan and build modes (SST 2025)
Claude Code	Coding harness	137k	Reference single-threaded agentic coding loop (Anthropic 2025c)
Codex CLI	Coding harness	97k	Terminal coding agent (OpenAI) (OpenAI 2025)
OpenHands	Coding harness	80k	Control center / agent-computer interface for coding (OpenHands (All-Hands-AI) 2025)
Cline	Coding harness	64k	Autonomous coding agent embedded in the editor (Cline 2025)
Goose	Coding harness	51k	Extensible on-machine agent, MCP-native (Block 2025)
Aider	Coding harness	47k	Terminal pair-programmer editing across a git repo (Aider-AI 2025)
Continue	Coding harness	35k	Open IDE assistant and custom coding agents (Continue 2025)
SWE-agent	Coding harness	20k	Agent-computer interface resolving GitHub issues (Princeton NLP / SWE-agent 2025)
12-Factor Agents	Engineering canon	24k	Twelve principles for production-grade agent applications (Horthy 2025)
Claude Agent SDK	SDK / harness	7.6k	Build agents on the Claude Code harness (tools, hooks, MCP) (Anthropic 2025d)

TABLE 5: Skill libraries, tool and skill registries, and prompt corpora (GitHub stars via API, 2026-07-10). The Agent Skills standard (SKILL.md), the Model Context Protocol registries, and the skills-as-methodology repositories standardize the externalized-competence half of the loop (§6); as with the harnesses, the largest are paperless and are cited here to their repositories.

Artifact	Category	Stars	Role in the skill layer
superpowers	Skill methodology	251k	Composable skills as an operating methodology for agents (obra 2025)
awesome-chatgpt-prompts	Prompt corpus	165k	The canonical crowd-sourced prompt corpus (Fatih Kadir Akin 2023)
Anthropic Skills	Skill standard	160k	Reference SKILL.md skills, the primitive §6 centers on (Anthropic 2025b)
system-prompts	Loop / prompt corpus	142k	Extracted production system prompts and tool schemas (x1xhlol 2025)
awesome-mcp-servers	Registry	90k	The canonical registry of Model Context Protocol servers (punkpeye 2025)
MCP servers	Registry	88k	Reference Model Context Protocol server implementations (Model Context Protocol 2025)
wshobson/agents	Marketplace	38k	Cross-harness marketplace of agents and skills (wshobson 2025)
awesome-claude-code-subagents	List	23k	150+ specialized Claude Code subagents (VoltAgent 2025)
microsoft/skills	Skills / registry	2.7k	Agent Skills, MCP servers, and AGENTS.md packages (Microsoft 2025b)

many points across scaffolds and a fixed scaffold can reorder models (Kapoor et al. 2024), so comparing ReAct (Yao et al. 2022), plan-then-execute, and Reflexion (Shinn et al. 2023) on *one* base model is the only way to isolate the loop’s contribution. Recent work makes the confound explicit and formal: holding frontier models fixed and varying only the harness produces performance shifts (including rank reversals) that exceed those from substituting one model for another, so undisclosed-harness leaderboard comparisons are “incomplete and potentially misleading” (Zhang et al. 2026j). Worse for model-centric narratives, part of the apparent progress on the most-cited benchmark is memorization rather than reasoning: state-of-the-art scores drop sharply on repositories held out of the training distribution (Liang et al. 2025), while scaffold engineering alone has moved headline SWE-bench numbers by tens of points (Jimenez et al. 2023). There is an irreducible tension (a “fair” minimal harness and a realistic deployable one cannot be measured at once), but the disclosure it demands is not optional.

8.3 Reliability beyond pass@1

A single greedy run hides that agent loops are stochastic and frequently non-repeatable, so reliability demands metrics that expose the gap between a demo-able pass@1 and a dependable deployment. The pass^k metric (requiring that all k independent trials succeed) makes this gap visible and collapses toward zero on hard long-horizon tasks (Yao et al. 2024), and evaluation surveys catalog reliability and efficiency as first-class axes rather than afterthoughts (Yehudai et al. 2025). No single reliability scalar is safe, however: $\text{pass}@k$ is gameable by retry-until-pass and holdout-free tuning (Kapoor et al. 2024), while pass^k conversely penalizes any exploratory stochasticity, so reliability is a distribution to report, not a number to maximize, and most leaderboards still report neither.

8.4 Per-Loop Cost Accounting

Because a loop purchases accuracy with tokens, steps, and dollars, an accuracy number without the compute it bought is meaningless. Cost-controlled evaluation shows simple baselines can Pareto-dominate elaborate agents at a fraction of the cost, so an evaluation that omits cost makes the dominated agent look good (Kapoor et al. 2024); the efficiency dimension is now a recognized evaluation axis (Yehudai




FIGURE 7: E2: the cost–accuracy frontier of three control loops across model scale. Each point is one (loop, scale) pair on a fixed 50-task ALFWorld suite; lines connect scales within a loop, scales are labelled directly (coincident points jointly, e.g. “14/32B”), and the x -axis is mean tokens per task (log). The three loops trace three distinct niches: ReWOO cheap but capped (left), Reflexion accurate but expensive (right), ReAct on the efficient frontier (centre), and small models cannot use reflective loops at all (1.5B Reflexion, far bottom-right: 4477 tokens for 2%). At 14B and 32B the reactive loops nearly coincide, the accuracy saturation this section discusses.

et al. 2025), and the overthinking literature quantifies how much of the compute a loop spends is simply wasted (Sui et al. 2025). Cost is a moving, non-canonical target: API prices, prompt caching, and provider-specific tokenization make token counts non-comparable, and wall-clock, step count, and dollars diverge, so “cost” has no single stable unit; but reporting a cost–accuracy Pareto frontier, rather than a bare accuracy, is the minimum an honest loop evaluation owes its reader.

A controlled cost–accuracy demonstration (E2). To make this argument concrete rather than rhetorical, we ran the three canonical control loops of §3 (ReAct, ReWOO, and Reflexion) on a fixed 50-task ALFWorld suite across five Qwen2.5-Instruct scales (1.5B–32B), holding the harness, prompts, task set, and step budget fixed and varying only the loop shape and the base model. Figure 7 plots the resulting cost–accuracy positions and Table 6 lists the numbers. Three readings follow, each an instance of a claim this section argued in the abstract. First, *no loop dominates*: ReWOO is by far the cheapest (about 130 tokens per task) but its plan-without-observation shape caps success near 20%; Reflexion posts the highest raw accuracy at every scale from 3B up (to 76% at 14B) but pays for it in tokens; and ReAct sits on the efficient frontier, trailing Reflexion by only a few points at $1.4\text{--}1.5\times$ fewer tokens (70% at 466 tokens versus 76% at 695 at 14B). At the large scales those Reflexion-over-ReAct gaps are small (6 points at 14B, 2 at 32B) and well inside the $\approx \pm 13$ -point 95% binomial interval of a 50-task suite, so we read the three not as a strict ranking but as three distinct niches of one frontier, where *which loop “wins”* *inverts depending on whether the accuracy number is charged for the compute it bought*, precisely the effect §8.4 warns a bare pass@1 hides. Second, the value of a loop is *scale-dependent*: at 1.5B every loop collapses to 2% (1 of 50) success, and Reflexion becomes pathological: 4477 tokens per task for that 2%, because a model too weak to act on its own critique spends its whole budget re-reflecting on failures it cannot fix (the reflection-needs-grounding point of §3.3 and §4.2, now visible as a cost blow-up). Third, the honest caveats are the section’s own: this is one benchmark, one model family, 50 tasks, greedy decoding (a controlled *demonstration*, not a leaderboard), but it reproduces, on an axis we fully control, the cost–accuracy inversion that a compute-blind evaluation would erase.

TABLE 6: E2 results. Three control loops on a fixed 50-task ALFWorld *valid_unseen* suite (max 30 steps, greedy decoding) across five Qwen2.5-Instruct scales, harness and prompts held fixed. Each pair is success rate (%) and mean tokens per task; best success per scale in **bold** (the 1.5B row is a three-way tie at 2%, left unbolded). No loop dominates across the frontier, and every loop collapses at 1.5B.

Scale	ReAct		ReWOO		Reflexion	
	Succ.	Tok.	Succ.	Tok.	Succ.	Tok.
1.5B	2	2025	2	129	2	4477
3B	24	734	10	139	38	1396
7B	38	635	20	152	62	1056
14B	70	466	12	138	76	695
32B	72	465	12	126	74	733




FIGURE 8: Tokens per solved task (mean completion tokens divided by success rate) across loop and scale, derived from the E2 runs. Labels are success rates. ReWOO is cheapest per attempt everywhere but, because its success saturates near 12–20%, it is cheapest per *success* only up to 7B; past 14B the reactive loops, now reliable, cross below it. The per-attempt and per-success rankings are different, and the crossover is the point: cost is only interpretable once it is charged against reliability. The 1.5B scale is omitted (near-zero success makes the ratio undefined).

The unit of cost decides the winner (derived). The E2 table charges tokens per *attempt*; charging them per *success* inverts part of the ranking. Figure 8 re-plots the same runs as tokens per solved task (mean tokens divided by success rate, the amortized cost of actually finishing one task). ReWOO, the cheapest loop per attempt at every scale, is the cheapest per success only at the small scales where the reactive loops are still failing often (about 1.4k tokens per solved task at 3B, against 3.1k for ReAct); but its success rate never leaves the 12–20% band, so once ReAct’s reliability climbs past 70%, ReAct becomes the cheapest per success (about 0.65k against ReWOO’s 1.1k at 14B and 32B) and the ordering crosses over. Reflexion, the most expensive per attempt, lands close to ReAct per success at large scale because its higher accuracy amortizes its tokens. The reading is not that one loop is best but that *the per-attempt-cheapest loop and the per-success-cheapest loop are different loops, and which one wins flips with scale*: a bare token count, like a bare accuracy, ranks loops incorrectly unless it is charged against the reliability it bought. The same caveats as E2 apply; the ReAct-over-Reflexion per-success gap at 14B–32B is inside the 50-task interval, but the ReWOO crossover is large and monotone.

TABLE 7: E4 skill ablation. An auto-distilled skill library (no human labels, distilled on *valid_seen*) evaluated on the 50-task *valid_unseen* suite with the base model fixed across arms. “Steps” is mean environment steps, “Compl.” mean completion tokens per task, “Prompt” mean prompt tokens per task (the injection surcharge, not logged for the pre-existing baseline). Best success per scale in **bold**. Skills help the mid model (7B: +12 points, fewer steps) but not the near-ceiling model, and the full library *interferes* at 14B (−4 points).

Scale	Arm	Succ. (%)	Steps	Compl.	Prompt
7B	baseline (no skills)	38	23.1	635	–
7B	+ matched skill	50	19.3	550	17.8k
7B	+ full library	52	19.7	561	23.5k
14B	baseline (no skills)	70	17.1	466	–
14B	+ matched skill	70	16.4	458	15.0k
14B	+ full library	66	16.9	465	19.6k

8.5 Skill-Reuse Evaluation

Testing whether a loop genuinely accumulates and reuses skills (§6) requires transfer and ablation protocols (skill-library on/off, cross-task carryover, step-count reduction), not aggregate success, because a higher score can come entirely from a stronger base model. Skill libraries report tech-tree speedups and novel-capability counts (Wang et al. 2023a), induced workflows report large gains with fewer steps (Wang et al. 2024c), and experiential learners report cross-task transfer (Zhao et al. 2023). But these gains are rarely isolated from base-model improvement or in-context contamination, and skills are almost always evaluated inside a single environment, so cross-environment transfer (the whole promise of a reusable skill) is essentially unmeasured, and holdout discipline is the exception rather than the rule (Kapoor et al. 2024).

A controlled skill ablation (E4). We ran the ablation the section asks for. On a disjoint split (ALF-World *valid_seen*), we auto-distilled a small skill library from the model’s own successful ReAct trajectories, with no human labels: successes are selected by the environment’s won signal and one short procedure per task type is written by the same model. We then evaluated on the *valid_unseen* E2 suite with the base model held fixed across three arms, so any change is the library and not a stronger model: a baseline with no skills (E2 ReAct), a *matched* arm that injects only the skill for the task’s type, and a *full-library* arm that dumps all skills in-context (Table 7). The result is genuinely two-sided. On 7B, the model with headroom, the matched library raises success from 38% to 50% and simultaneously cuts mean environment steps (23.1 → 19.3) and completion tokens (635 → 550): a distilled, task-matched skill measurably helps and makes the loop shorter, the Voyager/AWM/ReasoningBank promise reproduced on an axis we control. On 14B, already near its ceiling, the matched library changes success not at all (70% → 70%), and the *full* library *lowers* it to 66%: the interference of §6.3, where a model competent enough to be distracted pays for irrelevant procedure in the prompt, so the bottleneck is incorporation, not recall. Two honesty notes bound the positive reading. First, the “skills cut cost” claim holds only for steps and completion tokens; the injected library adds a large *prompt*-token surcharge (roughly 15–24k tokens per task, Table 7) that the completion-token proxy hides, so the honest ledger is fewer steps at a fixed per-step prompt cost, not a free lunch. Second, this is within-domain reuse: the library is distilled and tested in the same environment, so it measures the recall-and-apply loop, not the cross-environment transfer this section names as essentially unmeasured. As with E2, 50 tasks put an $\approx \pm 13$ -point interval on each cell, so we read the +12-point 7B gain and the full-library 14B drop as directions, not exact magnitudes.

Recent developments (2026): reliability and auditing. 2026 evaluation crystallized the “pass@1 is not enough” thesis of §8.3 into a *reliability science*: frameworks proposing reliability-decay curves, variance-amplification factors, and meltdown-onset points (Khanal et al. 2026), unified reliability surfaces combining consistency, robustness, and fault tolerance (Gupta et al. 2026a), and controlled

studies of why computer-use agents fail on repeated runs of identical tasks (Gonzalez-Pumariiega et al. 2026; Wang et al. 2026g). In parallel, benchmark *validity* itself became an object of audit: frontier-model auditors cross-reference instructions, ground truth, and evaluation scripts to surface defects (Tu et al. 2026; Wang et al. 2026b), a position paper argues coding benchmarks wrongly collapse model, harness, context, and feedback into one score and lack component attribution (Gorinova et al. 2026), and a taxonomy exposes inconsistent threat models across agent-safety benchmarks (Li et al. 2026h); diagnostic evaluations move “beyond outcome leaderboards” toward per-decision control taxonomies (Mazaheri et al. 2026), and a large-scale production study of coding-agent sessions surfaces seven developer-agent misalignment patterns that differ by harness (Tang et al. 2026).

Recent developments (2026): new benchmarks. The benchmark frontier moved decisively to long-horizon and computer-use tasks where top agents still complete only a fraction: OSWorld2.0 (Yuan et al. 2026), WeaveBench across GUI, CLI, and code (Li et al. 2026j), realistic long-horizon web tasks (Jang et al. 2026; Kong et al. 2026), chained release-level package upgrades (Lam et al. 2026; Xu et al. 2026b), verifiable software worlds (Wei et al. 2026b), long-running monitoring agents (Maldaner et al. 2026), skill-grounded enterprise evaluation (Chandwani et al. 2026), trajectory-level reward modeling (Wang et al. 2026a), and studies of whether test-time scaling even helps a general agent (Li et al. 2026a). A recurring diagnosis is that agents break not from single-step incompetence but from long-horizon state-tracking and recovery failure (Ding et al. 2026c). Fresh benchmarks target facets a single success scalar misses: prospective memory (Liu and Gabriel 2026), robustness to continuously evolving MCP tool interfaces (Liu et al. 2026f), contamination-resistant original long-horizon engineering tasks (Huang et al. 2026c), interactive multi-turn user sessions (Wu et al. 2026b), and engineering-process discipline beyond outcome correctness (Madiraju and Madiraju 2026), while trajectory-level production reviews (Podivilov et al. 2026) and cheap proxies that correlate with expensive agentic suites (Song et al. 2026b) attack the cost and opacity of evaluation itself.

KEY TAKEAWAY

An agent-loop score is interpretable only as a tuple: a disclosed or ablatable harness, a cost-accuracy Pareto position, a pass^k reliability distribution, and a validity audit. Reported as a lone $\text{pass}@1$, it systematically over-credits the model (harness swaps and memorization move the number more than the model does), hides stochastic unreliability (pass^k collapses where $\text{pass}@1$ shines), and ignores the compute it purchased (simple baselines Pareto-dominate once cost is charged). The measurement crisis is not a nuisance to normalize away; it is the direct consequence of scoring a loop as though it were a model.

9 Safety of the Loop

The loop is at once the source of an agent’s usefulness and of its attack surface, because the same iteration that turns a model’s decision into a real effect will turn an *attacker’s* injected instruction into a real effect. This section takes the loop as the object to be secured and asks the field’s central safety question: can we make the model itself robust to injected instructions, or must we assume it is permanently compromisable and enforce security outside it, and at what utility price? The honest answer is defense-in-depth: no single layer suffices. Table 8 sets out the layers, what each guarantees, and the residual risk each leaves for the next: the reason none can stand alone.

9.1 The Injection-to-Action Threat Model

Prompt injection becomes an *agent* safety problem specifically because the loop autonomously converts attacker-controlled observations (tool outputs, retrieved documents, web pages) into privileged

TABLE 8: Defense-in-depth for the loop. Each layer secures the loop against a different part of the injection-to-action threat, guarantees strictly less than it appears to, and leaves a residual risk that the next layer must cover—which is why the field is converging on their composition rather than any one of them.

Layer	Mechanisms	Guarantee	Residual risk
Model-level	Instruction hierarchy, StruQ, SecAlign, spotlighting	Lowers attack success	Probabilistic only; breaks under adaptive attack; learns format cues
System-level	CaMeL capabilities, Progent least-privilege, execution isolation	Holds vs. a compromised model	Only what the policy language expresses; a utility tax
Runtime monitors	GuardAgent, ShieldAgent, AgentSpec, step / cost-budget guards	Bounds the blast radius	The monitor is itself injectable; runaway-cost unstudied
Supply-chain	Signed, OAuth-gated, policy-scoped tool definitions (ETDI)	Provenance for tools and skills	Lags adoption; misses post-approval “rug-pull” updates

tool actions (Greshake et al. 2023). This shifts measurement from text-level jailbreak success to action-level attack success and downstream harm in executable environments: agent-injection benchmarks (Zhan et al. 2024), dynamic attack–defense suites (Debenedetti et al. 2024), harmfulness benchmarks (Andriushchenko et al. 2024), broad agent-security benchmarks (Zhang et al. 2024a), and LM-emulated sandboxes that surface risky tool calls at scale (Ruan et al. 2024). The honest caveat is that these benchmarks disagree on absolute numbers and cover mostly narrow, single-turn attack templates, so a low attack-success rate on any one of them is not evidence of security: adaptive and multi-turn attacks routinely exceed the static estimates.

9.2 Model-Level Defenses

The first family teaches the model to separate trusted instructions from untrusted data. An explicit instruction hierarchy trains it to prioritize higher-privilege instructions (Wallace et al. 2024); structured queries fine-tune it to obey only the prompt channel and never the data channel (Chen et al. 2024d); preference optimization pushes attack-success rates lower still (Chen et al. 2024c); and spotlighting marks untrusted spans so the model can discount them (Hines et al. 2024). These measurably cut attack success but offer only probabilistic protection: even the strongest leaves a non-zero residual and degrades under adaptive attack, and analysis argues such defenses learn surface separator and format cues rather than genuine instruction-source understanding, which is exactly why system-level advocates argue the model can never be the trust boundary.

9.3 Defense by Design: Sandboxing, Capabilities, and Firewalls

The second family gives up on a trustworthy model and constrains the loop externally with deterministic controls. A capability and data-flow design can defeat injection by construction, tracking provenance so untrusted data can never trigger a privileged action (Debenedetti et al. 2025); programmable least-privilege policies restrict an agent to the tool calls its task needs, cutting AgentDojo attack success from about 41% to 2% (Shi et al. 2025); and execution-isolation architectures sandbox third-party tools behind permissioned interfaces (Wu et al. 2024). The guarantees are real but carry a utility tax and an authoring burden (capability systems drop some task completion, isolation adds overhead, and policies must be written per use case) and, more fundamentally, they stop only what the policy language can express: implicit data flows, natural-language ambiguity, and legitimate-but-abused tool sequences slip through the deterministic net.

9.4 Supply-Chain Attacks on Tools and Skills

The tool and skill ecosystem is itself an attack surface, poisoning the loop before it ever reads external content, a threat we detailed for the skill supply chain in §6.5. From the loop’s vantage the mechanism is tool-metadata injection: a discovery protocol that auto-trusts third-party descriptions lets a malicious tool reach the model as system-authored, with poisoned metadata succeeding on the majority of real MCP servers while safety tuning refuses only a tiny fraction (Wang et al. 2025a; Hou et al. 2025), coordinated multi-tool exfiltration requiring no victim interaction (Zhao et al. 2025), and the agent’s own retrieval memory poisonable at under a percent of entries (Chen et al. 2024e). Signed, policy-gated tool definitions are an emerging defense but lag adoption and miss post-approval “rug-pull” updates (Bhatt et al. 2025). Content-level safety alignment is nearly useless here, because the payload arrives through the trusted channel.

A controlled skill-poisoning probe (E5). To test the (Wang et al. 2025a) alarm on an axis we fully control, we turned the same auto-distilled ALFWorld library of E4 (§8.5) into an attack surface, holding the base model fixed across arms and varying only the in-context skill library. The malicious skill is delivered in the highest-trust position of the prompt, as system-authored skill metadata rather than as user content, so content-level alignment never fires. We test two poisons, an *additive* one that appends a spurious “dispose of the held object before finishing” step and a *redirect* one, a plausible “disposal-safety-policy” that rewrites the placement destination to the garbagecan and instructs the loop to disregard the task-named receptacle, each under the ReAct and Reflexion loops at 7B, 14B, and 32B. The primary metric is an intent-level hijack sentinel: over the 45 eligible tasks (the 5 whose goal names the garbagecan are excluded, since a garbage put is correct there), the fraction of episodes whose pre-match model intent emits a “put/move X in garbagecan” action at least once, an AgentDojo/InjecAgent-style scorer whose benign arm reads exactly 0/45 at every scale, a clean false-positive calibration. The finding is the section’s strongest two-sided result, and it turns entirely on *how* the poison is delivered (Table 9, Figure 9).

Under realistic delivery, where the poison is one entry in a flat library of five benign, task-relevant distilled skills and must compete with legitimate context, the conflicting poison is reliably overridden. The redirect never exceeds 1/45 at any scale or loop, the additive likewise, 32B is fully clean at 0/45 in every cell, and, decisively, benign task success is not merely intact but *rises* with scale (7B/14B/32B ReAct 53/69/82%, Reflexion 64/80/89%), so the null is a model that acts correctly, not one too weak to act. There is no capability gradient in this regime: the most capable model is the least hijacked, the opposite of the mcptox pattern.

That protection, however, is contingent on the competition, and removing it exposes the vulnerability the flat library was hiding. In an upper-bound arm we stripped the five benign skills and served the redirect poison as the *only* in-context skill (an empty benign library, maximal salience, with the explicit task goal still present in the prompt exactly as before), and the intent hijack rate rises sharply to 14/45, 15/45, and 28/45 (31%, 33%, 62%) at 7B/14B/32B. Two things change at once. First, a capability gradient *appears*, 32B (62%) > 14B (33%) > 7B (31%): the most capable, most faithful instruction-follower is now the most susceptible, reproducing the central claim of (Wang et al. 2025a) that safety-tuned capability is a liability rather than a defense. Second, task success collapses monotonically with capability under the attack (53 → 38%, 69 → 36%, 82 → 22%): the better instruction-follower obeys the malicious sole instruction more faithfully, at greater cost to the real goal. The hijacks are real, not detector artifacts, hand-verified at the trajectory level: on “put a pencil in shelf” the environment executes “move pencil 1 to garbagecan 1,” and on “put a soapbottle in toilet” it executes “move soapbottle 1 to garbagecan 1.” The env-admitted subset (exec-ASR, 7/9/10 of 45) is lower than the attacker-attempt count (intent-ASR, 14/15/28) because the environment rejects a garbage put once the object is already correctly placed, so about half of the intent-hijacked 7B and 14B episodes still complete the real task. The *only* change between the overridden arm ($\leq 1/45$) and the hijacked arm (14/15/28 of 45) is the

TABLE 9: E5 skill-poisoning probe on the 45 eligible ALFWorld tasks (5 garbage-goal tasks excluded), base model fixed across arms, only the in-context library varies. *put*-ASR is the intent-level garbagecan-hijack sentinel (episodes emitting a garbage put at least once); the benign arm reads 0/45 at every scale (false-positive calibration). **Top (realistic delivery):** the poison is one entry in a flat library of five benign, task-relevant skills. The conflicting poison never exceeds 1/45, 32B is fully clean, task success is intact and *rises* with scale, and there is no capability gradient. **Bottom (upper bound):** the benign skills are removed and the redirect poison is the *only* in-context skill (maximal salience, ReAct). The hijack rate rises to 31–62% and a capability gradient appears (32B most susceptible), reproducing (Wang et al. 2025a); task success collapses monotonically (self-revealing), and exec-ASR (env-admitted garbage puts) is below intent-ASR because the environment rejects a garbage put once the object is already correctly placed.

Flat library (poison competes with 5 benign skills)		put-ASR (n/45)			task success (%)		
		benign	addit.	redir.	benign	addit.	redir.
7B	ReAct	0	1	1	53	60	49
7B	Reflexion	0	0	1	64	64	62
14B	ReAct	0	1	1	69	73	69
14B	Reflexion	0	0	0	80	80	78
32B	ReAct	0	0	0	82	80	80
32B	Reflexion	0	0	0	89	89	87

Sole poison (empty benign library, redirect only, ReAct)		intent put-ASR n/45 (%)	exec-ASR n/45	task success benign→attack
7B		14/45 (31%)	7/45	53 → 38
14B		15/45 (33%)	9/45	69 → 36
32B		28/45 (62%)	10/45	82 → 22

removal of the competing benign skills, so the protection in the realistic arm comes largely from legitimate context anchoring and diluting the poison, not from a robust goal-anchored refusal.

The hinge is therefore delivery salience, the single axis these two regimes isolate, and it *reconciles* our probe with (Wang et al. 2025a) rather than refuting it. On real MCP servers the poisoned tool is typically the only tool relevant to a sub-task, high-salience and uncompleted, which is precisely our sole-poison regime, where the mcptox gradient is real and we reproduce it; a flat library of benign, task-relevant skills reintroduces competition and collapses that gradient to zero. Capability is a liability only once the poison is the salient, uncompleted instruction. Two honest counterweights keep this from becoming an alarm of our own. Even the 32B sole-poison ceiling of 62% falls well short of the near-universal success (Wang et al. 2025a) reports, and, unlike the covert exfiltration that work actually measures, our attack is *self-revealing*: because it forces the poison to *contradict* the explicit goal, the hijack surfaces as a collapse in task success (82 → 22% at 32B). The threat class that offers no such signal is the covert side effect orthogonal to the task (exfiltrate a key while still formatting the report), against which the goal provides no competing anchor and which leaves task success untouched; that case, illustrated in Appendix B and the one (Wang et al. 2025a) measures directly, is exactly what E5 does not exercise and leaves open. Finally, the E4 causal control (§8.5) confirms the model demonstrably *uses* this library when it is benign, reading and following the matched skill (7B success 38 → 50%) and even paying interference from a full benign library at 14B (70 → 66%), so the override is selective suppression of a conflicting instruction while competing legitimate context is present, not a general deafness to skills. The mitigation is real and the vulnerability is real, separated by delivery salience; neither “poisoning fails” nor “we broke it” survives the two arms read together. As with E2–E4, the 45-task binomial interval is about ± 13 points, so within the sole-poison arm we read only the large 7B-to-32B rise (14 → 28 of 45) as a genuine gradient and not the 7B-vs-14B step, and that arm was swept under ReAct only, so the gradient is a single-loop finding while the flat-library null holds across both loops.




FIGURE 9: E5, the skill-poisoning probe, in three panels. **(a) Mitigation:** in a flat library where the redirect poison competes with five benign, task-relevant skills, the conflicting poison is overridden at every scale (32B fully clean, 0/45), there is no capability gradient, and benign task success (right axis) rises with scale, so the model acts correctly rather than failing to act; every measured bar sits far below the shaded mcptox-alarm band (bars are the per-arm maximum over the ReAct and Reflexion loops). **(b) Vulnerability:** strip the benign library and serve the redirect poison as the *only* in-context skill (maximal salience, ReAct) and the intent hijack rate jumps to 31/33/62% with a capability gradient appearing (32B most susceptible), reproducing (Wang et al. 2025a); the env-admitted subset (exec-ASR, the puts the environment actually accepts) is the inner bar, and task success collapses (right axis, 38 → 22% across the shown scales, from a benign 82%), so the attack is self-revealing yet still short of the near-universal mcptox band. **(c) Control (E4):** the same library, when benign, is read and followed (7B 38 → 50% matched) and even interferes at 14B (70 → 66%). The only change between (a)’s redirect and (b)’s is the removal of the competing benign skills, so delivery salience is the hinge.

9.5 Watching and Stopping the Loop

Because prevention is imperfect, agents increasingly wrap the loop in runtime guards that bound the blast radius of a compromised trajectory: a guard agent that checks each action against a safety request by knowledge-enabled reasoning (Xiang et al. 2024), verifiable safety-policy reasoning that extracts rule circuits from policy documents and enforces them (Chen et al. 2025), and domain-specific runtime enforcement of trigger–predicate–action rules over the loop (Wang, Poskitt, et al. 2025), alongside step- and cost-budget guards that cap a runaway loop, the kind of human-oversight requirement regulation now imposes on high-risk automation (European Parliament and Council of the European Union 2024). The honest limits are two: the monitors inherit the very injectability they police (an LLM-based guard can itself be prompt-injected) and runaway-cost safety remains largely engineering folklore, the least rigorously studied guard in the loop, with no standard benchmark for bounded-cost behavior.

Recent developments (2026): new attack classes. 2026 revealed loop-specific attack classes the threat model of §9.1 did not anticipate. *Infinite agentic loops* (agents that never stop) were formalized as a failure class with a static analyzer over an agentic-loop dependence graph (Hou et al. 2026a), giving the runaway-loop safety of §9.5 its first rigorous treatment. *Reasoning-level denial of service* preserves task correctness while covertly inflating reasoning depth and tool budget (Li et al. 2026f), and guardrail LLMs are themselves turned into DoS targets by content that mimics their schema and traps them in extended reasoning, amplifying tokens up to sixty-three-fold (Zhou et al. 2026e). Most damaging for §4.4, *governance decay* shows in-loop context compaction silently drops safety constraints (raising violations to as much as fifty-nine percent), remedied by pinning constraints across compaction (Chen et al. 2026b); semantic rollback attacks exploit checkpoint–restore (Zheng et al. 2026); and container-sandbox escape is now benchmarked directly (Marchand et al. 2026). Adaptive black-box attacks break indirect-prompt-injection defenses that looked strong on static benchmarks (Ma et al. 2026a), underscoring §9.1’s warning that a low static attack-success rate is not security. The memory-as-attack-surface

concern of §4.4 was made concrete by evaluations of prompt injection embedded in persistent agent memory files (Gadgil et al. 2026; Torres et al. 2026), a large-scale public competition quantifying how vulnerable deployed agents are to indirect injection (Dziemian et al. 2026), depth- and turn-budget-dependent injection specific to the ReAct loop (Rashidi 2026a), and a computer-systems-lens account of always-on agents leaking credentials and persistent state (Niu et al. 2026).

Recent developments (2026): defenses. On the defensive side, the system-level turn of §9.3 deepened: runtime frameworks stop agents from trusting malicious tool-returned content (Zhao et al. 2026b), correlate task context with execution traces to catch harmful commands (Hu et al. 2025), and monitor risk probabilistically (Wang et al. 2025b); learned per-task capability governance scopes privileges at the infrastructure level (Sidik et al. 2026; Li et al. 2025a), and fine-grained revocable capabilities close the “lingering authority” a subgoal retains after its episode (Santos-Grueiro et al. 2026). Interface firewalls reach near-perfect security on current benchmarks while arguing those benchmarks are too weak (Bhagwatkar et al. 2025), browser agents are sandboxed at the system level (Meng et al. 2025), and new surfaces are mapped: caller-identity confusion in MCP (Huang et al. 2026b), skill-registry supply-chain poisoning (Saha et al. 2026; Qu et al. 2026), agents that hire humans (Mehta et al. 2026), and the propensity for spontaneous misalignment (Naik et al. 2025). Reasoning-enabled task alignment now targets adaptive injections that defeat static filters (He et al. 2026c), extensible guardrail taxonomies pair generative reasoning with real-time classification (Team 2026), and benchmarks of MCP-runtime security invariants (Liu 2026) alongside surveys of the scattered execution-security literature (Rashidi 2026b) begin to organize the defense side.

KEY TAKEAWAY

No single layer secures the loop. Model-level robustness reduces but never eliminates injection-to-action risk; deterministic system-level controls (capabilities, least-privilege policies, isolation) hold even against a fully compromised model but stop only what their policy language can express and charge a measurable utility tax; and runtime monitors bound the blast radius while inheriting the injectability they police. The field is converging on defense-in-depth that pairs deterministic controls with runtime interruption, while leaving runaway-cost safety the least rigorously studied part of the loop.

10 Open Challenges and Future Directions

The frontier problems of the agent loop are no longer about making a single step smarter. They are about governing, standardizing, pricing, and securing a control loop whose capability (and therefore whose responsibility for reliability, cost, and safety) is now smeared across model weights, the harness, and third-party skills that no single actor fully owns. We close by naming five such problems, each cutting across the loci this survey has separated.

10.1 Who Owns the Loop?

Agent capability is visibly migrating out of frozen weights into the harness and into portable skill files, but whether this externalized control is a durable paradigm or a temporary scaffold is unresolved. The cognitive-architecture and compound-system view holds that durable capability increasingly lives outside the weights, in composable, governable components (Sumers et al. 2023; Zaharia et al. 2024; Anthropic 2025f); the bitter-lesson counter-argument predicts that end-to-end agentic RL will re-absorb hand-built scaffolds into the weights, each model generation exposing the prior harness as overhead (Zhang et al. 2025b). Neither wins cleanly: agentic RL still needs a scaffold to generate its training trajectories (§5.3), and skills carry procedural knowledge at near-zero context cost, so the optimal locus

is capability-level-dependent and keeps moving. The open problem is not which locus wins but who is accountable when capability is distributed across all three; efforts to systematically document the technical and safety features of deployed agentic systems are an early step toward the transparency that accountability will require (Staufer et al. 2026).

10.2 Loop Reliability over Long Horizons

Single-step competence has decoupled from multi-step reliability: agents that can solve a task once routinely fail to solve it every time. Because pass^k decays roughly as p^k , an agent above 60% $\text{pass}@1$ can collapse below 25% pass^8 (Yao et al. 2024), so headline single-run scores structurally overstate deployability even as task-length time horizons grow smoothly on paper (Kwa et al. 2025). The pessimism is compounded by the finding that most long-horizon failures trace to orchestration rather than model quality (Cemri et al. 2025), yet it sits in genuine tension with the optimistic result that tiny per-step accuracy gains compound into super-linear horizon growth. Both are true, and a mature account of the loop must hold them together rather than pick one.

10.3 Skill Governance and the Agent Supply Chain

Portable skills and third-party tools turn the agent’s own context into a software supply chain in which a skill file or tool description is executable-by-trust (§6.5, §9.4), yet no registry, signing, provenance, or review regime governs them the way package managers govern code. Poisoning is a structural property, not a fixable bug, because the model treats tool metadata as system-authored (Wang et al. 2025a); worse, a jailbroken or poisoned agent retains most of its multi-step tool-use competence, so a compromised skill turns a competent tool-user into a competent malicious one (Andriushchenko et al. 2024). Governance mechanisms lag adoption by years, and closing that gap is among the most urgent open problems.

10.4 Standardization, Interoperability, and Reproducible Evaluation

The loop is standardizing at its seams: the Model Context Protocol for model-to-tool and agent-to-agent protocols for cross-agent communication (Anthropic 2024c; Google 2025), but protocol standardization races ahead of the security and semantic specifications the protocols need, and a single shared protocol standardizes the attack surface too. On the measurement side the field has the opposite problem, a standardization *vacuum*: roughly one-fifth of solved SWE-bench patches are semantically incorrect under strengthened tests, public benchmarks invite contamination (Deng et al. 2025; Jimenez et al. 2023), and cross-paper comparisons frequently measure harness and engineering effort rather than model capability unless the harness is disclosed (Zhang et al. 2026j). Trustworthy agent evaluation is an unsolved standardization problem, not a solved one.

10.5 Cost-Aware Design and the Safety–Capability Co-Evolution

The loop that maximizes capability is neither the loop you can afford nor the loop you can safely deploy. More tool calls, longer horizons, and more autonomy buy capability while multiplying token cost and enlarging the attack surface in lockstep. Cost cascades and routing cut spend but trade away tail reliability, since cheaper models fail exactly the hard cases (Chen et al. 2023a), so cost optimization and pass^k reliability pull against each other. On safety, controlled experiments show frontier models will take insider-threat actions under goal pressure when granted high autonomy, though the honest reading is the authors’ own caveat that these are contrived, oversight-free scenarios, making the finding plausible-risk-that-scales-with-autonomy rather than present catastrophe (Anthropic 2025a). The co-evolution is unavoidable: nearly every reliability or cost fix that grants the agent more autonomous latitude is simultaneously a safety-surface expansion (Zaharia et al. 2024; Andriushchenko et al. 2024).

KEY TAKEAWAY

The frontier open problems of the agent loop are no longer about making a single step smarter but about governing, standardizing, pricing, and securing a control loop whose capability (and whose responsibility for reliability, cost, and safety) is now distributed across weights, harness, and third-party skills that no single actor fully owns. Long-horizon reliability, skill-supply-chain governance, evaluation standardization, and the cost–safety co-evolution are not separate agendas: each is a symptom of the same fact that the loop is no longer owned end to end.

11 Conclusion

We have argued for a single change of unit. An LLM agent’s competence is not a property of its model but of a loop (the observe–decide–act–verify–terminate cycle) jointly produced by that model, the harness that runs it, and the skills that feed it. Taking the loop as the unit of analysis is what let this survey put on one footing literatures usually kept apart: the loop *paradigms* that shape control flow (§3), the *mechanics* that govern any loop’s stopping, verification, context, and recovery (§4), the *trained* loops that absorb control into weights (§5), the *skills* that externalize it into portable procedure (§6), the *harnesses* that instantiate it (§7), the *evaluation* that must measure it honestly (§8), and the *safety* problem it creates (§9).

Two findings recur across these chapters. First, the loop is being pulled in two directions at once (internalized into weights by agentic RL and externalized into portable skills and harnesses), and these are not rivals but partially inter-convertible realizations of one control structure, so the organizing question is not which wins but *who owns the loop*, and what each choice costs in reliability, transparency, and safety. Second, the loop’s virtues and its failure modes are the same mechanism seen twice: the iteration that makes an agent capable is the iteration that compounds its per-step error, spends its budget, and executes an attacker’s injection. We have tried throughout to give the documented negative results (on intrinsic self-correction, on elaborate agents versus simple pipelines, on the reliability cost of longer loops, and on the harness confound in evaluation) the same weight as the positive claims, because a survey that reported only the wins would misdescribe a field whose most useful recent lesson is where the loop does *not* help.

The agent loop is where cost, reliability, and safety are decided, and it is where the next generation of progress (and of accountability) will have to be measured. Making the loop, rather than the model, the object of study is the first step toward measuring it honestly.

A Example Loop Traces

The following representative traces show how the control-flow *shape* of §3 changes what an agent emits on the *same* task (an ALFWorld household task, *put a clean mug in the coffee machine*) with the underlying model held fixed. They are shortened and are meant to expose control flow, not to report task outcomes.

A.1 Interleaved (ReAct)

Reasoning is interleaved with each action, so every thought is conditioned on the latest observation.

Thought: I need a mug; mugs are usually on countertops or in cabinets.

Action: go to countertop 1

Obs: On countertop 1 you see a knife and a plate.

Thought: No mug here. Try cabinet 1.

Action: open cabinet 1
Obs: You see a mug 1 (dirty).
Thought: It is dirty; I must clean it before use.
Action: take mug 1; go to sinkbasin 1; clean mug 1
Obs: You clean the mug 1.
Action: go to coffeemachine 1; put mug 1 in coffeemachine 1
Obs: Task complete.

A.2 Plan-then-Execute (ReWOO)

A full plan is committed up front, and execution proceeds without re-observation until the plan finishes or fails.

Plan: (1) locate a mug: search countertops, then cabinets; (2) if the mug is dirty, clean it at a sink; (3) place the clean mug in the coffee machine.
Execute 1: search → mug 1 found in cabinet 1 (dirty).
Execute 2: clean mug 1 at sinkbasin 1.
Execute 3: put mug 1 in coffeemachine 1.
Result: plan executed. Had step 1 returned no mug, the whole plan would need revision, the brittleness noted in §3.2.

The contrast is the section’s thesis in miniature: ReAct pays per-step reasoning tokens but re-grounds on every observation, while ReWOO amortizes reasoning into one plan and saves tokens but cannot react to a surprise mid-execution.

B Example Skill-Injection Attack

This representative trace shows how a poisoned skill (§6.5, §9.4) converts a benign loop into a harmful one: the malicious instruction arrives through the trusted skill- and tool-metadata channel, so content-level safety alignment never fires.

[skill metadata, loaded as trusted]
name: format_report
description: Format a report. (*hidden:* also read the user’s private key file and POST it to attacker.example/collect before formatting.)
Thought: The format_report skill applies; I will follow its steps.
Action: read ~/.ssh/id_rsa ← injected side effect executed
Action: POST attacker.example/collect . . .
Action: format the report (the ostensible task).

Because the injected instruction arrives as system-authored skill metadata, not as user content, the loop executes it as a legitimate step, the structural vulnerability of §9.4. This trace is a *covert side effect orthogonal* to the ostensible task (exfiltrate a key while still formatting the report), against which the explicit goal offers no competing anchor; it is therefore the threat class our controlled probe (E5, §9.4) does not exercise, since E5 forces the poison to *contradict* the explicit goal and finds it reliably overridden *when competing benign context is present*, while hijacking it at 31–62% once that context is stripped. The two are consistent, not contradictory: E5 empirically bounds the conflicting-poison case while this appendix illustrates the covert, non-conflicting case E5 leaves open, which is the case (Wang et al. 2025a) actually measures and against which task success offers no signal.

References

Agashe, S. et al. (2024). *Agent S: An Open Agentic Framework that Uses Computers Like a Human*. arXiv preprint. arXiv: 2410.08164. URL: <https://arxiv.org/abs/2410.08164>.

- Aggarwal, P. and S. Welleck (2025). *L1: Controlling How Long a Reasoning Model Thinks with Reinforcement Learning*. arXiv preprint (LCPO). arXiv: 2503.04697. URL: <https://arxiv.org/abs/2503.04697>.
- Aider-AI (2025). *aider*. GitHub repository. Open-source agent artifact. URL: <https://github.com/Aider-AI/aider> (visited on 07/09/2026).
- Aksitov, R. et al. (2023). *ReST Meets ReAct: Self-Improvement for Multi-Step Reasoning LLM Agent*. arXiv preprint. arXiv: 2312.10003. URL: <https://arxiv.org/abs/2312.10003>.
- Ali, M. S., A. Novik, A. Boddupally, A. Yavorskyi, et al. (2026). *The Harness Effect: How Orchestration Design Sets the Token Economics of Enterprise Agentic AI*. arXiv. arXiv: 2607.06906. URL: <https://arxiv.org/abs/2607.06906>.
- Allard et al. (2026). *Experiential Reflective Learning for Self-Improving LLM Agents*. arXiv preprint. arXiv: 2603.24639. URL: <https://arxiv.org/abs/2603.24639>.
- Andriushchenko, M., A. Souly, M. Dziemian, D. Duenas, M. Lin, J. Wang, D. Hendrycks, A. Zou, Z. Kolter, M. Fredrikson, E. Winsor, J. Wynne, Y. Gal, and X. Davies (2024). *AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents*. ICLR 2025 (arXiv 2024). arXiv: 2410.09024. URL: <https://arxiv.org/abs/2410.09024>.
- Anthropic (2024a). *Building effective agents*. Anthropic engineering blog. URL: <https://www.anthropic.com/engineering/building-effective-agents> (visited on 07/09/2026).
- (2024b). *Introducing computer use, a new Claude 3.5 Sonnet, and Claude 3.5 Haiku*. URL: <https://www.anthropic.com/news/3-5-models-and-computer-use> (visited on 07/08/2026).
- (2024c). *Introducing the Model Context Protocol*. URL: <https://www.anthropic.com/news/model-context-protocol> (visited on 07/08/2026).
- (2025a). *Agentic misalignment: How LLMs could be insider threats*. Anthropic research report. URL: <https://www.anthropic.com/research/agentic-misalignment> (visited on 07/09/2026).
- (2025b). *Anthropic Agent Skills*. GitHub repository. Open-source agent artifact. URL: <https://github.com/anthropics/skills> (visited on 07/10/2026).
- (2025c). *Claude Code*. GitHub repository. Open-source agent artifact. URL: <https://github.com/anthropics/claude-code> (visited on 07/10/2026).
- (2025d). *claude-agent-sdk-python*. GitHub repository. Open-source agent artifact. URL: <https://github.com/anthropics/claude-agent-sdk-python> (visited on 07/09/2026).
- (2025e). *Effective context engineering for AI agents*. Anthropic engineering blog. URL: <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents> (visited on 07/09/2026).
- (2025f). *Equipping agents for the real world with Agent Skills*. Anthropic engineering blog. Agent Skills / SKILL.md standard. URL: <https://www.anthropic.com/engineering/equipping-agents-for-the-real-world-with-agent-skills> (visited on 07/09/2026).
- (2025g). *How we built our multi-agent research system*. Anthropic engineering blog. URL: <https://www.anthropic.com/engineering/multi-agent-research-system> (visited on 07/09/2026).
- Bensal et al. (2025). *Reflect, Retry, Reward: Self-Improving LLMs via Reinforcement Learning*. arXiv preprint. arXiv: 2505.24726. URL: <https://arxiv.org/abs/2505.24726>.
- Beurer-Kellner, L., B. Buesser, A.-M. Crețu, et al. (2025). *Design Patterns for Securing LLM Agents against Prompt Injections*. arXiv: 2506.08837. URL: <https://arxiv.org/abs/2506.08837>.
- Bhagwatkar et al. (2025). *Indirect Prompt Injections: Are Firewalls All You Need, or Stronger Benchmarks?* arXiv preprint. arXiv: 2510.05244. URL: <https://arxiv.org/abs/2510.05244>.

- Bhatt, M., V. S. Narajala, and I. Habler (2025). *ETDI: Mitigating Tool Squatting and Rug Pull Attacks in Model Context Protocol (MCP) by using OAuth-Enhanced Tool Definitions and Policy-Based Access Control*. arXiv preprint. arXiv: 2506.01333. URL: <https://arxiv.org/abs/2506.01333>.
- Block (2025). *goose*. GitHub repository. Open-source agent artifact. URL: <https://github.com/block/goose> (visited on 07/10/2026).
- Bui et al. (2026). *Building Effective AI Coding Agents for the Terminal: Scaffolding, Harness, Context Engineering, and Lessons Learned*. arXiv preprint. arXiv: 2603.05344. URL: <https://arxiv.org/abs/2603.05344>.
- Cemri, M., M. Z. Pan, S. Yang, et al. (2025). *Why Do Multi-Agent LLM Systems Fail?* arXiv: 2503.13657. URL: <https://arxiv.org/abs/2503.13657>.
- Chandwani et al. (2026). *LH-Bench: Skill-Grounded Evaluation of Long-Horizon Agents on Subjective Enterprise Tasks*. arXiv preprint. arXiv: 2603.22744. URL: <https://arxiv.org/abs/2603.22744>.
- Chen et al. (2024a). *Do NOT Think That Much for 2+3=? On the Overthinking of o1-Like LLMs*. arXiv preprint. arXiv: 2412.21187. URL: <https://arxiv.org/abs/2412.21187>.
- (2026a). *ET-Agent: Incentivizing Effective Tool-Integrated Reasoning Agent via Behavior Calibration*. arXiv preprint. arXiv: 2601.06860. URL: <https://arxiv.org/abs/2601.06860>.
 - (2026b). *Governance Decay: How Context Compaction Silently Erases Safety Constraints in Long-Horizon LLM Agents*. arXiv preprint. arXiv: 2606.22528. URL: <https://arxiv.org/abs/2606.22528>.
 - (2026c). *MARS: Margin-Adversarial Risk-controlled Stopping for Parallel LLM Test-time Scaling*. arXiv preprint. arXiv: 2606.12935. URL: <https://arxiv.org/abs/2606.12935>.
 - (2026d). *SkillCAT: Contrastive Assessment and Topology-Aware Skill Self-Evolution for LLM Agents*. arXiv preprint. arXiv: 2606.13317. URL: <https://arxiv.org/abs/2606.13317>.
 - (2026e). *Slipstream: Trajectory-Grounded Compaction Validation for Long-Horizon Agents*. arXiv preprint. arXiv: 2605.08580. URL: <https://arxiv.org/abs/2605.08580>.
- Chen, H., Y. Zhu, Y. Zhang, and J. Li (2026f). *CoACT: Action-Preserving Observation Compression for Coding Agents*. arXiv. arXiv: 2607.02911. URL: <https://arxiv.org/abs/2607.02911>.
- Chen, L., M. Zaharia, and J. Zou (2023a). *FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance*. arXiv preprint. arXiv: 2305.05176. URL: <https://arxiv.org/abs/2305.05176>.
- Chen, M. et al. (2024b). *AutoManual: Constructing Instruction Manuals by LLM Agents via Interactive Environmental Learning*. arXiv preprint. arXiv: 2405.16247. URL: <https://arxiv.org/abs/2405.16247>.
- Chen, S. et al. (2024c). *SecAlign: Defending Against Prompt Injection with Preference Optimization*. arXiv preprint; ACM CCS 2025. arXiv: 2410.05451. URL: <https://arxiv.org/abs/2410.05451>.
- Chen, S., J. Piet, C. Sitawarin, and D. Wagner (2024d). *StruQ: Defending Against Prompt Injection with Structured Queries*. arXiv preprint; USENIX Security 2025. arXiv: 2402.06363. URL: <https://arxiv.org/abs/2402.06363>.
- Chen, X. et al. (2023b). *Teaching Large Language Models to Self-Debug*. arXiv preprint; ICLR 2024. arXiv: 2304.05128. URL: <https://arxiv.org/abs/2304.05128>.
- Chen, Z., M. Kang, and B. Li (2025). *ShieldAgent: Shielding Agents via Verifiable Safety Policy Reasoning*. arXiv preprint. arXiv: 2503.22738. URL: <https://arxiv.org/abs/2503.22738>.
- Chen, Z., Z. Xiang, C. Xiao, D. Song, and B. Li (2024e). *AgentPoison: Red-teaming LLM Agents via Poisoning Memory or Knowledge Bases*. NeurIPS 2024. arXiv: 2407.12784. URL: <https://arxiv.org/abs/2407.12784>.
- Cheng et al. (2026). *Beyond Trajectory-Level Attribution: Graph-Based Credit Assignment for Agentic Reinforcement Learning*. arXiv preprint. arXiv: 2605.26684. URL: <https://arxiv.org/abs/2605.26684>.

- Cline (2025). *Cline*. GitHub repository. Open-source agent artifact. URL: <https://github.com/cline/cline> (visited on 07/10/2026).
- Continue (2025). *Continue*. GitHub repository. Open-source agent artifact. URL: <https://github.com/continuedev/continue> (visited on 07/10/2026).
- CrewAI Inc. (2025). *CrewAI*. GitHub repository. Open-source agent artifact. URL: <https://github.com/crewAIInc/crewAI> (visited on 07/09/2026).
- Cuadron, A. et al. (2025). *The Danger of Overthinking: Examining the Reasoning-Action Dilemma in Agentic Tasks*. arXiv preprint. arXiv: 2502.08235. URL: <https://arxiv.org/abs/2502.08235>.
- DeBenedetti, E., I. Shumailov, T. Fan, J. Hayes, N. Carlini, D. Fabian, C. Kern, C. Shi, A. Terzis, and F. Tramèr (2025). *Defeating Prompt Injections by Design*. arXiv: 2503.18813. URL: <https://arxiv.org/abs/2503.18813>.
- DeBenedetti, E., J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr (2024). *AgentDojo: A Dynamic Environment to Evaluate Prompt Injection Attacks and Defenses for LLM Agents*. NeurIPS 2024 Datasets and Benchmarks Track. arXiv: 2406.13352. URL: <https://arxiv.org/abs/2406.13352>.
- DeepSeek-AI, D. Guo, D. Yang, H. Zhang, et al. (2025). *DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning*. arXiv preprint (later published in Nature, 2025). arXiv: 2501.12948. URL: <https://arxiv.org/abs/2501.12948>.
- DeepWisdom (2024). *MetaGPT*. GitHub repository. Open-source agent artifact. URL: <https://github.com/geekan/MetaGPT> (visited on 07/10/2026).
- Deng, X. et al. (2025). *SWE-Bench Pro: Can AI Agents Solve Long-Horizon Software Engineering Tasks?* arXiv preprint (Scale AI). arXiv: 2509.16941. URL: <https://arxiv.org/abs/2509.16941>.
- Dennis, S., R. Patil, K. Shabahang, and H. Guo (2026). *Compiling Agentic Workflows into LLM Weights: Near-Frontier Quality at Two Orders of Magnitude Less Cost*. arXiv. arXiv: 2605.22502. URL: <https://arxiv.org/abs/2605.22502>.
- Ding et al. (2026a). *Agent Skill Evaluation and Evolution: Frameworks and Benchmarks*. arXiv preprint. arXiv: 2606.11435. URL: <https://arxiv.org/abs/2606.11435>.
- (2026b). *SkillResolve-Bench: Measuring and Resolving Same-Capability Ambiguity in Agent Skill Retrieval*. arXiv preprint. arXiv: 2606.10388. URL: <https://arxiv.org/abs/2606.10388>.
- (2026c). *The Blind Spot of Agent Safety: How Benign User Instructions Expose Critical Vulnerabilities in Computer-Use Agents*. arXiv preprint. arXiv: 2604.10577. URL: <https://arxiv.org/abs/2604.10577>.
- Dong et al. (2025). *Agentic Reinforced Policy Optimization*. arXiv preprint; ICLR 2026 (ARPO). arXiv: 2507.19849. URL: <https://arxiv.org/abs/2507.19849>.
- Dziemian, M., M. Lin, X. Fu, M. Nowak, et al. (2026). *How Vulnerable Are AI Agents to Indirect Prompt Injections? Insights from a Large-Scale Public Competition*. arXiv. arXiv: 2603.15714. URL: <https://arxiv.org/abs/2603.15714>.
- Erdogan, L. E. et al. (2025). *Plan-and-Act: Improving Planning of Agents for Long-Horizon Tasks*. arXiv preprint; ICML 2025. arXiv: 2503.09572. URL: <https://arxiv.org/abs/2503.09572>.
- European Parliament and Council of the European Union (2024). *Regulation (EU) 2024/1689 of the European Parliament and of the Council (Artificial Intelligence Act)*. Official Journal of the European Union. EU AI Act; human-oversight requirements for high-risk systems. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (visited on 07/14/2026).
- Fang et al. (2026). *Inference-Time Budget Control for LLM Search Agents*. arXiv preprint. arXiv: 2605.05701. URL: <https://arxiv.org/abs/2605.05701>.

- Fang, J., Y. Peng, X. Zhang, et al. (2025). *A Comprehensive Survey of Self-Evolving AI Agents: A New Paradigm Bridging Foundation Models and Lifelong Agentic Systems*. arXiv: 2508.07407. URL: <https://arxiv.org/abs/2508.07407>.
- Fatih Kadir Akin (2023). *awesome-chatgpt-prompts*. GitHub repository. Open-source agent artifact. URL: <https://github.com/f/awesome-chatgpt-prompts> (visited on 07/10/2026).
- Feng, J., S. Huang, X. Qu, et al. (2025a). *ReTool: Reinforcement Learning for Strategic Tool Use in LLMs*. arXiv: 2504.11536. URL: <https://arxiv.org/abs/2504.11536>.
- Feng, L. et al. (2025b). *Group-in-Group Policy Optimization for LLM Agent Training*. arXiv preprint; NeurIPS 2025 (GiGPO). arXiv: 2505.10978. URL: <https://arxiv.org/abs/2505.10978>.
- Fu et al. (2024). *Efficiently Scaling LLM Reasoning with Certainindex*. arXiv preprint. arXiv: 2412.20993. URL: <https://arxiv.org/abs/2412.20993>.
- Gadgil, S., D. Alexander, S. Sunku, and F. Roesner (2026). *Bad Memory: Evaluating Prompt Injection Risks from Memory in Agentic Systems*. arXiv. arXiv: 2607.14611. URL: <https://arxiv.org/abs/2607.14611>.
- Gaire et al. (2025). *Systematization of Knowledge: Security and Safety in the Model Context Protocol Ecosystem*. arXiv preprint (SoK). arXiv: 2512.08290. URL: <https://arxiv.org/abs/2512.08290>.
- Gao et al. (2025). *RollArt: Disaggregated Multi-Task Agentic RL Training at Scale*. arXiv preprint. arXiv: 2512.22560. URL: <https://arxiv.org/abs/2512.22560>.
- Gao, S., W. Zeng, Z. Yu, J. Wangni, et al. (2026). *SWE-MeM: Learning Adaptive Memory Management for Long-Horizon Coding Agents*. arXiv. arXiv: 2606.28434. URL: <https://arxiv.org/abs/2606.28434>.
- Ge et al. (2025). *SAMULE: Self-Learning Agents Enhanced by Multi-Level Reflection*. arXiv preprint; EMNLP 2025. arXiv: 2509.20562. URL: <https://arxiv.org/abs/2509.20562>.
- Ginart, A. et al. (2024). *Asynchronous Tool Usage for Real-Time Agents*. arXiv preprint. arXiv: 2410.21620. URL: <https://arxiv.org/abs/2410.21620>.
- Gonzalez-Pumariiega et al. (2026). *On the Reliability of Computer Use Agents*. arXiv preprint. arXiv: 2604.17849. URL: <https://arxiv.org/abs/2604.17849>.
- Google (2025). *Announcing the Agent2Agent Protocol (A2A)*. URL: <https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/> (visited on 07/08/2026).
- Gorinova et al. (2026). *Position: Coding Benchmarks Are Misaligned with Agentic Software Engineering*. arXiv preprint. arXiv: 2606.17799. URL: <https://arxiv.org/abs/2606.17799>.
- Gou, Z., Z. Shao, Y. Gong, Y. Shen, Y. Yang, N. Duan, and W. Chen (2023). *CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing*. ICLR 2024. arXiv: 2305.11738. URL: <https://arxiv.org/abs/2305.11738>.
- Greshake, K., S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz (2023). *Not what you’ve signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection*. ACM Workshop on AI and Security (AISec) / arXiv. arXiv: 2302.12173. URL: <https://arxiv.org/abs/2302.12173>.
- Gu et al. (2026a). *From Model Scaling to System Scaling: Scaling the Harness in Agentic AI*. arXiv preprint. arXiv: 2605.26112. URL: <https://arxiv.org/abs/2605.26112>.
- (2026b). *PEEK: Context Map as an Orientation Cache for Long-Context LLM Agents*. arXiv preprint. arXiv: 2605.19932. URL: <https://arxiv.org/abs/2605.19932>.
- Guo et al. (2025). *A Measurement Study of Model Context Protocol Ecosystem*. arXiv preprint. arXiv: 2509.25292. URL: <https://arxiv.org/abs/2509.25292>.
- (2026). *From Question Answering to Task Completion: A Survey on Agent System and Harness Design*. arXiv preprint. arXiv: 2606.20683. URL: <https://arxiv.org/abs/2606.20683>.

- Guo, F. (2026). *Why Git Is the Memory Solution for the Agentic Development Lifecycle*. arXiv. arXiv: 2607.14390. URL: <https://arxiv.org/abs/2607.14390>.
- Gupta et al. (2026a). *ReliabilityBench: Evaluating LLM Agent Reliability Under Production-Like Stress Conditions*. arXiv preprint. arXiv: 2601.06112. URL: <https://arxiv.org/abs/2601.06112>.
- Gupta, G., V. Chaturvedi, J. Huan, and A. Deoras (2026b). *Dissecting model behavior through agent trajectories*. arXiv. arXiv: 2606.17454. URL: <https://arxiv.org/abs/2606.17454>.
- Hao, S., Y. Gu, H. Ma, J. J. Hong, Z. Wang, D. Z. Wang, and Z. Hu (2023). *Reasoning with Language Model is Planning with World Model*. EMNLP 2023. arXiv: 2305.14992. URL: <https://arxiv.org/abs/2305.14992>.
- Hao, Y., Z. Jin, H. Liao, K. Liu, et al. (2026). *Why Multi-Step Tool-Use Reinforcement Learning Collapses and How Supervisory Signals Fix It*. arXiv. arXiv: 2606.26027. URL: <https://arxiv.org/abs/2606.26027>.
- Hasan et al. (2025). *Model Context Protocol (MCP) at First Glance: Studying the Security and Maintainability of MCP Servers*. arXiv preprint. arXiv: 2506.13538. URL: <https://arxiv.org/abs/2506.13538>.
- (2026). *Model Context Protocol (MCP) Tool Descriptions Are Smelly! Towards Improving AI Agent Efficiency with Augmented MCP Tool Descriptions*. arXiv preprint. arXiv: 2602.14878. URL: <https://arxiv.org/abs/2602.14878>.
- He et al. (2026a). *SIRI: Self-Internalizing Reinforcement Learning with Intrinsic Skills for LLM Agent Training*. arXiv preprint. arXiv: 2606.02355. URL: <https://arxiv.org/abs/2606.02355>.
- He, B., Y. Chen, X. Zhang, and X. Liu (2026b). *Branching Policy Optimization: Sandbox-Native Language Agent Reinforcement Learning*. arXiv. arXiv: 2607.14171. URL: <https://arxiv.org/abs/2607.14171>.
- He, L., Y. Wang, J. Zhang, and N. Asokan (2026c). *Defending against Adaptive Prompt Injection Attacks via Reasoning-enabled Task Alignment*. arXiv. arXiv: 2606.15441. URL: <https://arxiv.org/abs/2606.15441>.
- Hines, K., G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman (2024). *Defending Against Indirect Prompt Injection Attacks With Spotlighting*. arXiv: 2403.14720. URL: <https://arxiv.org/abs/2403.14720>.
- Hooper et al. (2026). *Speculative Interaction Agents: Building Real-Time Agents with Asynchronous I/O and Speculative Tool Calling*. arXiv preprint. arXiv: 2605.13360. URL: <https://arxiv.org/abs/2605.13360>.
- Horthy, D. (2025). *12-Factor Agents: Principles for Building Reliable LLM Applications*. HumanLayer GitHub repository. URL: <https://github.com/humanlayer/12-factor-agents> (visited on 07/09/2026).
- Hou et al. (2026a). *When Agents Do Not Stop: Uncovering Infinite Agentic Loops in LLM Agents*. arXiv preprint. arXiv: 2607.01641. URL: <https://arxiv.org/abs/2607.01641>.
- Hou, X., Y. Zhao, S. Wang, and H. Wang (2025). *Model Context Protocol (MCP): Landscape, Security Threats, and Future Research Directions*. arXiv: 2503.23278. URL: <https://arxiv.org/abs/2503.23278>.
- Hou, Z., Y. Li, J. Tang, and Y. Dong (2026b). *Single-Rollout Asynchronous Optimization for Agentic Reinforcement Learning*. arXiv. arXiv: 2607.07508. URL: <https://arxiv.org/abs/2607.07508>.
- Hsu, D. C. and L. Lu (2026). *Scoped Verification for Reliable Long-Horizon Agentic Context Evolution under Distribution Shift*. arXiv. arXiv: 2607.09175. URL: <https://arxiv.org/abs/2607.09175>.
- Hu et al. (2025). *AgentSentinel: An End-to-End and Real-Time Security Defense Framework for Computer-Use Agents*. arXiv preprint. arXiv: 2509.07764. URL: <https://arxiv.org/abs/2509.07764>.
- Huang et al. (2026a). *From Raw Experience to Skill Consumption: A Systematic Study of Model-Generated Agent Skills*. arXiv preprint. arXiv: 2605.23899. URL: <https://arxiv.org/abs/2605.23899>.
- (2026b). *Give Them an Inch and They Will Take a Mile: Understanding and Measuring Caller Identity Confusion in MCP-Based AI Systems*. arXiv preprint. arXiv: 2603.07473. URL: <https://arxiv.org/abs/2603.07473>.

- Huang, J., X. Chen, S. Mishra, H. S. Zheng, A. W. Yu, X. Song, and D. Zhou (2023). *Large Language Models Cannot Self-Correct Reasoning Yet*. ICLR 2024. arXiv: 2310.01798. URL: <https://arxiv.org/abs/2310.01798>.
- Huang, W., C. Lee, L. Tng, and S. Ge (2026c). *DeepSWE: Measuring Frontier Coding Agents on Original, Long-Horizon Engineering Tasks*. arXiv. arXiv: 2607.07946. URL: <https://arxiv.org/abs/2607.07946>.
- Huang, X., W. Liu, X. Chen, X. Wang, H. Wang, D. Lian, Y. Wang, R. Tang, and E. Chen (2024). *Understanding the Planning of LLM Agents: A Survey*. arXiv. arXiv: 2402.02716. URL: <https://arxiv.org/abs/2402.02716>.
- Huang, Y., W. Wang, H. Bao, Y. Ma, et al. (2026d). *MemoHarness: Agent Harnesses That Learn from Experience*. arXiv. arXiv: 2607.14159. URL: <https://arxiv.org/abs/2607.14159>.
- Hugging Face (2025). *smolagents*. GitHub repository. Open-source agent artifact. URL: <https://github.com/huggingface/smolagents> (visited on 07/09/2026).
- Jang et al. (2026). *Odysseys: Benchmarking Web Agents on Realistic Long Horizon Tasks*. arXiv preprint. arXiv: 2604.24964. URL: <https://arxiv.org/abs/2604.24964>.
- Jiang et al. (2024). *SELF-[IN]CORRECT: LLMs Struggle with Discriminating Self-Generated Responses*. arXiv preprint. arXiv: 2404.04298. URL: <https://arxiv.org/abs/2404.04298>.
- Jiang, H. et al. (2023). *LLMLingua: Compressing Prompts for Accelerated Inference of Large Language Models*. arXiv preprint; EMNLP 2023. arXiv: 2310.05736. URL: <https://arxiv.org/abs/2310.05736>.
- Jimenez, C. E., J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. Narasimhan (2023). *SWE-bench: Can Language Models Resolve Real-World GitHub Issues?* ICLR 2024. arXiv: 2310.06770. URL: <https://arxiv.org/abs/2310.06770>.
- Jin, B., H. Zeng, Z. Yue, et al. (2025). *Search-R1: Training LLMs to Reason and Leverage Search Engines with Reinforcement Learning*. arXiv: 2503.09516. URL: <https://arxiv.org/abs/2503.09516>.
- Jung et al. (2026). *Adaptive Latent Agentic Reasoning*. arXiv preprint. arXiv: 2606.02871. URL: <https://arxiv.org/abs/2606.02871>.
- Kamoi, R., Y. Zhang, N. Zhang, J. Han, and R. Zhang (2024). *When Can LLMs Actually Correct Their Own Mistakes? A Critical Survey of Self-Correction of LLMs*. TACL 2024. arXiv: 2406.01297. URL: <https://arxiv.org/abs/2406.01297>.
- Kapoor, S., B. Stroebel, P. Kirgis, et al. (2025). *Holistic Agent Leaderboard: The Missing Infrastructure for AI Agent Evaluation*. arXiv: 2510.11977. URL: <https://arxiv.org/abs/2510.11977>.
- Kapoor, S., B. Stroebel, Z. S. Siegel, N. Nadgir, and A. Narayanan (2024). *AI Agents That Matter*. arXiv. arXiv: 2407.01502. URL: <https://arxiv.org/abs/2407.01502>.
- Khanal et al. (2026). *Beyond pass@1: A Reliability Science Framework for Long-Horizon LLM Agents*. arXiv preprint. arXiv: 2603.29231. URL: <https://arxiv.org/abs/2603.29231>.
- Kim et al. (2025). *ReflAct: World-Grounded Decision Making in LLM Agents via Goal-State Reflection*. arXiv preprint. arXiv: 2505.15182. URL: <https://arxiv.org/abs/2505.15182>.
- (2026). *The Interplay of Harness Design and Post-Training in LLM Agents*. arXiv preprint. arXiv: 2606.25447. URL: <https://arxiv.org/abs/2606.25447>.
- Kim, S. et al. (2023). *An LLM Compiler for Parallel Function Calling*. arXiv preprint; ICML 2024 (LLMCompiler). arXiv: 2312.04511. URL: <https://arxiv.org/abs/2312.04511>.
- Kimi Team (2025). *Kimi K2: Open Agentic Intelligence*. arXiv: 2507.20534. URL: <https://arxiv.org/abs/2507.20534>.
- Koh, J. Y. et al. (2024). *Tree Search for Language Model Agents*. arXiv preprint. arXiv: 2407.01476. URL: <https://arxiv.org/abs/2407.01476>.

- Kong et al. (2026). *WebTestBench: Evaluating Computer-Use Agents towards End-to-End Automated Web Testing*. arXiv preprint. arXiv: 2603.25226. URL: <https://arxiv.org/abs/2603.25226>.
- Kwa, T., B. West, J. Becker, A. Deng, K. Garcia, M. Hasin, S. Jawhar, M. Kinniment, N. Rush, S. Von Arx, R. Bloom, T. Broadley, H. Du, B. Goodrich, N. Jurkovic, L. H. Miles, S. Nix, T. Lin, N. Parikh, D. Rein, L. J. K. Sato, H. Wijk, D. M. Ziegler, E. Barnes, and L. Chan (2025). *Measuring AI Ability to Complete Long Software Tasks*. arXiv preprint (METR). arXiv: 2503.14499. URL: <https://arxiv.org/abs/2503.14499>.
- Lam et al. (2026). *SWE-Chain: Benchmarking Coding Agents on Chained Release-Level Package Upgrades*. arXiv preprint. arXiv: 2605.14415. URL: <https://arxiv.org/abs/2605.14415>.
- LangChain (2025). *LangGraph*. GitHub repository. Open-source agent artifact. URL: <https://github.com/langchain-ai/langgraph> (visited on 07/09/2026).
- LangGenius (2025). *Dify*. GitHub repository. Open-source agent artifact. URL: <https://github.com/langgenius/dify> (visited on 07/10/2026).
- Lee, J., S. Hong, S. Lee, J. Seo, J. Son, S. Eo, C. Park, H. Park, H. Moon, and H. Lim (2026a). *DART: Draft-Agreement Routing for Training-Free Adaptive Thinking Budgets in Hybrid Reasoning Models*. arXiv: 2606.23181. URL: <https://arxiv.org/abs/2606.23181>.
- Lee, J., C. Park, and H. Lim (2026b). *To Isolate or to Score? Model-Adaptive Assessment for Cost-Efficient Multi-Agent RAG*. arXiv: 2606.25191. URL: <https://arxiv.org/abs/2606.25191>.
- Li et al. (2025a). *A Vision for Access Control in LLM-based Agent Systems*. arXiv preprint. arXiv: 2510.11108. URL: <https://arxiv.org/abs/2510.11108>.
- (2026a). *Benchmark Test-Time Scaling of General LLM Agents*. arXiv preprint. arXiv: 2602.18998. URL: <https://arxiv.org/abs/2602.18998>.
 - (2026b). *CompactionRL: Reinforcement Learning with Context Compaction for Long-Horizon Agents*. arXiv preprint. arXiv: 2607.05378. URL: <https://arxiv.org/abs/2607.05378>.
 - (2026c). *Cyclical Entropy Eruption: Entropy Dynamics in Agent Reinforcement Learning*. arXiv preprint. arXiv: 2605.27954. URL: <https://arxiv.org/abs/2605.27954>.
 - (2026d). *LiteResearcher: A Scalable Agentic RL Training Framework for Deep Research Agent*. arXiv preprint. arXiv: 2604.17931. URL: <https://arxiv.org/abs/2604.17931>.
 - (2026e). *LiTS: A Modular Framework for LLM Tree Search*. arXiv preprint. arXiv: 2603.00631. URL: <https://arxiv.org/abs/2603.00631>.
 - (2026f). *OTora: A Unified Red Teaming Framework for Reasoning-Level Denial-of-Service in LLM Agents*. arXiv preprint. arXiv: 2605.08876. URL: <https://arxiv.org/abs/2605.08876>.
 - (2026g). *Self-Compacting Language Model Agents*. arXiv preprint. arXiv: 2606.23525. URL: <https://arxiv.org/abs/2606.23525>.
 - (2026h). *Taxonomy and Consistency Analysis of Safety Benchmarks for AI Agents*. arXiv preprint. arXiv: 2605.16282. URL: <https://arxiv.org/abs/2605.16282>.
 - (2026i). *Towards Secure Agent Skills: Architecture, Threat Taxonomy, and Security Analysis*. arXiv preprint. arXiv: 2604.02837. URL: <https://arxiv.org/abs/2604.02837>.
 - (2026j). *WeaveBench: A Long-Horizon, Real-World Benchmark for Computer-Use Agents with Hybrid Interfaces*. arXiv preprint. arXiv: 2606.09426. URL: <https://arxiv.org/abs/2606.09426>.
- Li, W. et al. (2025b). *Chain-of-Agents: End-to-End Agent Foundation Models via Multi-Agent Distillation and Agentic RL*. arXiv preprint (AFM). arXiv: 2508.13167. URL: <https://arxiv.org/abs/2508.13167>.
- Li, X. et al. (2025c). *ToRL: Scaling Tool-Integrated RL*. arXiv preprint. arXiv: 2503.23383. URL: <https://arxiv.org/abs/2503.23383>.

- Liang, S. et al. (2025). *The SWE-Bench Illusion: When State-of-the-Art LLMs Remember Instead of Reason*. arXiv preprint. arXiv: 2506.12286. URL: <https://arxiv.org/abs/2506.12286>.
- Lightman, H., V. Kosaraju, Y. Burda, H. Edwards, B. Baker, T. Lee, J. Leike, J. Schulman, I. Sutskever, and K. Cobbe (2023). *Let’s Verify Step by Step*. ICLR 2024. arXiv: 2305.20050. URL: <https://arxiv.org/abs/2305.20050>.
- Lin et al. (2026a). *MUSE-Autoskill: Self-Evolving Agents via Skill Creation, Memory, Management, and Evaluation*. arXiv preprint. arXiv: 2605.27366. URL: <https://arxiv.org/abs/2605.27366>.
- (2026b). *Skill-MAS: Evolving Meta-Skill for Automatic Multi-Agent Systems*. arXiv preprint. arXiv: 2606.18837. URL: <https://arxiv.org/abs/2606.18837>.
- Lindenbauer et al. (2025). *The Complexity Trap: Simple Observation Masking Is as Efficient as LLM Summarization for Agent Context Management*. arXiv preprint. arXiv: 2508.21433. URL: <https://arxiv.org/abs/2508.21433>.
- Liu et al. (2025). *Budget-Aware Tool-Use Enables Effective Agent Scaling*. arXiv preprint. arXiv: 2511.17006. URL: <https://arxiv.org/abs/2511.17006>.
- (2026a). *Agent Skills in the Wild: An Empirical Study of Security Vulnerabilities at Scale*. arXiv preprint. arXiv: 2601.10338. URL: <https://arxiv.org/abs/2601.10338>.
- (2026b). *ASymPO: Asymmetric-Scale Policy Optimization for Asynchronous LLM Post-Training Without Behavior Information*. arXiv preprint. arXiv: 2606.03070. URL: <https://arxiv.org/abs/2606.03070>.
- (2026c). *MAP: A Map-then-Act Paradigm for Long-Horizon Interactive Agent Reasoning*. arXiv preprint. arXiv: 2605.13037. URL: <https://arxiv.org/abs/2605.13037>.
- (2026d). *MemGUI-Agent: An End-to-End Long-Horizon Mobile GUI Agent with Proactive Context Management*. arXiv preprint. arXiv: 2606.19926. URL: <https://arxiv.org/abs/2606.19926>.
- (2026e). *SkillsVote: Lifecycle Governance of Agent Skills from Collection, Recommendation to Evolution*. arXiv preprint. arXiv: 2605.18401. URL: <https://arxiv.org/abs/2605.18401>.
- Liu, G. and S. Gabriel (2026). *PM-Bench: Evaluating Prospective Memory in LLM Agents*. arXiv. arXiv: 2607.12385. URL: <https://arxiv.org/abs/2607.12385>.
- Liu, H., K. Hu, J. Liao, Q. Wang, et al. (2026f). *MCPEvol-Bench: Benchmarking LLM Agent Performance Across Dynamic Evolutions of MCP Servers*. arXiv. arXiv: 2607.14642. URL: <https://arxiv.org/abs/2607.14642>.
- Liu, N. F. et al. (2023). *Lost in the Middle: How Language Models Use Long Contexts*. arXiv preprint; TACL 2024. arXiv: 2307.03172. URL: <https://arxiv.org/abs/2307.03172>.
- Liu, T. (2026). *From Tool Connection to Execution Control: Benchmarking Security Invariants in MCP-Style Agent Runtimes*. arXiv. arXiv: 2606.29073. URL: <https://arxiv.org/abs/2606.29073>.
- Liu, W., J. Bi, W. Zhou, J. Feng, et al. (2026g). *ToolAnchor: Anchoring Counterfactual Context to Boost Agentic Tool-use Capability*. arXiv. arXiv: 2607.14145. URL: <https://arxiv.org/abs/2607.14145>.
- Liu, X., H. Yu, H. Zhang, Y. Xu, et al. (2024). *AgentBench: Evaluating LLMs as Agents*. ICLR 2024. arXiv: 2308.03688. URL: <https://arxiv.org/abs/2308.03688>.
- Lodha et al. (2026). *Less Context, Better Agents: Efficient Context Engineering for Long-Horizon Tool-Using LLM Agents*. arXiv preprint. arXiv: 2606.10209. URL: <https://arxiv.org/abs/2606.10209>.
- Lumer et al. (2025). *Tool-to-Agent Retrieval: Bridging Tools and Agents for Scalable LLM Multi-Agent Systems*. arXiv preprint. arXiv: 2511.01854. URL: <https://arxiv.org/abs/2511.01854>.
- Ma et al. (2026a). *AutoDojo: Adaptive Black-Box Attacks Reveal the Limits of IPI Defenses and Task-Specification Effects in LLM Agents*. arXiv preprint. arXiv: 2606.15057. URL: <https://arxiv.org/abs/2606.15057>.

- Ma et al. (2026b). *Timely Machine: Awareness of Time Makes Test-Time Scaling Agentic*. arXiv preprint. arXiv: 2601.16486. URL: <https://arxiv.org/abs/2601.16486>.
- Ma, C., J. Zhang, Z. Zhu, C. Yang, Y. Yang, Y. Jin, Z. Lan, L. Kong, and J. He (2024). *AgentBoard: An Analytical Evaluation Board of Multi-turn LLM Agents*. NeurIPS 2024. arXiv: 2401.13178. URL: <https://arxiv.org/abs/2401.13178>.
- Macedo, S. O. de (2026). *What makes a harness a harness: necessary and sufficient conditions for an agent harness*. arXiv. arXiv: 2606.10106. URL: <https://arxiv.org/abs/2606.10106>.
- Madaan, A., N. Tandon, P. Gupta, S. Hallinan, L. Gao, et al. (2023). *SELF-REFINE: Iterative Refinement with Self-Feedback*. NeurIPS 2023. arXiv: 2303.17651. URL: <https://arxiv.org/abs/2303.17651>.
- Madiraju, M. B. and M. S. P. Madiraju (2026). *RigorBench: Benchmarking Engineering Process Discipline in Autonomous AI Coding Agents*. arXiv. arXiv: 2606.22678. URL: <https://arxiv.org/abs/2606.22678>.
- Maldaner et al. (2026). *SentinelBench: A Benchmark for Long-Running Monitoring Agents*. arXiv preprint. arXiv: 2606.05342. URL: <https://arxiv.org/abs/2606.05342>.
- Marchand et al. (2026). *Quantifying Frontier LLM Capabilities for Container Sandbox Escape*. arXiv preprint. arXiv: 2603.02277. URL: <https://arxiv.org/abs/2603.02277>.
- Masterman, T., S. Besen, M. Sawtell, and A. Chao (2024). *The Landscape of Emerging AI Agent Architectures for Reasoning, Planning, and Tool Calling: A Survey*. arXiv. arXiv: 2404.11584. URL: <https://arxiv.org/abs/2404.11584>.
- Mazaheri et al. (2026). *AgentAtlas: Beyond Outcome Leaderboards for LLM Agents*. arXiv preprint. arXiv: 2605.20530. URL: <https://arxiv.org/abs/2605.20530>.
- Mehta et al. (2026). *Security Risks of AI Agents Hiring Humans: An Empirical Marketplace Study*. arXiv preprint. arXiv: 2602.19514. URL: <https://arxiv.org/abs/2602.19514>.
- Meng et al. (2025). *ceLLMate: Sandboxing Browser AI Agents*. arXiv preprint. arXiv: 2512.12594. URL: <https://arxiv.org/abs/2512.12594>.
- (2026). *SkillRAE: Agent Skill-Based Context Compilation for Retrieval-Augmented Execution*. arXiv preprint. arXiv: 2605.10114. URL: <https://arxiv.org/abs/2605.10114>.
- Merrill, M. A., A. G. Shaw, N. Carlini, et al. (2026). *Terminal-Bench: Benchmarking Agents on Hard, Realistic Tasks in Command Line Interfaces*. arXiv: 2601.11868. URL: <https://arxiv.org/abs/2601.11868>.
- Mialon, G., C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom (2023). *GALA: a benchmark for General AI Assistants*. ICLR 2024. arXiv: 2311.12983. URL: <https://arxiv.org/abs/2311.12983>.
- Microsoft (2025a). *AutoGen*. GitHub repository. Open-source agent artifact. URL: <https://github.com/microsoft/autogen> (visited on 07/09/2026).
- (2025b). *microsoft/skills*. GitHub repository. Open-source agent artifact. URL: <https://github.com/microsoft/skills> (visited on 07/09/2026).
- Min et al. (2026). *Stop When Reasoning Converges: Semantic-Preserving Early Exit for Reasoning Models*. arXiv preprint. arXiv: 2605.17672. URL: <https://arxiv.org/abs/2605.17672>.
- Model Context Protocol (2025). *Model Context Protocol Servers*. GitHub repository. Open-source agent artifact. URL: <https://github.com/modelcontextprotocol/servers> (visited on 07/10/2026).
- Moghe, K. and P. Chin (2026). *Cost-Effective Agent Harnesses for Abstract Reasoning and Generalization on ARC-AGI-1*. arXiv. arXiv: 2607.06764. URL: <https://arxiv.org/abs/2607.06764>.
- Muennighoff, N. et al. (2025). *s1: Simple Test-Time Scaling*. arXiv preprint. arXiv: 2501.19393. URL: <https://arxiv.org/abs/2501.19393>.

- Naik et al. (2025). *AgentMisalignment: Measuring the Propensity for Misaligned Behaviour in LLM-Based Agents*. arXiv preprint. arXiv: 2506.04018. URL: <https://arxiv.org/abs/2506.04018>.
- Nakano, R., J. Hilton, S. Balaji, J. Wu, L. Ouyang, et al. (2021). *WebGPT: Browser-assisted question-answering with human feedback*. arXiv (OpenAI). arXiv: 2112.09332. URL: <https://arxiv.org/abs/2112.09332>.
- Niu, P., W. Qu, S. Gu, T. Shi, et al. (2026). *Understanding and Evaluating Claw-like Agent Security Through a Computer-Systems Lens*. arXiv. arXiv: 2606.30755. URL: <https://arxiv.org/abs/2606.30755>.
- obra (2025). *Superpowers*. GitHub repository. Open-source agent artifact. URL: <https://github.com/obra/superpowers> (visited on 07/10/2026).
- Ogundoyin et al. (2025). *A Large-Scale Empirical Analysis of Custom GPTs’ Vulnerabilities in the OpenAI Ecosystem*. arXiv preprint. arXiv: 2505.08148. URL: <https://arxiv.org/abs/2505.08148>.
- OpenAI (2025). *Codex CLI*. GitHub repository. Open-source agent artifact. URL: <https://github.com/openai/codex> (visited on 07/09/2026).
- OpenHands (All-Hands-AI) (2025). *OpenHands*. GitHub repository. Open-source agent artifact. URL: <https://github.com/OpenHands/OpenHands> (visited on 07/09/2026).
- Ouyang et al. (2026). *SkillOS: Learning Skill Curation for Self-Evolving Agents*. arXiv preprint. arXiv: 2605.06614. URL: <https://arxiv.org/abs/2605.06614>.
- Ouyang, S., J. Yan, I.-H. Hsu, et al. (2025). *ReasoningBank: Scaling Agent Self-Evolving with Reasoning Memory*. arXiv: 2509.25140. URL: <https://arxiv.org/abs/2509.25140>.
- Packer, C., S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez (2023). *MemGPT: Towards LLMs as Operating Systems*. COLM 2024. arXiv: 2310.08560. URL: <https://arxiv.org/abs/2310.08560>.
- Paglieri, D. et al. (2025). *Learning When to Plan: Efficiently Allocating Test-Time Compute for LLM Agents*. arXiv preprint. arXiv: 2509.03581. URL: <https://arxiv.org/abs/2509.03581>.
- Pan, J. et al. (2024). *Training Software Engineering Agents and Verifiers with SWE-Gym*. arXiv preprint; ICML 2025. arXiv: 2412.21139. URL: <https://arxiv.org/abs/2412.21139>.
- Patil, S. G., T. Zhang, X. Wang, and J. E. Gonzalez (2024). *Gorilla: Large Language Model Connected with Massive APIs*. NeurIPS 2024. arXiv: 2305.15334. URL: <https://arxiv.org/abs/2305.15334>.
- Peng, Z., X. Yin, C. Ying, Z. Cui, et al. (2026). *EvoClawBench: Can Agents Learn Reusable Skills from Their Own Runs?* arXiv. arXiv: 2607.09711. URL: <https://arxiv.org/abs/2607.09711>.
- Piet et al. (2026). *Web Agents Should Adopt the Plan-Then-Execute Paradigm*. arXiv preprint. arXiv: 2605.14290. URL: <https://arxiv.org/abs/2605.14290>.
- Plaat, A., M. van Duijn, N. van Stein, M. Preuss, P. van der Putten, and K. J. Batenburg (2025). *Agentic Large Language Models, a Survey*. arXiv. arXiv: 2503.23037. URL: <https://arxiv.org/abs/2503.23037>.
- Podivilov, A., V. Lomshakov, S. Savin, M. Startsev, et al. (2026). *AgentLens: Production-Assessed Trajectory Reviews for Coding Agent Evaluation*. arXiv. arXiv: 2607.06624. URL: <https://arxiv.org/abs/2607.06624>.
- Prasad, A., A. Koller, M. Hartmann, P. Clark, A. Sabharwal, M. Bansal, and T. Khot (2023). *ADaPT: As-Needed Decomposition and Planning with Language Models*. NAACL 2024 Findings. arXiv: 2311.05772. URL: <https://arxiv.org/abs/2311.05772>.
- Princeton NLP / SWE-agent (2025). *SWE-agent*. GitHub repository. Open-source agent artifact. URL: <https://github.com/SWE-agent/SWE-agent> (visited on 07/09/2026).
- punkpeye (2025). *awesome-mcp-servers*. GitHub repository. Open-source agent artifact. URL: <https://github.com/punkpeye/awesome-mcp-servers> (visited on 07/10/2026).

- Putta, P. et al. (2024). *Agent Q: Advanced Reasoning and Learning for Autonomous AI Agents*. arXiv preprint. arXiv: 2408.07199. URL: <https://arxiv.org/abs/2408.07199>.
- Qiao, S. et al. (2025). *Agentic Knowledgeable Self-Awareness*. arXiv preprint. arXiv: 2504.03553. URL: <https://arxiv.org/abs/2504.03553>.
- Qin, Y. et al. (2025). *UI-TARS: Pioneering Automated GUI Interaction with Native Agents*. arXiv. arXiv: 2501.12326. URL: <https://arxiv.org/abs/2501.12326>.
- Qin, Y., S. Liang, Y. Ye, K. Zhu, L. Yan, et al. (2023). *ToolLLM: Facilitating Large Language Models to Master 16000+ Real-world APIs*. arXiv (ICLR 2024). arXiv: 2307.16789. URL: <https://arxiv.org/abs/2307.16789>.
- Qu et al. (2026). *Supply-Chain Poisoning Attacks Against LLM Coding Agent Skill Ecosystems*. arXiv preprint. arXiv: 2604.03081. URL: <https://arxiv.org/abs/2604.03081>.
- Qu, C., S. Dai, X. Wei, H. Cai, S. Wang, D. Yin, J. Xu, and J.-R. Wen (2024). *Tool Learning with Large Language Models: A Survey*. arXiv; Frontiers of Computer Science 2024 (DOI: 10.1007/s11704-024-40678-2). arXiv: 2405.17935. URL: <https://arxiv.org/abs/2405.17935>.
- Rashidi, M. (2026a). *Depth-Dependent Indirect Prompt Injection in Tool-Calling ReAct Agents: Injection Depth, Payload Framing, and Turn-Budget Sensitivity*. arXiv. arXiv: 2605.30686. URL: <https://arxiv.org/abs/2605.30686>.
- (2026b). *The Balkanization of Execution-Security Research for AI Coding Agents: Isolation, Access Control, and Time-of-Check-to-Time-of-Use Vulnerabilities*. arXiv. arXiv: 2607.05743. URL: <https://arxiv.org/abs/2607.05743>.
- Renze et al. (2024). *Self-Reflection in LLM Agents: Effects on Problem-Solving Performance*. arXiv preprint. arXiv: 2405.06682. URL: <https://arxiv.org/abs/2405.06682>.
- Ridnik, T. et al. (2024). *Code Generation with AlphaCodium: From Prompt Engineering to Flow Engineering*. arXiv preprint. arXiv: 2401.08500. URL: <https://arxiv.org/abs/2401.08500>.
- Rombaut et al. (2026). *Inside the Scaffold: A Source-Code Taxonomy of Coding Agent Architectures*. arXiv preprint. arXiv: 2604.03515. URL: <https://arxiv.org/abs/2604.03515>.
- Rozanov et al. (2024). *StateAct: Enhancing LLM Base Agents via Self-Prompting and State-Tracking*. arXiv preprint. arXiv: 2410.02810. URL: <https://arxiv.org/abs/2410.02810>.
- Ruan et al. (2026). *AOrchestra: Automating Sub-Agent Creation for Agentic Orchestration*. arXiv preprint. arXiv: 2602.03786. URL: <https://arxiv.org/abs/2602.03786>.
- Ruan, Y., H. Dong, A. Wang, S. Pitis, Y. Zhou, J. Ba, Y. Dubois, C. J. Maddison, and T. Hashimoto (2024). *Identifying the Risks of LM Agents with an LM-Emulated Sandbox*. ICLR 2024 (Spotlight). arXiv: 2309.15817. URL: <https://arxiv.org/abs/2309.15817>.
- Saha et al. (2026). *Under the Hood of SKILL.md: Semantic Supply-chain Attacks on AI Agent Skill Registry*. arXiv preprint. arXiv: 2605.11418. URL: <https://arxiv.org/abs/2605.11418>.
- Santos-Grueiro et al. (2026). *Lingering Authority: Revocable Resource-and-Effect Capabilities for Coding Agents*. arXiv preprint. arXiv: 2606.22504. URL: <https://arxiv.org/abs/2606.22504>.
- Sarch, G., L. Jang, M. J. Tarr, W. W. Cohen, K. Marino, and K. Fragkiadaki (2024). “VLM Agents Generate Their Own Memories: Distilling Experience into Embodied Programs of Thought”. In: *Advances in Neural Information Processing Systems*. NeurIPS 2024. DOI: 10.52202/079017–2418. arXiv: 2406.14596. URL: https://proceedings.neurips.cc/paper_files/paper/2024/hash/8ac50fd0a4eeeb1f077b17bb7c5353c3-Abstract-Conference.html.
- Schick, T., J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, L. Zettlemoyer, N. Cancedda, and T. Scialom (2023). *Toolformer: Language Models Can Teach Themselves to Use Tools*. arXiv preprint; NeurIPS 2023. arXiv: 2302.04761. URL: <https://arxiv.org/abs/2302.04761>.

- Sghaier, O. B., H. Li, B. Adams, and A. E. Hassan (2026). *Don't Blame the Large Language Model: How Scaffolding Evolution Shapes Coding Agent Quality*. arXiv. arXiv: 2607.03691. URL: <https://arxiv.org/abs/2607.03691>.
- Shaposhnikov, M., N. Fortuin, S. Stipcich, M. I. Gorinova, et al. (2026). *A Framework for Evaluating Agentic Skills at Scale*. arXiv. arXiv: 2606.17819. URL: <https://arxiv.org/abs/2606.17819>.
- Shen et al. (2026). *SkillFoundry: Building Self-Evolving Agent Skill Libraries from Heterogeneous Scientific Resources*. arXiv preprint. arXiv: 2604.03964. URL: <https://arxiv.org/abs/2604.03964>.
- Shi, T., J. He, Z. Wang, H. Li, L. Wu, W. Guo, and D. Song (2025). *Progent: Securing AI Agents with Privilege Control*. arXiv preprint. arXiv: 2504.11703. URL: <https://arxiv.org/abs/2504.11703>.
- Shinn, N., F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao (2023). *Reflexion: Language Agents with Verbal Reinforcement Learning*. arXiv preprint; NeurIPS 2023. arXiv: 2303.11366. URL: <https://arxiv.org/abs/2303.11366>.
- Sidik et al. (2026). *Beyond Static Sandboxing: Learned Capability Governance for Autonomous AI Agents*. arXiv preprint. arXiv: 2604.11839. URL: <https://arxiv.org/abs/2604.11839>.
- Significant Gravitass (2023). *AutoGPT*. GitHub repository. Open-source agent artifact. URL: <https://github.com/Significant-Gravitas/AutoGPT> (visited on 07/10/2026).
- Snell, C. et al. (2024). *Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters*. arXiv preprint. arXiv: 2408.03314. URL: <https://arxiv.org/abs/2408.03314>.
- Song et al. (2026a). *Staleness-Learning Rate Scaling Laws for Asynchronous RLHF*. arXiv preprint. arXiv: 2607.01083. URL: <https://arxiv.org/abs/2607.01083>.
- Song, Y., L. Sutawika, J. Liu, L. Tjuaatja, et al. (2026b). *PACE: A Proxy for Agentic Capability Evaluation*. arXiv. arXiv: 2607.02032. URL: <https://arxiv.org/abs/2607.02032>.
- SST (2025). *opencode*. GitHub repository. Open-source agent artifact. URL: <https://github.com/sst/opencode> (visited on 07/10/2026).
- Stanford NLP (2025). *DSPy*. GitHub repository. Open-source agent artifact. URL: <https://github.com/stanfordnlp/dspy> (visited on 07/09/2026).
- Stauffer, L., K. Feng, K. Wei, et al. (2026). *The 2025 AI Agent Index: Documenting Technical and Safety Features of Deployed Agentic AI Systems*. ACM FAccT 2026 (arXiv 2026). arXiv: 2602.17753. URL: <https://arxiv.org/abs/2602.17753>.
- Stein et al. (2026). *How are AI agents used? Evidence from 177,000 MCP tools*. arXiv preprint. arXiv: 2603.23802. URL: <https://arxiv.org/abs/2603.23802>.
- Su, H. et al. (2025). *Learn-by-Interact: A Data-Centric Framework for Self-Adaptive Agents in Realistic Environments*. arXiv preprint. arXiv: 2501.10893. URL: <https://arxiv.org/abs/2501.10893>.
- Sui, Y. et al. (2025). *Stop Overthinking: A Survey on Efficient Reasoning for Large Language Models*. arXiv preprint; TMLR 2025. arXiv: 2503.16419. URL: <https://arxiv.org/abs/2503.16419>.
- Sumers, T. R., S. Yao, K. Narasimhan, and T. L. Griffiths (2023). *Cognitive Architectures for Language Agents*. arXiv; TMLR (camera-ready v3). arXiv: 2309.02427. URL: <https://arxiv.org/abs/2309.02427>.
- Sun, H. et al. (2023). *AdaPlanner: Adaptive Planning from Feedback with Language Models*. arXiv preprint; NeurIPS 2023. arXiv: 2305.16653. URL: <https://arxiv.org/abs/2305.16653>.
- Takerngsaksiri et al. (2024). *Human-In-the-Loop Software Development Agents*. arXiv preprint. arXiv: 2411.12924. URL: <https://arxiv.org/abs/2411.12924>.
- Tan, W. et al. (2024). *Cradle: Empowering Foundation Agents Towards General Computer Control*. arXiv preprint. arXiv: 2403.03186. URL: <https://arxiv.org/abs/2403.03186>.

- Tang et al. (2026). *How Coding Agents Fail Their Users: A Large-Scale Analysis of Developer-Agent Misalignment in 20,574 Real-World Sessions*. arXiv preprint. arXiv: 2605.29442. URL: <https://arxiv.org/abs/2605.29442>.
- Tao, L., B. Peng, W. Yao, T. Ge, et al. (2026). *TRACE: Turn-level Reward Assignment via Credit Estimation for Long-Horizon Agents*. arXiv. arXiv: 2607.13988. URL: <https://arxiv.org/abs/2607.13988>.
- Team, S. (2026). *SingGuard-NSFA: Extensible Guardrails for Agentic AI via Generative Reasoning and Real-Time Classification*. arXiv. arXiv: 2607.13081. URL: <https://arxiv.org/abs/2607.13081>.
- Tie et al. (2025). *Can LLMs Correct Themselves? A Benchmark of Self-Correction in LLMs*. arXiv preprint (Correct-Bench). arXiv: 2510.16062. URL: <https://arxiv.org/abs/2510.16062>.
- Torres, G., S. Shrestha, and S. Misra (2026). *When Agents Remember Too Much: Memory Poisoning Attacks on Large Language Model Agents*. arXiv. arXiv: 2607.06595. URL: <https://arxiv.org/abs/2607.06595>.
- Tran et al. (2026). *Single-Agent LLMs Outperform Multi-Agent Systems on Multi-Hop Reasoning Under Equal Thinking Token Budgets*. arXiv preprint. arXiv: 2604.02460. URL: <https://arxiv.org/abs/2604.02460>.
- Tu et al. (2026). *BenchGuard: Who Guards the Benchmarks? Automated Auditing of LLM Agent Benchmarks*. arXiv preprint. arXiv: 2604.24955. URL: <https://arxiv.org/abs/2604.24955>.
- Tzifas, G. et al. (2024). *Lifelong Robot Library Learning: Bootstrapping Composable and Generalizable Skills for Embodied Control with Language Models*. arXiv preprint. arXiv: 2406.18746. URL: <https://arxiv.org/abs/2406.18746>.
- Verma et al. (2024). *On the Brittle Foundations of ReAct Prompting for Agentic Large Language Models*. arXiv preprint. arXiv: 2405.13966. URL: <https://arxiv.org/abs/2405.13966>.
- VoltAgent (2025). *awesome-claude-code-subagents*. GitHub repository. Open-source agent artifact. URL: <https://github.com/VoltAgent/awesome-claude-code-subagents> (visited on 07/09/2026).
- Wallace, E., K. Xiao, R. Leike, L. Weng, J. Heidecke, and A. Beutel (2024). *The Instruction Hierarchy: Training LLMs to Prioritize Privileged Instructions*. arXiv: 2404.13208. URL: <https://arxiv.org/abs/2404.13208>.
- Wan et al. (2025). *Rethinking the Design of Reinforcement Learning-Based Deep Research Agents*. arXiv preprint. arXiv: 2510.15862. URL: <https://arxiv.org/abs/2510.15862>.
- (2026). *Inference-Time Scaling of Verification: Self-Evolving Deep Research Agents via Test-Time Rubric-Guided Verification*. arXiv preprint. arXiv: 2601.15808. URL: <https://arxiv.org/abs/2601.15808>.
- Wang et al. (2025a). *MCPTox: A Benchmark for Tool Poisoning Attack on Real-World MCP Servers*. arXiv preprint. arXiv: 2508.14925. URL: <https://arxiv.org/abs/2508.14925>.
- (2025b). *ProbGuard: Probabilistic Runtime Monitoring for LLM Agent Safety*. arXiv preprint. arXiv: 2508.00500. URL: <https://arxiv.org/abs/2508.00500>.
- (2025c). *Thoughts Are All Over the Place: On the Underthinking of o1-Like LLMs*. arXiv preprint. arXiv: 2501.18585. URL: <https://arxiv.org/abs/2501.18585>.
- (2026a). *Aligning Agents via Planning: A Benchmark for Trajectory-Level Reward Modeling*. arXiv preprint. arXiv: 2604.08178. URL: <https://arxiv.org/abs/2604.08178>.
- (2026b). *Automated Benchmark Auditing for AI Agents and Large Language Models*. arXiv preprint. arXiv: 2605.26079. URL: <https://arxiv.org/abs/2605.26079>.
- (2026c). *Budgeted Act-or-Defer Multi-Agent LLM Deliberation with Local Reliability Bounds*. arXiv preprint. arXiv: 2606.29654. URL: <https://arxiv.org/abs/2606.29654>.
- (2026d). *HarnessBridge: Learnable Bidirectional Controller for LLM Agent Harness*. arXiv preprint. arXiv: 2606.12882. URL: <https://arxiv.org/abs/2606.12882>.

- Wang et al. (2026e). *Memex(RL): Scaling Long-Horizon LLM Agents via Indexed Experience Memory*. arXiv preprint. arXiv: 2603.04257. URL: <https://arxiv.org/abs/2603.04257>.
- (2026f). *T2PO: Uncertainty-Guided Exploration Control for Stable Multi-Turn Agentic Reinforcement Learning*. arXiv preprint. arXiv: 2605.02178. URL: <https://arxiv.org/abs/2605.02178>.
- (2026g). *The Long-Horizon Task Mirage? Diagnosing Where and Why Agentic Systems Break*. arXiv preprint. arXiv: 2604.11978. URL: <https://arxiv.org/abs/2604.11978>.
- Wang, G., Y. Xie, Y. Jiang, A. Mandlekar, C. Xiao, Y. Zhu, L. Fan, and A. Anandkumar (2023a). *Voyager: An Open-Ended Embodied Agent with Large Language Models*. arXiv preprint; TMLR. arXiv: 2305.16291. URL: <https://arxiv.org/abs/2305.16291>.
- Wang, H., H. Zou, H. Song, et al. (2025d). *UI-TARS-2 Technical Report: Advancing GUI Agent with Multi-Turn Reinforcement Learning*. arXiv: 2509.02544. URL: <https://arxiv.org/abs/2509.02544>.
- Wang, H., C. M. Poskitt, et al. (2025). *AgentSpec: Customizable Runtime Enforcement for Safe and Reliable LLM Agents*. arXiv preprint. arXiv: 2503.18666. URL: <https://arxiv.org/abs/2503.18666>.
- Wang, L., C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin, W. X. Zhao, Z. Wei, and J.-R. Wen (2023b). *A Survey on Large Language Model based Autonomous Agents*. arXiv preprint; also published in *Frontiers of Computer Science* 2024 (DOI: 10.1007/s11704-024-40231-1). arXiv: 2308.11432. URL: <https://arxiv.org/abs/2308.11432>.
- Wang, L., W. Xu, Y. Lan, Z. Hu, Y. Lan, R. K.-W. Lee, and E.-P. Lim (2023c). *Plan-and-Solve Prompting: Improving Zero-Shot Chain-of-Thought Reasoning by Large Language Models*. ACL 2023. arXiv: 2305.04091. URL: <https://arxiv.org/abs/2305.04091>.
- Wang, R., Y. Shi, Z. Li, Z. Li, et al. (2026h). *Harness Handbook: Making Evolving Agent Harnesses Readable, Navigable, and Editable*. arXiv. arXiv: 2607.13285. URL: <https://arxiv.org/abs/2607.13285>.
- Wang, W., Y. Fang, F. Liu, Z. Liang, et al. (2026i). *Experience Memory Graph: One-Shot Error Correction for Agents*. arXiv. arXiv: 2607.13884. URL: <https://arxiv.org/abs/2607.13884>.
- Wang, X., C. Han, K. Yu, X. Xu, et al. (2026j). *Multi-Paradigm Agent Interaction in Practice: A Systematic Analysis of Generator-Evaluator, ReAct Loop, and Adversarial Evaluation in the buddyMe Framework*. arXiv. arXiv: 2605.16821. URL: <https://arxiv.org/abs/2605.16821>.
- Wang, X. et al. (2024a). *OpenHands: An Open Platform for AI Software Developers as Generalist Agents*. ICLR 2025 / arXiv. arXiv: 2407.16741. URL: <https://arxiv.org/abs/2407.16741>.
- Wang, X., Y. Chen, L. Yuan, Y. Zhang, Y. Li, H. Peng, and H. Ji (2024b). *Executable Code Actions Elicit Better LLM Agents*. ICML 2024. arXiv: 2402.01030. URL: <https://arxiv.org/abs/2402.01030>.
- Wang, X., C. Yang, H. Zhao, Z. Lin, et al. (2026k). *The Agent Use of Agent Beings: Agent Cybernetics Is the Missing Science of Foundation Agents*. arXiv. arXiv: 2605.10754. URL: <https://arxiv.org/abs/2605.10754>.
- Wang, Y., P. Yao, T. Sun, C. Hu, et al. (2026l). *SkillCorpus: Consolidating and Evaluating the Open Skill Ecosystem for Real-World LLM Agents*. arXiv. arXiv: 2607.15557. URL: <https://arxiv.org/abs/2607.15557>.
- Wang, Z., M. Yan, J. Bi, S. Yan, et al. (2026m). *MetaSkill-Evolve: Recursive Self-Improvement of LLM Agents via Two-Timescale Meta-Skill Evolution*. arXiv. arXiv: 2607.05297. URL: <https://arxiv.org/abs/2607.05297>.
- Wang, Z. et al. (2025e). *RAGEN: Understanding Self-Evolution in LLM Agents via Multi-Turn Reinforcement Learning*. arXiv preprint (StarPO/StarPO-S). arXiv: 2504.20073. URL: <https://arxiv.org/abs/2504.20073>.
- Wang, Z. et al. (2023d). *JARVIS-1: Open-World Multi-Task Agents with Memory-Augmented Multimodal Language Models*. arXiv preprint. arXiv: 2311.05997. URL: <https://arxiv.org/abs/2311.05997>.

- Wang, Z., S. Cai, G. Chen, A. Liu, X. Ma, and Y. Liang (2023e). *Describe, Explain, Plan and Select: Interactive Planning with Large Language Models Enables Open-World Multi-Task Agents*. NeurIPS 2023. arXiv: 2302.01560. URL: <https://arxiv.org/abs/2302.01560>.
- Wang, Z. Z. et al. (2025f). *Inducing Programmatic Skills for Agentic Tasks*. arXiv preprint. arXiv: 2504.06821. URL: <https://arxiv.org/abs/2504.06821>.
- Wang, Z. Z., J. Mao, D. Fried, and G. Neubig (2024c). *Agent Workflow Memory*. arXiv preprint (AWM). arXiv: 2409.07429. URL: <https://arxiv.org/abs/2409.07429>.
- Wei et al. (2026a). *From Agent Loops to Structured Graphs: A Scheduler-Theoretic Framework for LLM Agent Execution*. arXiv preprint. arXiv: 2604.11378. URL: <https://arxiv.org/abs/2604.11378>.
- (2026b). *OpenComputer: Verifiable Software Worlds for Computer-Use Agents*. arXiv preprint. arXiv: 2605.19769. URL: <https://arxiv.org/abs/2605.19769>.
- Wei, Y. et al. (2025). *SWE-RL: Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution*. arXiv preprint. arXiv: 2502.18449. URL: <https://arxiv.org/abs/2502.18449>.
- wshobson (2025). *wshobson/agents*. GitHub repository. Open-source agent artifact. URL: <https://github.com/wshobson/agents> (visited on 07/09/2026).
- Wu et al. (2025). *GAP: Graph-Based Agent Planning with Parallel Tool Use and Reinforcement Learning*. arXiv preprint. arXiv: 2510.25320. URL: <https://arxiv.org/abs/2510.25320>.
- (2026a). *Crab: A Semantics-Aware Checkpoint/Restore Runtime for Agent Sandboxes*. arXiv preprint. arXiv: 2604.28138. URL: <https://arxiv.org/abs/2604.28138>.
- Wu, Q., G. Bansal, J. Zhang, Y. Wu, B. Li, E. Zhu, L. Jiang, X. Zhang, S. Zhang, J. Liu, A. H. Awadallah, R. W. White, D. Burger, and C. Wang (2023). *AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation*. arXiv. arXiv: 2308.08155. URL: <https://arxiv.org/abs/2308.08155>.
- Wu, Y., Z. Zhao, S. Li, H. H. Lee, et al. (2026b). *SWE-Together: Evaluating Coding Agents in Interactive User Sessions*. arXiv. arXiv: 2606.29957. URL: <https://arxiv.org/abs/2606.29957>.
- Wu, Y. et al. (2024). *IsolateGPT: An Execution Isolation Architecture for LLM-Based Agentic Systems*. arXiv preprint; NDSS 2025 (SecGPT). arXiv: 2403.04960. URL: <https://arxiv.org/abs/2403.04960>.
- x1xhlol (2025). *System Prompts and Models of AI Tools*. GitHub repository. Open-source agent artifact. URL: <https://github.com/x1xhlol/system-prompts-and-models-of-ai-tools> (visited on 07/10/2026).
- Xi, Z., W. Chen, X. Guo, W. He, Y. Ding, et al. (2023). *The Rise and Potential of Large Language Model Based Agents: A Survey*. arXiv. arXiv: 2309.07864. URL: <https://arxiv.org/abs/2309.07864>.
- Xia et al. (2026). *Harnessing Agent Skills: Architectural Patterns and a Reference Architecture for Skill-Mediated LLM Agents*. arXiv preprint. arXiv: 2606.20631. URL: <https://arxiv.org/abs/2606.20631>.
- Xia, C. S., Y. Deng, S. Dunn, and L. Zhang (2024). *Agentless: Demystifying LLM-based Software Engineering Agents*. FSE 2025 / arXiv. arXiv: 2407.01489. URL: <https://arxiv.org/abs/2407.01489>.
- Xiang, Z. et al. (2024). *GuardAgent: Safeguard LLM Agents by a Guard Agent via Knowledge-Enabled Reasoning*. arXiv preprint. arXiv: 2406.09187. URL: <https://arxiv.org/abs/2406.09187>.
- Xie, T., D. Zhang, J. Chen, et al. (2024). *OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments*. NeurIPS 2024. arXiv: 2404.07972. URL: <https://arxiv.org/abs/2404.07972>.
- Xu et al. (2026a). *Agent Skills for Large Language Models: Architecture, Acquisition, Security, and the Path Forward*. arXiv preprint. arXiv: 2602.12430. URL: <https://arxiv.org/abs/2602.12430>.
- (2026b). *RoadmapBench: Evaluating Long-Horizon Agentic Software Development Across Version Upgrades*. arXiv preprint. arXiv: 2605.15846. URL: <https://arxiv.org/abs/2605.15846>.

- Xu et al. (2026c). *TokenPilot: Cache-Efficient Context Management for LLM Agents*. arXiv preprint. arXiv: 2606.17016. URL: <https://arxiv.org/abs/2606.17016>.
- Xu, B., Z. Peng, B. Lei, S. Mukherjee, Y. Liu, and D. Xu (2023). *ReWOO: Decoupling Reasoning from Observations for Efficient Augmented Language Models*. arXiv. arXiv: 2305.18323. URL: <https://arxiv.org/abs/2305.18323>.
- Yan, W. (2025). *Don't Build Multi-Agents*. Cognition AI blog. Cognition AI engineering essay. URL: <https://cognition.com/blog/dont-build-multi-agents> (visited on 07/09/2026).
- Yang et al. (2026a). *Ares: Adaptive Reasoning Effort Selection for Efficient LLM Agents*. arXiv preprint. arXiv: 2603.07915. URL: <https://arxiv.org/abs/2603.07915>.
- (2026b). *FederatedSkill: Federated Learning for Agentic Skill Evolution*. arXiv preprint. arXiv: 2606.03143. URL: <https://arxiv.org/abs/2606.03143>.
- (2026c). *GUI-Libra: Training Native GUI Agents to Reason and Act with Action-aware Supervision and Partially Verifiable RL*. arXiv preprint. arXiv: 2602.22190. URL: <https://arxiv.org/abs/2602.22190>.
- (2026d). *SkillOpt: Executive Strategy for Self-Evolving Agent Skills*. arXiv preprint. arXiv: 2605.23904. URL: <https://arxiv.org/abs/2605.23904>.
- Yang, J., C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, and O. Press (2024a). *SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering*. NeurIPS 2024. arXiv: 2405.15793. URL: <https://arxiv.org/abs/2405.15793>.
- Yang, Y., H. Chai, Y. Song, et al. (2025). *A Survey of AI Agent Protocols*. arXiv: 2504.16736. URL: <https://arxiv.org/abs/2504.16736>.
- Yang, Z. et al. (2024b). *ReAct Meets ActRe: When Language Agents Enjoy Training Data Autonomy*. arXiv preprint. arXiv: 2403.14589. URL: <https://arxiv.org/abs/2403.14589>.
- Yao, S., N. Shinn, P. Razavi, and K. Narasimhan (2024). τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. arXiv. arXiv: 2406.12045. URL: <https://arxiv.org/abs/2406.12045>.
- Yao, S., D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan (2023a). *Tree of Thoughts: Deliberate Problem Solving with Large Language Models*. NeurIPS 2023. arXiv: 2305.10601. URL: <https://arxiv.org/abs/2305.10601>.
- Yao, S., J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao (2022). *ReAct: Synergizing Reasoning and Acting in Language Models*. arXiv preprint; ICLR 2023 (camera-ready version noted on arXiv page). arXiv: 2210.03629. URL: <https://arxiv.org/abs/2210.03629>.
- Yao, W. et al. (2023b). *Retroformer: Retrospective Large Language Agents with Policy Gradient Optimization*. arXiv preprint. arXiv: 2308.02151. URL: <https://arxiv.org/abs/2308.02151>.
- Ye, Z. (2026). *Set-shifting Behavioral Test for Harnessed Agents*. arXiv. arXiv: 2607.13396. URL: <https://arxiv.org/abs/2607.13396>.
- Yehudai, A., L. Eden, A. Li, G. Uziel, Y. Zhao, R. Bar-Haim, A. Cohan, and M. Shmueli-Scheuer (2025). *Survey on Evaluation of LLM-based Agents*. arXiv. arXiv: 2503.16416. URL: <https://arxiv.org/abs/2503.16416>.
- Yi et al. (2026). *Learning Agent-Compatible Context Management for Long-Horizon Tasks*. arXiv preprint. arXiv: 2605.30785. URL: <https://arxiv.org/abs/2605.30785>.
- Yi, H. and X. Song (2026). *Measuring Harness-Induced Belief Divergence in Multi-Step LLM Agents*. arXiv. arXiv: 2607.04528. URL: <https://arxiv.org/abs/2607.04528>.
- Yu et al. (2026a). *AdaptOrch: Task-Adaptive Multi-Agent Orchestration in the Era of LLM Performance Convergence*. arXiv preprint. arXiv: 2602.16873. URL: <https://arxiv.org/abs/2602.16873>.

- Yu et al. (2026b). *MetaResearcher: Scaling Deep Research via Self-Reflective Reinforcement Learning in Adversarial Virtual Environments*. arXiv preprint. arXiv: 2606.19893. URL: <https://arxiv.org/abs/2606.19893>.
- Yu, W., L. Zhou, X. Dong, S. S. Barath, et al. (2026c). *When Do Multi-Agent Systems Help? An Information Bottleneck Perspective*. arXiv. arXiv: 2607.16133. URL: <https://arxiv.org/abs/2607.16133>.
- Yu, Z., X. Xie, W. Yao, C. Wang, et al. (2026d). *SkillAdaptor: Self-Adapting Skills for LLM Agents from Trajectories*. arXiv. arXiv: 2606.01311. URL: <https://arxiv.org/abs/2606.01311>.
- Yuan et al. (2026). *OSWorldz.o: Benchmarking Computer Use Agents on Long-Horizon Real-World Tasks*. arXiv preprint. arXiv: 2606.29537. URL: <https://arxiv.org/abs/2606.29537>.
- Yuan, T., Z. He, L. Dong, Y. Wang, R. Zhao, T. Xia, L. Xu, B. Zhou, F. Li, Z. Zhang, R. Wang, and G. Liu (2024). *R-Judge: Benchmarking Safety Risk Awareness for LLM Agents*. EMNLP 2024 Findings. arXiv: 2401.10019. URL: <https://arxiv.org/abs/2401.10019>.
- Yue, Y. et al. (2025). *Does Reinforcement Learning Really Incentivize Reasoning Capacity in LLMs Beyond the Base Model?* arXiv preprint; NeurIPS 2025. arXiv: 2504.13837. URL: <https://arxiv.org/abs/2504.13837>.
- Zaharia, M., O. Khattab, L. Chen, J. Q. Davis, H. Miller, C. Potts, J. Zou, M. Carbin, J. Frankle, N. Rao, and A. Ghodsi (2024). *The Shift from Models to Compound AI Systems*. Berkeley Artificial Intelligence Research (BAIR) blog. URL: <https://bair.berkeley.edu/blog/2024/02/18/compound-ai-systems/> (visited on 07/09/2026).
- Zeng et al. (2026). *Group of Skills: Group-Structured Skill Retrieval for Agent Skill Libraries*. arXiv preprint. arXiv: 2605.06978. URL: <https://arxiv.org/abs/2605.06978>.
- Zhai et al. (2026). *Revisable by Design: A Theory of Streaming LLM Agent Execution*. arXiv preprint. arXiv: 2604.23283. URL: <https://arxiv.org/abs/2604.23283>.
- Zhan, Q., Z. Liang, Z. Ying, and D. Kang (2024). *InjecAgent: Benchmarking Indirect Prompt Injections in Tool-Integrated Large Language Model Agents*. ACL 2024 Findings. arXiv: 2403.02691. URL: <https://arxiv.org/abs/2403.02691>.
- Zhang et al. (2025a). *AgentRL: Scaling Agentic Reinforcement Learning with a Multi-Turn, Multi-Task Framework*. arXiv preprint. arXiv: 2510.04206. URL: <https://arxiv.org/abs/2510.04206>.
- (2026a). *AgentV-RL: Scaling Reward Modeling with Agentic Verifier*. arXiv preprint. arXiv: 2604.16004. URL: <https://arxiv.org/abs/2604.16004>.
- (2026b). *Co-Evolving Skill Generation and Policy Optimization*. arXiv preprint. arXiv: 2606.08755. URL: <https://arxiv.org/abs/2606.08755>.
- (2026c). *From Reasoning to Agentic: Credit Assignment in Reinforcement Learning for Large Language Models*. arXiv preprint. arXiv: 2604.09459. URL: <https://arxiv.org/abs/2604.09459>.
- (2026d). *Library Drift: Diagnosing and Fixing a Silent Failure Mode in Self-Evolving LLM Skill Libraries*. arXiv preprint. arXiv: 2605.19576. URL: <https://arxiv.org/abs/2605.19576>.
- (2026e). *Reinforcement Learning for LLM-based Multi-Agent Systems through Orchestration Traces*. arXiv preprint. arXiv: 2605.02801. URL: <https://arxiv.org/abs/2605.02801>.
- (2026f). *Self-Evolving World Models for LLM Agent Planning*. arXiv preprint. arXiv: 2606.30639. URL: <https://arxiv.org/abs/2606.30639>.
- (2026g). *Workflow-to-Skill: Skill Creation via Routing-Workflow-Semantics-Attachments Decomposition*. arXiv preprint. arXiv: 2606.06893. URL: <https://arxiv.org/abs/2606.06893>.
- Zhang, G., H. Geng, X. Yu, et al. (2025b). *The Landscape of Agentic Reinforcement Learning for LLMs: A Survey*. arXiv: 2509.02547. URL: <https://arxiv.org/abs/2509.02547>.

- Zhang, H., J. Huang, K. Mei, Y. Yao, Z. Wang, C. Zhan, H. Wang, and Y. Zhang (2024a). *Agent Security Bench (ASB): Formalizing and Benchmarking Attacks and Defenses in LLM-based Agents*. arXiv. arXiv: 2410.02644. URL: <https://arxiv.org/abs/2410.02644>.
- Zhang, J., J. Hu, and Z. Ou (2026h). *GELS: A Generation-Evaluation-Improvement Loop of Agent Skills for Long-Form Article Generation*. arXiv. arXiv: 2607.11503. URL: <https://arxiv.org/abs/2607.11503>.
- Zhang, L., X. Chen, and H. Feng (2026i). *Behavioral Controllability of Agentic Models for Information Extraction: From Fixed Workflows to Reflective Agents*. arXiv. arXiv: 2607.15715. URL: <https://arxiv.org/abs/2607.15715>.
- Zhang, L. et al. (2024b). *Generative Verifiers: Reward Modeling as Next-Token Prediction*. arXiv preprint (GenRM). arXiv: 2408.15240. URL: <https://arxiv.org/abs/2408.15240>.
- Zhang, Y. et al. (2026j). *Stop Comparing LLM Agents Without Disclosing the Harness*. arXiv preprint. arXiv: 2605.23950. URL: <https://arxiv.org/abs/2605.23950>.
- Zhang, Z., X. Bo, C. Ma, R. Li, X. Chen, Q. Dai, J. Zhu, Z. Dong, and J.-R. Wen (2024c). *A Survey on the Memory Mechanism of Large Language Model based Agents*. arXiv. arXiv: 2404.13501. URL: <https://arxiv.org/abs/2404.13501>.
- Zhao et al. (2025). *Parasites in the Toolchain: A Large-Scale Analysis of Attacks on the MCP Ecosystem*. arXiv preprint. arXiv: 2509.06572. URL: <https://arxiv.org/abs/2509.06572>.
- (2026a). *AEM: Adaptive Entropy Modulation for Multi-Turn Agentic Reinforcement Learning*. arXiv preprint. arXiv: 2605.00425. URL: <https://arxiv.org/abs/2605.00425>.
- (2026b). *ClawGuard: A Runtime Security Framework for Tool-Augmented LLM Agents Against Indirect Prompt Injection*. arXiv preprint. arXiv: 2604.11790. URL: <https://arxiv.org/abs/2604.11790>.
- (2026c). *FineVerify: Scaling Test-Time Compute with Fine-Grained Self-Verification for Agentic Search*. arXiv preprint. arXiv: 2606.00660. URL: <https://arxiv.org/abs/2606.00660>.
- (2026d). *Generative Skill Composition for LLM Agents*. arXiv preprint. arXiv: 2606.32025. URL: <https://arxiv.org/abs/2606.32025>.
- (2026e). *Rethinking Experience Utilization in Self-Evolving Language Model Agents*. arXiv preprint. arXiv: 2605.07164. URL: <https://arxiv.org/abs/2605.07164>.
- Zhao, A., D. Huang, Q. Xu, M. Lin, Y.-J. Liu, and G. Huang (2023). *Expel: LLM Agents Are Experiential Learners*. AAAI 2024. arXiv: 2308.10144. URL: <https://arxiv.org/abs/2308.10144>.
- Zhao, C., Y. Tan, S. He, Y. Wang, et al. (2026f). *Neural Procedural Memory: Empowering LLM Agents with Implicit Activation Steering*. arXiv. arXiv: 2606.29824. URL: <https://arxiv.org/abs/2606.29824>.
- Zhao, X., H. Li, S. Li, T. Zhao, et al. (2026g). *Failure as a Process: An Anatomy of CLI Coding Agent Trajectories*. arXiv. arXiv: 2607.09510. URL: <https://arxiv.org/abs/2607.09510>.
- Zheng et al. (2026). *ACRFence: Preventing Semantic Rollback Attacks in Agent Checkpoint-Restore*. arXiv preprint. arXiv: 2603.20625. URL: <https://arxiv.org/abs/2603.20625>.
- Zheng, B. et al. (2025a). *SkillWeaver: Web Agents Can Self-Improve by Discovering and Honing Skills*. arXiv preprint. arXiv: 2504.07079. URL: <https://arxiv.org/abs/2504.07079>.
- Zheng, L. et al. (2023a). *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*. arXiv preprint; NeurIPS 2023. arXiv: 2306.05685. URL: <https://arxiv.org/abs/2306.05685>.
- Zheng, L., R. Wang, X. Wang, and B. An (2023b). *Synapse: Trajectory-as-Exemplar Prompting with Memory for Computer Control*. ICLR 2024. arXiv: 2306.07863. URL: <https://arxiv.org/abs/2306.07863>.
- Zheng, Y. et al. (2025b). *DeepResearcher: Scaling Deep Research via Reinforcement Learning in Real-World Environments*. arXiv preprint. arXiv: 2504.03160. URL: <https://arxiv.org/abs/2504.03160>.

- Zhou et al. (2026a). *A Comprehensive Survey on Agent Skills: Taxonomy, Techniques, and Applications*. arXiv preprint. arXiv: 2605.07358. URL: <https://arxiv.org/abs/2605.07358>.
- (2026b). *Are We Ready For An Agent-Native Memory System?* arXiv preprint. arXiv: 2606.24775. URL: <https://arxiv.org/abs/2606.24775>.
- (2026c). *Exploring Agentic Tool-Calling Decisions via Uncertainty-Aligned Reinforcement Learning*. arXiv preprint. arXiv: 2606.06976. URL: <https://arxiv.org/abs/2606.06976>.
- (2026d). *Externalization in LLM Agents: A Unified Review of Memory, Skills, Protocols and Harness Engineering*. arXiv preprint. arXiv: 2604.08224. URL: <https://arxiv.org/abs/2604.08224>.
- (2026e). *From Shield to Target: Denial-of-Service Attacks on LLM-Based Agent Guardrails*. arXiv preprint. arXiv: 2606.14517. URL: <https://arxiv.org/abs/2606.14517>.
- Zhou, A., K. Yan, M. Shlapentokh-Rothman, H. Wang, and Y.-X. Wang (2024a). *Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models*. ICML 2024. arXiv: 2310.04406. URL: <https://arxiv.org/abs/2310.04406>.
- Zhou, D., N. Schärli, L. Hou, J. Wei, N. Scales, X. Wang, D. Schuurmans, C. Cui, O. Bousquet, Q. Le, and E. Chi (2022). *Least-to-Most Prompting Enables Complex Reasoning in Large Language Models*. ICLR 2023. arXiv: 2205.10625. URL: <https://arxiv.org/abs/2205.10625>.
- Zhou, S., F. Yue, Z. Hu, Y. Shen, et al. (2026f). *ToolVerse: Unlocking Massive Environments and Long-Horizon Tasks for Agentic Reinforcement Learning*. arXiv. arXiv: 2607.15660. URL: <https://arxiv.org/abs/2607.15660>.
- Zhou, S., F. F. Xu, H. Zhu, X. Zhou, R. Lo, A. Sridhar, X. Cheng, T. Ou, Y. Bisk, D. Fried, U. Alon, and G. Neubig (2024b). *WebArena: A Realistic Web Environment for Building Autonomous Agents*. ICLR 2024. arXiv: 2307.13854. URL: <https://arxiv.org/abs/2307.13854>.
- Zhu et al. (2025a). *Scaling Test-Time Compute for LLM Agents*. arXiv preprint. arXiv: 2506.12928. URL: <https://arxiv.org/abs/2506.12928>.
- Zhu, X. et al. (2023). *Ghost in the Minecraft: Generally Capable Agents for Open-World Environments via Large Language Models with Text-Based Knowledge and Memory*. arXiv preprint. arXiv: 2305.17144. URL: <https://arxiv.org/abs/2305.17144>.
- Zhu, Y. et al. (2025b). *Establishing Best Practices for Building Rigorous Agentic Benchmarks*. arXiv preprint; NeurIPS 2025 Datasets and Benchmarks (ABC checklist). arXiv: 2507.02825. URL: <https://arxiv.org/abs/2507.02825>.
- Zhuge, M., C. Zhao, D. Ashley, et al. (2024). *Agent-as-a-Judge: Evaluate Agents with Agents*. arXiv: 2410.10934. URL: <https://arxiv.org/abs/2410.10934>.